

CS186 Discussion #3

(Joins)

Keys

superkey

inst letters	course	first name	last name	grade
ab	cs186	Bob	Smith	B
ac	cs186	Joe	Hellerstein	A+
ab	cs162	Eric	Brewer	A
ad	cs160	Toby	Brown	C

Keys

superkey

inst letters	course	first name	last name	grade
ab	cs186	Bob	Smith	B
ac	cs186	Joe	Hellerstein	A+
ab	cs162	Eric	Brewer	A
ad	cs160	Toby	Brown	C

Keys

superkey
candidate key

inst letters	course	first name	last name	grade
ab	cs186	Bob	Smith	B
ac	cs186	Joe	Hellerstein	A+
ab	cs162	Eric	Brewer	A
ad	cs160	Toby	Brown	C

Keys

superkey
candidate key

inst letters	course	first name	last name	grade
ab	cs186	Bob	Smith	B
ac	cs186	Joe	Hellerstein	A+
ab	cs162	Eric	Brewer	A
ad	cs160	Toby	Brown	C

* This is only true if no two people have the same first and last name!

Last week...

SELECT S.name
FROM Students S, Grades G;

S.name	S.id	S.gpa	G.id	G.class	G.grade
Bob	1	3.7	1	cs186	C
Bob	1	3.7	1	cs164	A
Bob	1	3.7	3	cs70	B
Bob	1	3.7	4	cs61a	A
Bob	1	3.7	2	cs61c	D
Bob	1	3.7	4	cs170	A
Sue	2	2.9	1	cs186	C
Sue	2	2.9	1	cs164	A
Sue	2	2.9	3	cs70	B
Sue	2	2.9	4	cs61a	A

Last week...

```
SELECT S.name  
FROM Students S, Grades G  
WHERE S.id = G.id;
```

S.name	S.id	S.gpa	G.id	G.class	G.grade
Bob	1	3.7	1	cs186	C
Bob	1	3.7	1	cs164	A
Sue	2	2.9	2	cs61c	D
Ron	3	1.2	3	cs70	B
Al	4	4.0	4	cs61a	A
Al	4	4.0	4	cs170	A

```
SELECT S.name  
FROM Students S INNER JOIN Grades G  
ON S.id = G.id;
```


Students

name	id	gpa
Bob	1	3.7
Sue	2	2.9
Ron	3	1.2
Al	4	4.0
Sally	5	3.6
Bob	6	2.1
Joe	9	4.0

Grades

id	class	grade
1	cs186	C
1	cs164	A
3	cs70	B
4	cs61a	A
2	cs61c	D
4	cs170	A
7	cs61C	F

SELECT S.name
FROM Students S **INNER JOIN** Grades G
ON S.id = G.id;

Students

name	id	gpa
Bob	1	3.7
Sue	2	2.9
Ron	3	1.2
Al	4	4.0
Sally	5	3.6
Bob	6	2.1
Joe	9	4.0

Grades

id	class	grade
1	cs186	C
1	cs164	A
3	cs70	B
4	cs61a	A
2	cs61c	D
4	cs170	A
7	cs61C	F

SELECT S.name
FROM Students S **INNER JOIN** Grades G
ON S.id = G.id;

name	id	gpa	id	class	grade
Bob	1	3.7	1	cs186	C
Bob	1	3.7	1	cs164	A
Ron	3	1.2	3	cs70	B
Al	4	4.0	4	cs61a	A
Sue	2	2.9	2	cs61c	D
Al	4	4.0	4	cs170	A

SELECT S.name
FROM Students S **LEFT OUTER JOIN** Grades G
ON S.id = G.id;

name	id	gpa	id	class	grade
Bob	1	3.7	1	cs186	C
Bob	1	3.7	1	cs164	A
Ron	3	1.2	3	cs70	B
Al	4	4.0	4	cs61a	A
Sue	2	2.9	2	cs61c	D
Al	4	4.0	4	cs170	A
Sally	5	3.6			
Bob	6	2.1			
Joe	9	4.0			

SELECT S.name
FROM Students S **RIGHT OUTER JOIN** Grades G
ON S.id = G.id;

name	id	gpa	id	class	grade
Bob	1	3.7	1	cs186	C
Bob	1	3.7	1	cs164	A
Ron	3	1.2	3	cs70	B
Al	4	4.0	4	cs61a	A
Sue	2	2.9	2	cs61c	D
Al	4	4.0	4	cs170	A
			7	cs61C	F

SELECT S.name
FROM Students S **FULL OUTER JOIN** Grades G
ON S.id = G.id;

name	id	gpa	id	class	grade
Bob	1	3.7	1	cs186	C
Bob	1	3.7	1	cs164	A
Ron	3	1.2	3	cs70	B
Al	4	4.0	4	cs61a	A
Sue	2	2.9	2	cs61c	D
Al	4	4.0	4	cs170	A
Sally	5	3.6			
Bob	6	2.1			
Joe	9	4.0			
			7	cs61C	F

Simplifying Subqueries

```
SELECT sname
```

```
FROM
```

```
    (SELECT sid
```

```
        FROM Reserves
```

```
    EXCEPT
```

```
        (SELECT sid
```

```
            FROM
```

```
                (SELECT Reserves.sid, PinkBoats.bid
```

```
                FROM Reserves,
```

```
                    (SELECT bid
```

```
                    FROM Boats
```

```
                    WHERE color='pink') PinkBoats
```

```
                EXCEPT SELECT sid, bid
```

```
                FROM Reserves)))
```

```
R, Sailors S
```

```
WHERE R.sid = S.sid;
```

Simplifying Subqueries

```
WITH Res(sid, bid) AS (SELECT Reserves.sid, PinkBoats.bid
                        FROM Reserves,
                             (SELECT bid
                              FROM Boats
                              WHERE color='pink') PinkBoats
                        EXCEPT SELECT sid, bid
                        FROM Reserves))
```

```
SELECT sname
FROM
    (SELECT sid
     FROM Reserves
     EXCEPT
        (SELECT sid
         FROM Res))
R, Sailors S
WHERE R.sid = S.sid;
```


Worksheet: Advanced SQL

Questions: 1, 3, 5

Find all album id's and names for every artist active since 2000 or later. If an artist does not have any albums, you should still include the artist's information in your output.

Find all album id's and names for every artist active since 2000 or later. If an artist does not have any albums, you should still include the artist's information in your output.

```
SELECT Ar.artist_id, Ar.artist_name,  
Al.album_id, Al.album_name
```

```
FROM Artists Ar LEFT OUTER JOIN Albums  
Al
```

```
ON Ar.artist_id=Al.artist_id
```

```
WHERE Ar.first_year_active >= 2000;
```

Find the id and name for each artist who has albums of genre “pop” and “rock”.

Find the id and name for each artist who has albums of genre “pop” and “rock”.

```
SELECT Ar.artist_id, Ar.artist_name
```

```
FROM Albums A11 INNER JOIN Albums A12 ON  
A11.artist_id=A12.artist_id INNER JOIN  
Artists Ar ON  
A12.artist_id=Ar.artist_id;
```

```
WHERE A11.genre="pop" AND  
A12.genre="rock"
```

Find all artists who released songs in 2014 or released songs that spent more than 10 weeks in the top 40.

Find all artists who released songs in 2014 or released songs that spent more than 10 weeks in the top 40.

```
WITH Temp(artist_id) AS
    (SELECT Al.artist_id
     FROM Albums Al
     WHERE Al.year_released=2014
     UNION
     SELECT Al.artist_id
     FROM Albums Al
     INNER JOIN Songs S ON Al.album_id=S.album_id
     WHERE S.weeks_in_top_40>10)
SELECT Ar.artist_id, Ar.artist_name
FROM Artists Ar
INNER JOIN Temp T ON Ar.artist_id=T.artist_id;
```


Join Algorithms

(Slides courtesy of CS186 Fall 2013)


```
SELECT * FROM Sailors S, Reserves R
WHERE S.sid = R.sid;
```

Visualizations

Sailors



Visualizations

Sailors



Visualizations

Sailors

Record 1
Record 2
Record 3
Record 4
Record 5

Page 2

Page 3

Page 4

Visualizations

Sailors

Reserves

Record 1

Record 2

Record 3




Record 4

Record 5

Page 2

Page 3

Page 4



Simple Nested Loops Join

Sailors

Reserves

Key idea:

Take each record of S and match it with each record of R.

Steps:

1. Get tuple of S.
2. Iterate through each tuple in R.

Simple Nested Loops Join

Sailors

Reserves

(name = Bob, sid = 1)

Key idea:

Take each record of S and match it with each record of R.


Steps:

1. Get tuple of S.
2. Iterate through each tuple in R.

Simple Nested Loops Join

Sailors

Reserves



Key idea:

Take each record of S and match it with each record of R.

Steps:

1. Get tuple of S.
2. Iterate through each tuple in R.

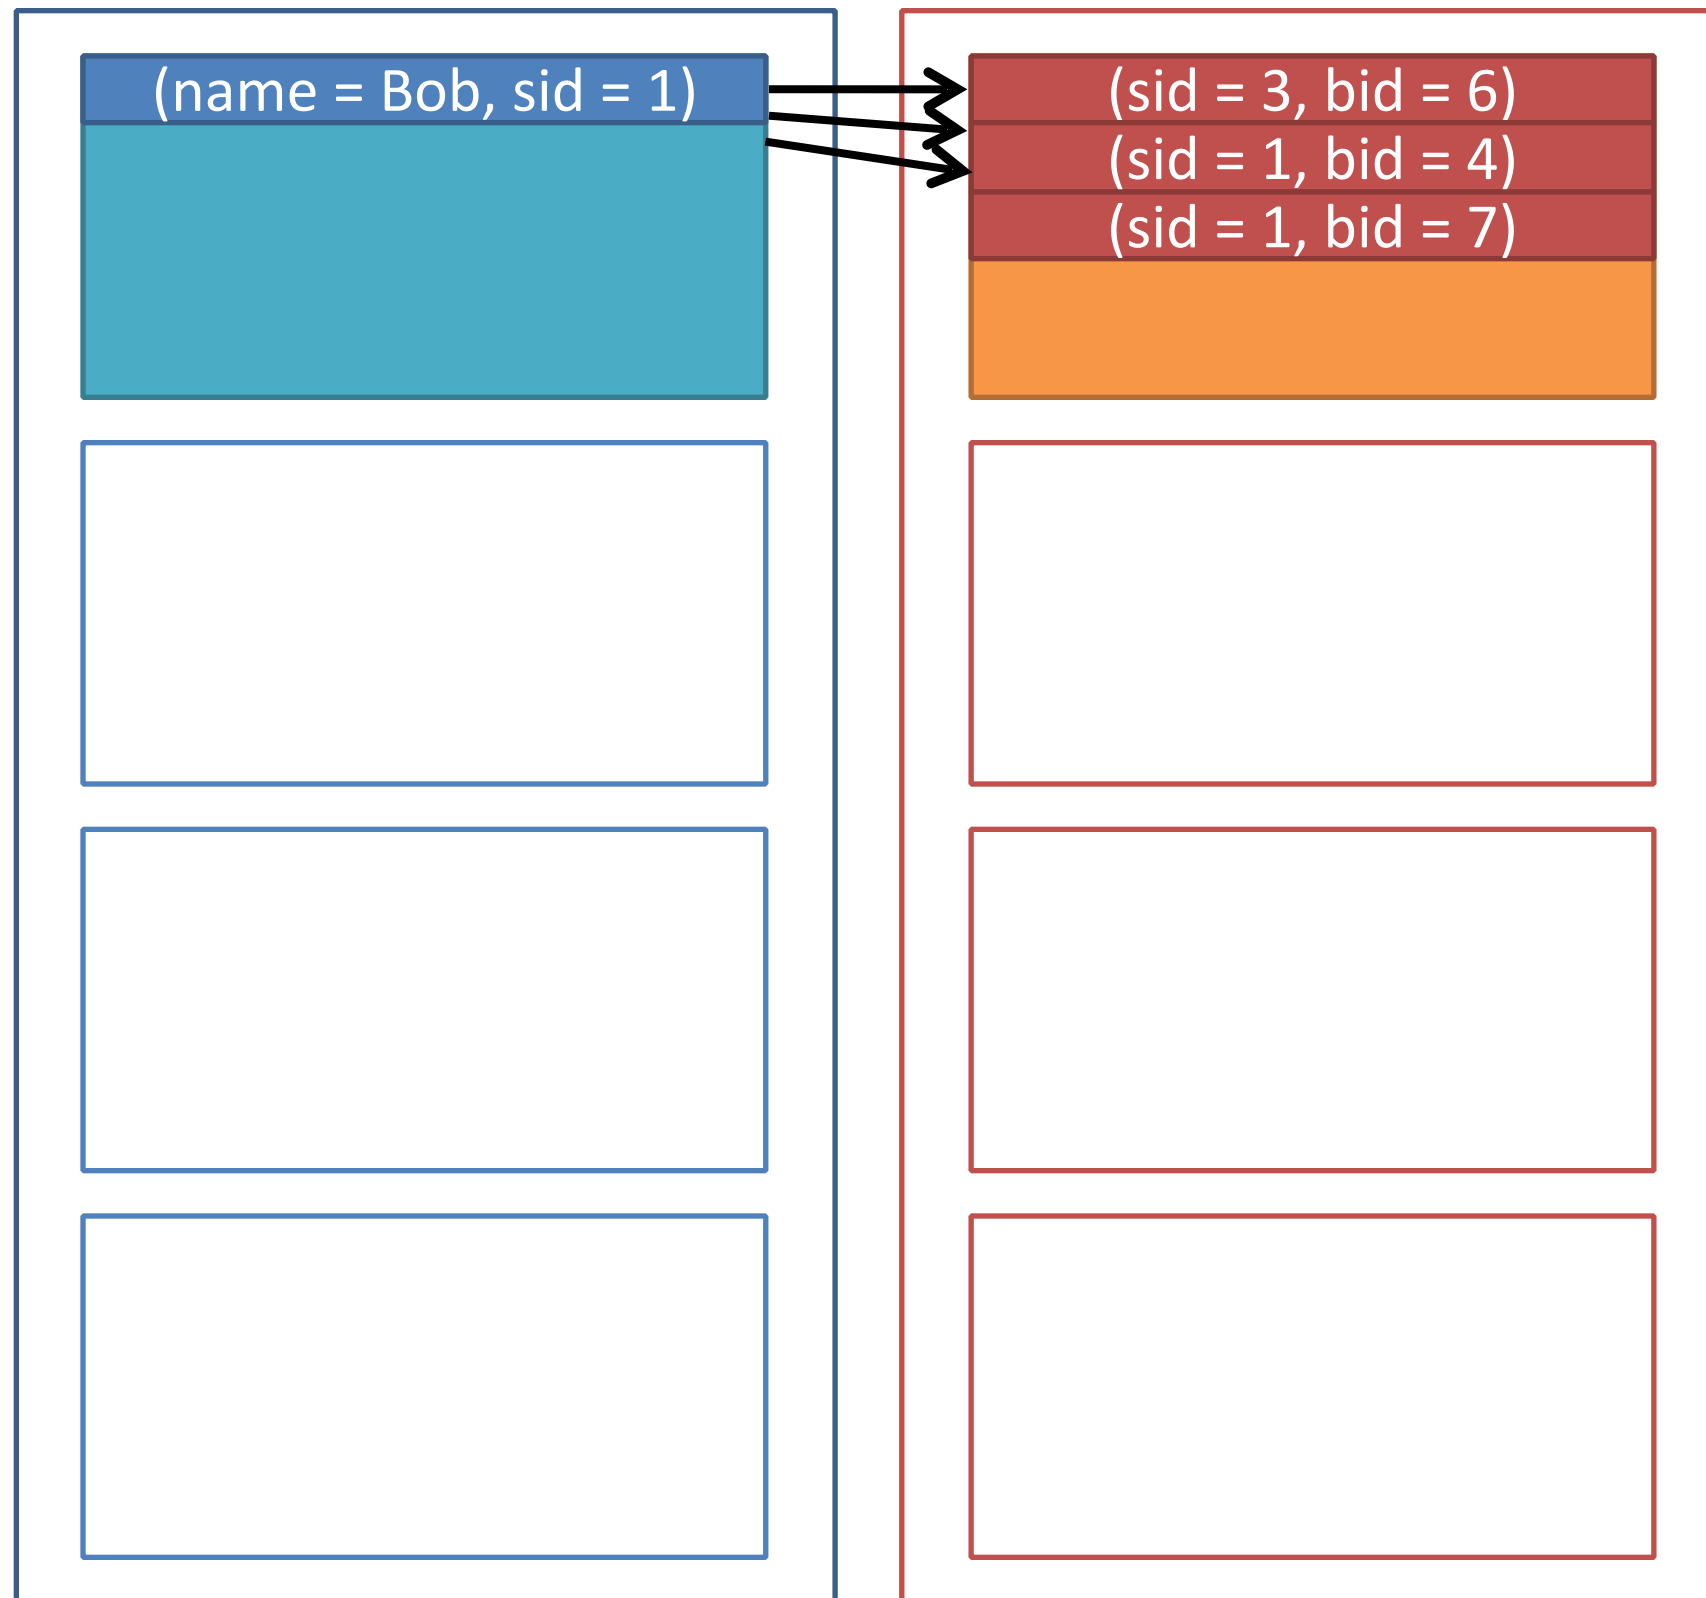
Output:

`(name = Bob, sid = 1, bid = 4)`

Simple Nested Loops Join

Sailors

Reserves



Key idea:

Take each record of S and match it with each record of R.

Steps:

1. Get tuple of S.
2. Iterate through each tuple in R.

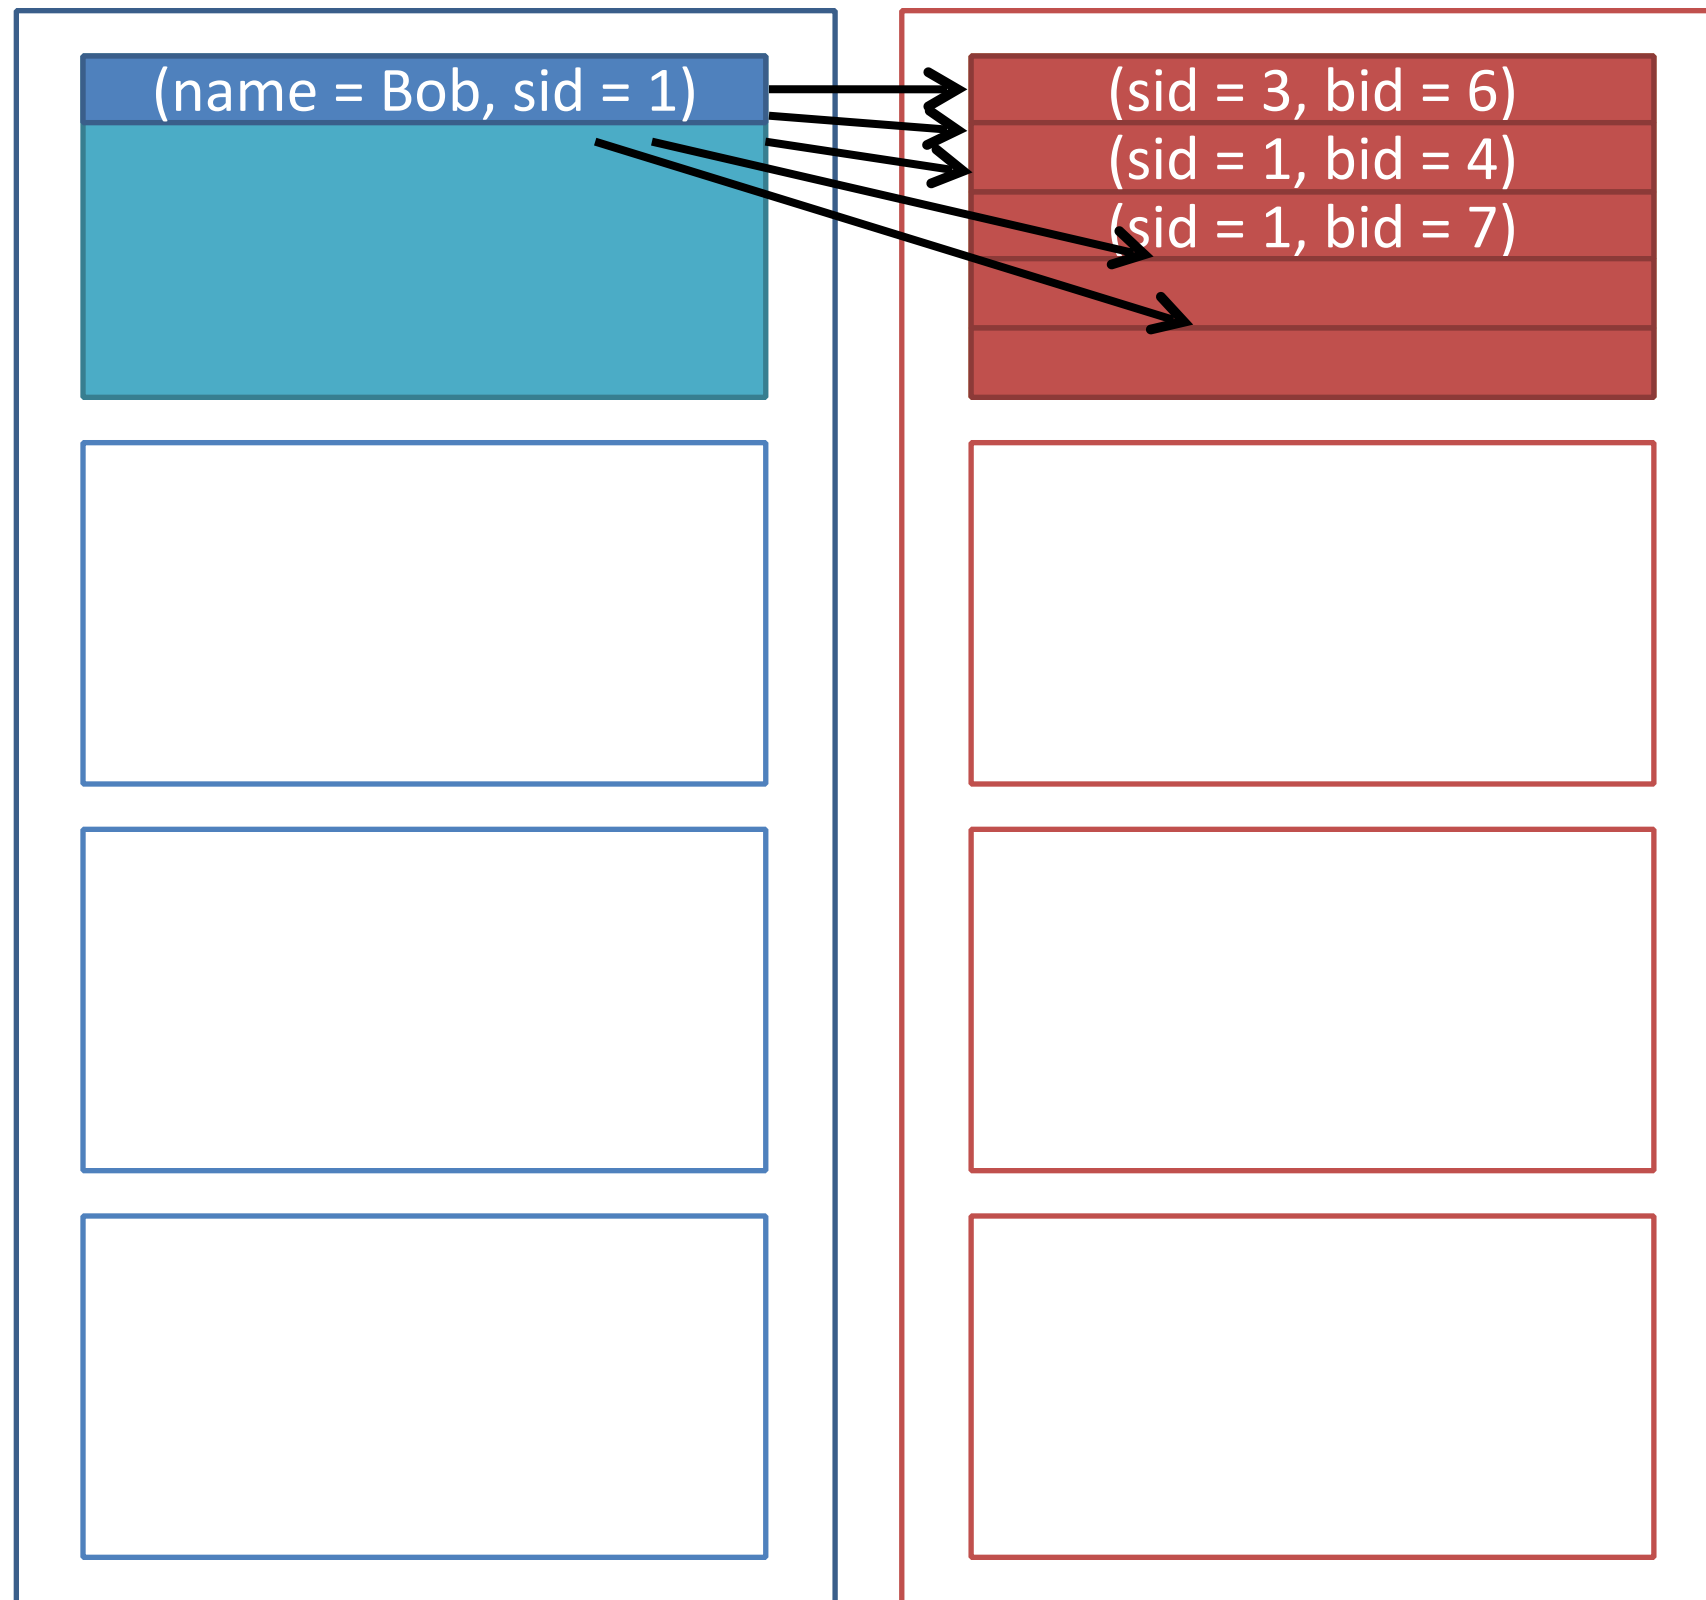
Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)

Simple Nested Loops Join

Sailors

Reserves



Key idea:

Take each record of S and match it with each record of R.

Steps:

1. Get tuple of S.
2. Iterate through each tuple in R.

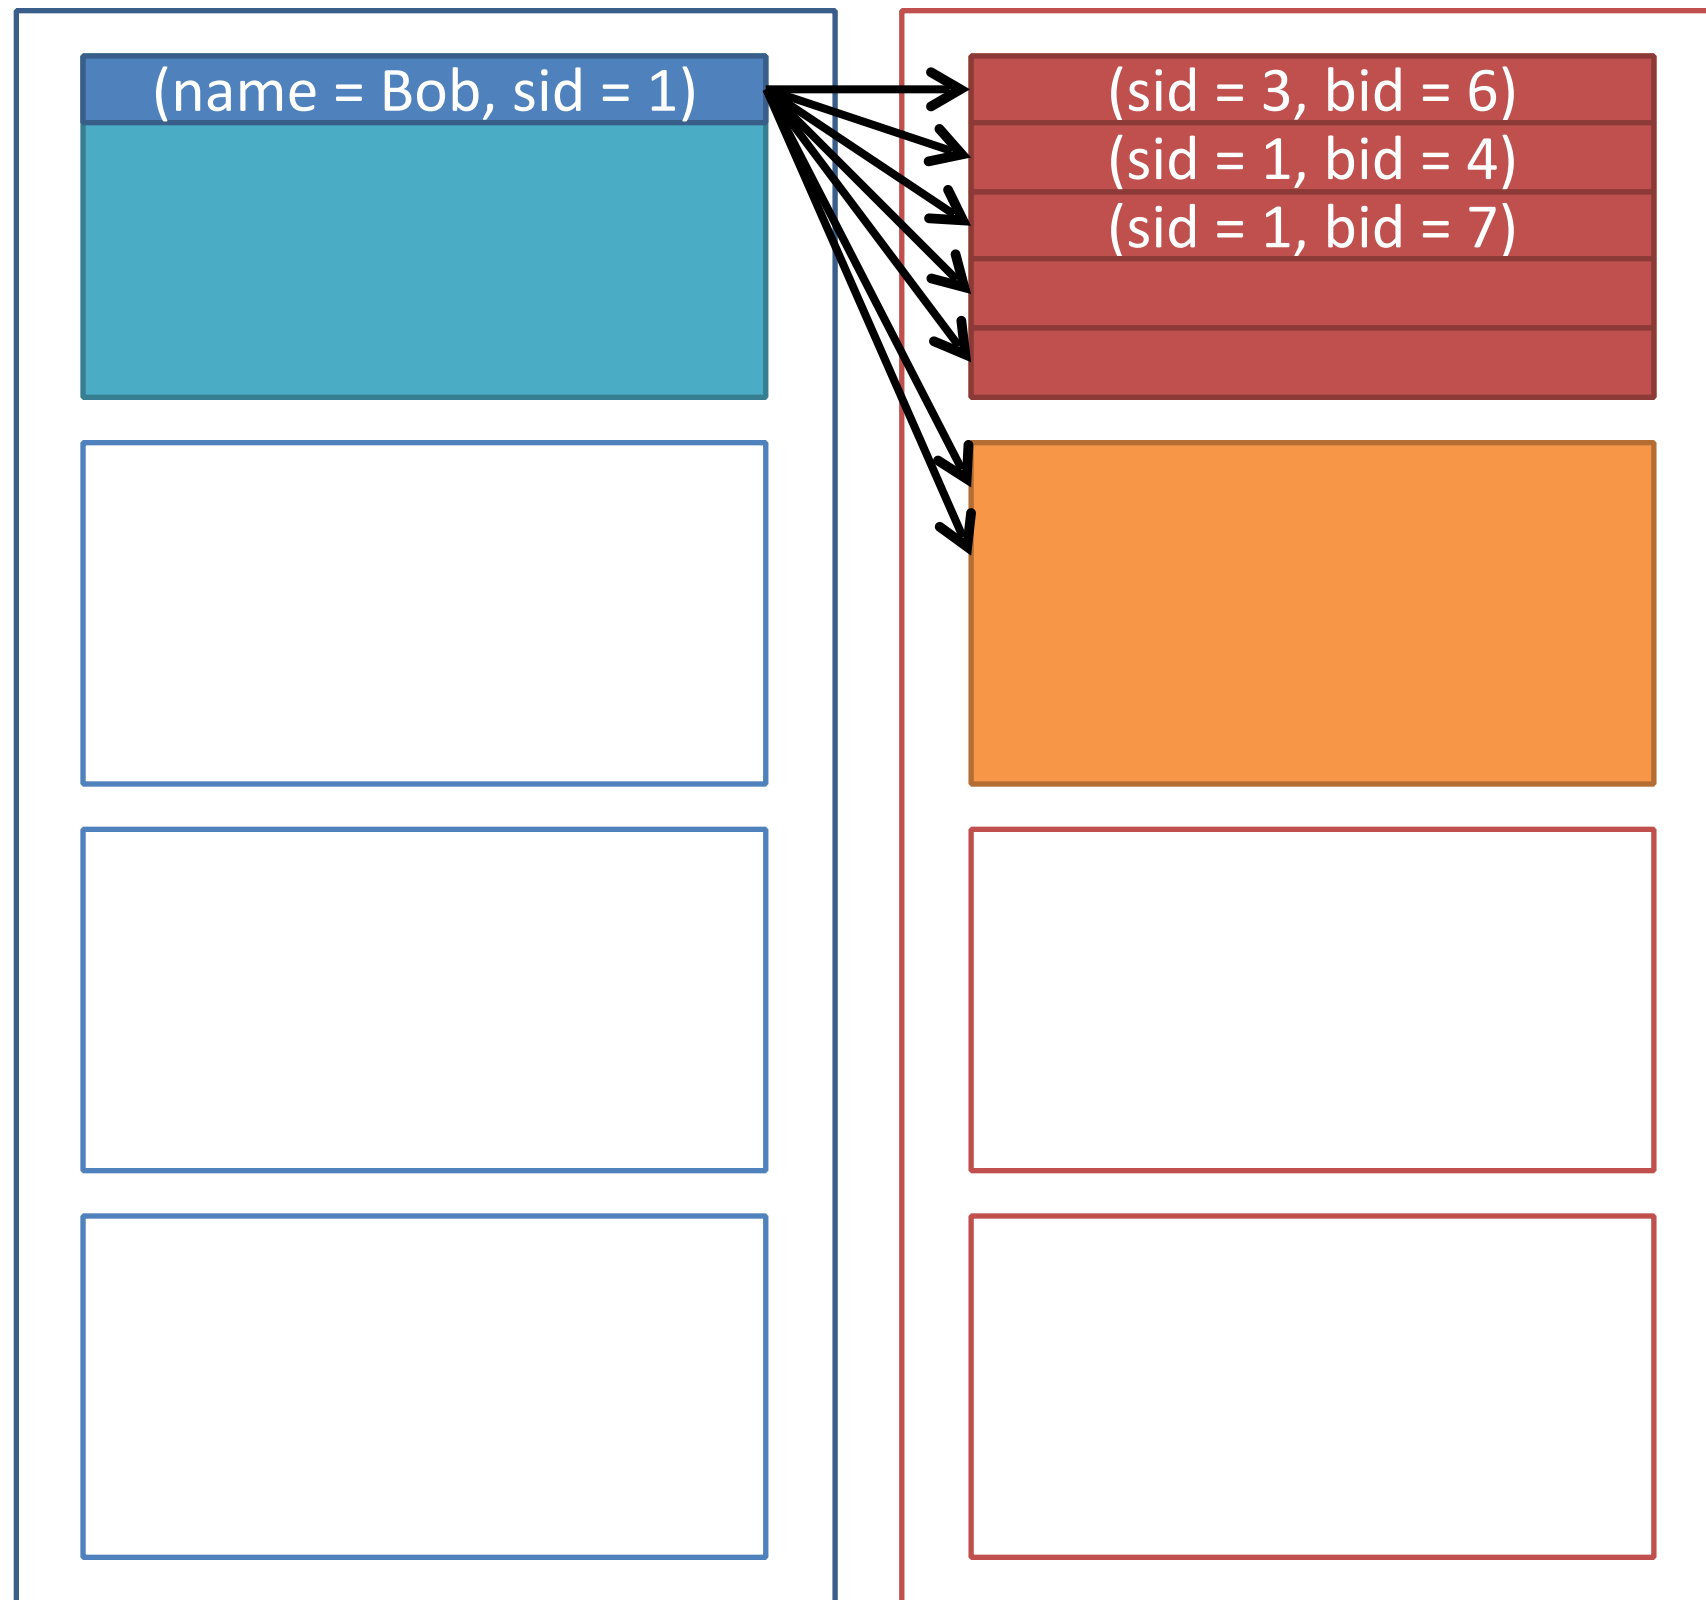
Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)

Simple Nested Loops Join

Sailors

Reserves



Key idea:

Take each record of S and match it with each record of R.

Steps:

1. Get tuple of S.
2. Iterate through each tuple in R.

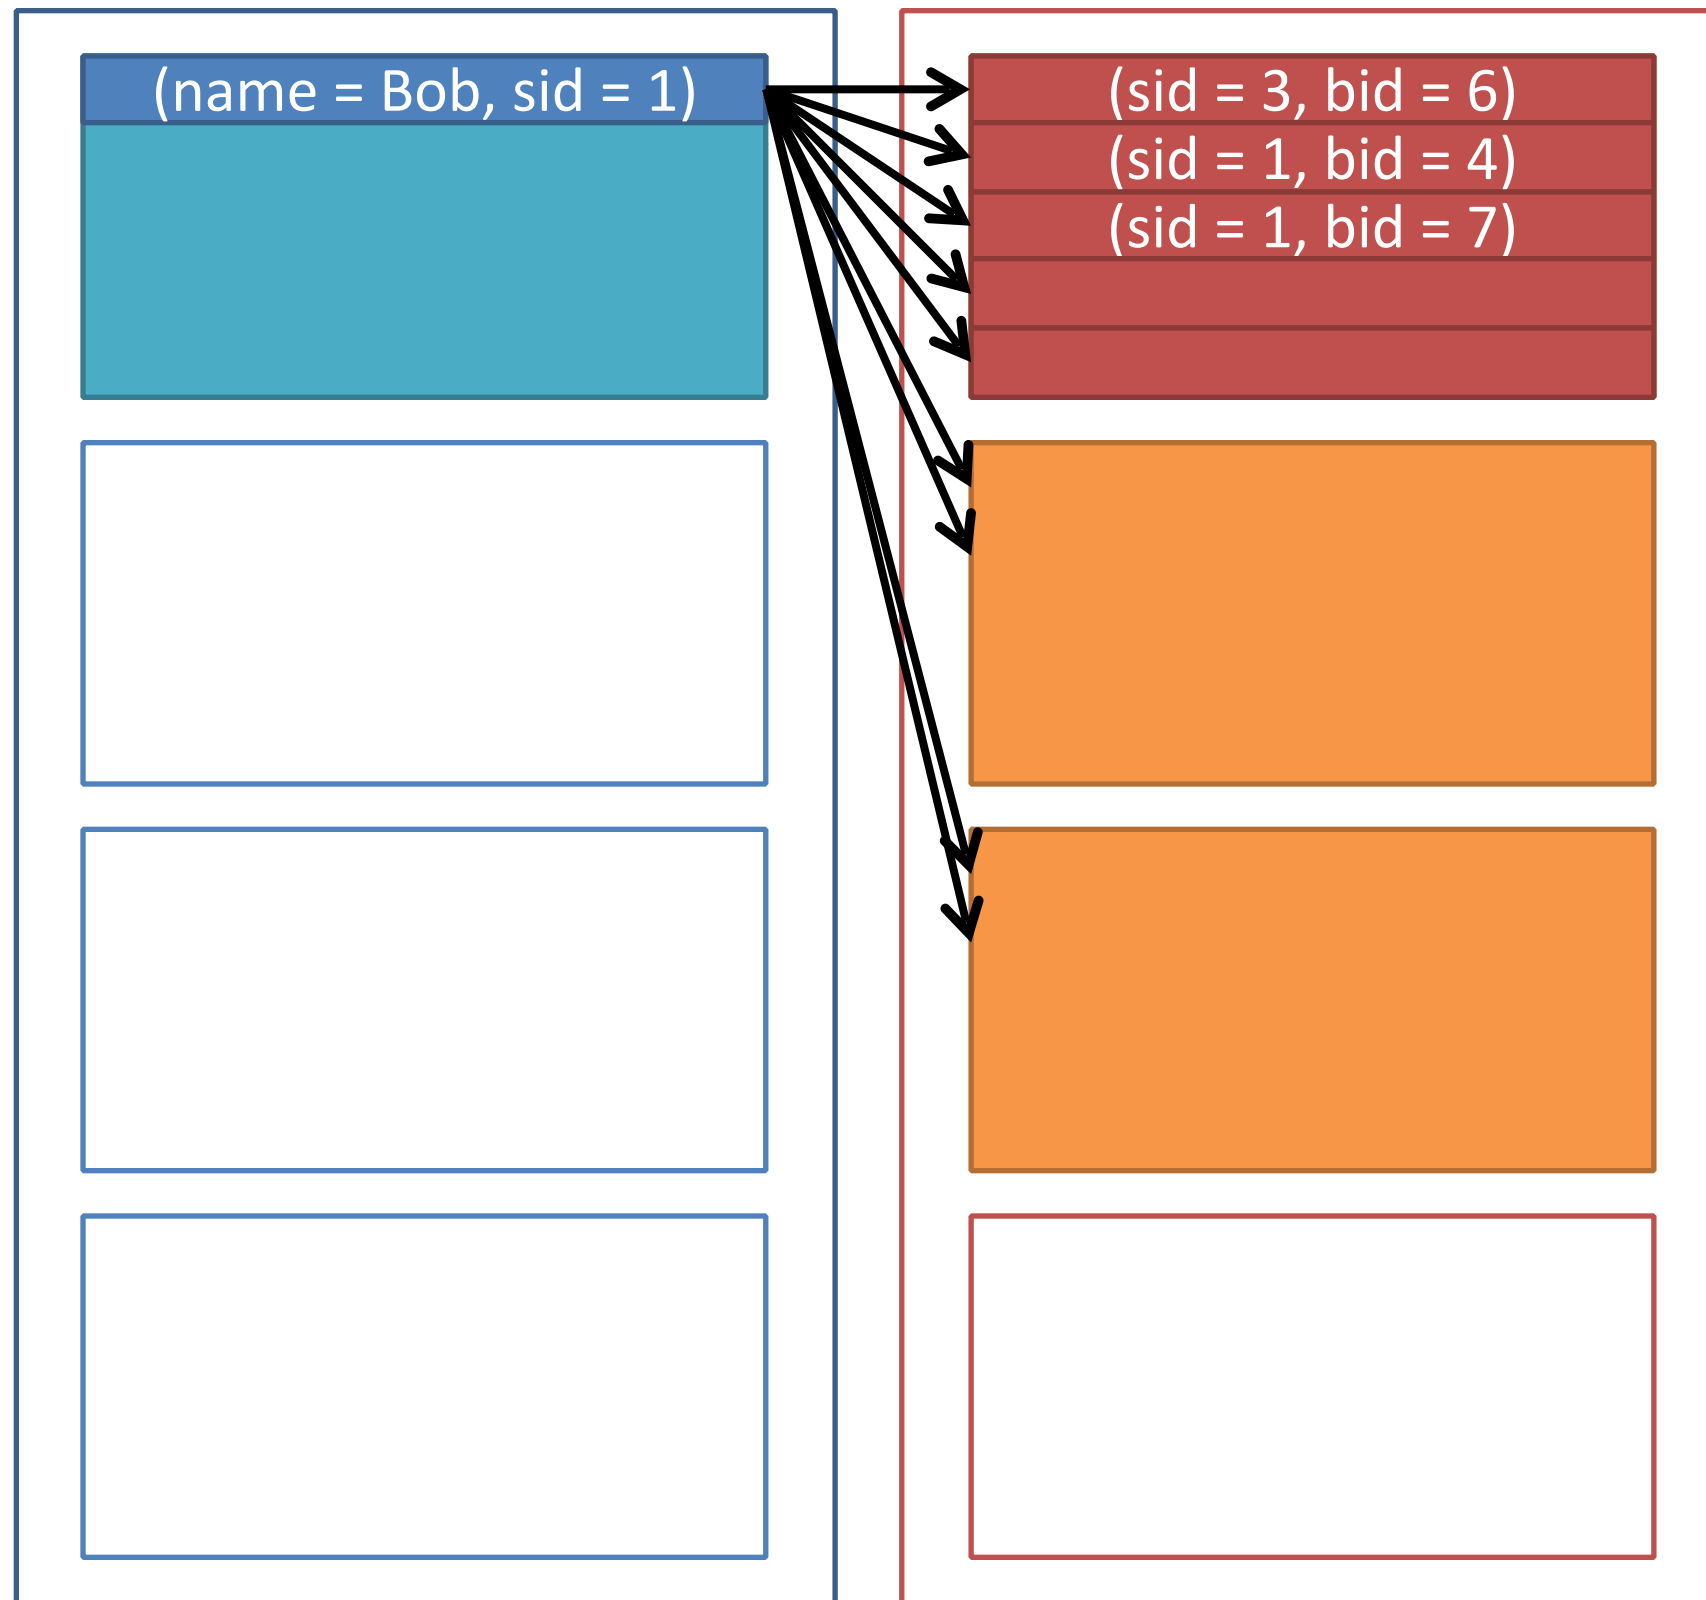
Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)

Simple Nested Loops Join

Sailors

Reserves



Key idea:

Take each record of S and match it with each record of R.

Steps:

1. Get tuple of S.
2. Iterate through each tuple in R.


Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)

Simple Nested Loops Join

Sailors

Reserves



Key idea:

Take each record of S and match it with each record of R.

Steps:

1. Get tuple of S.
2. Iterate through each tuple in R.

Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)

Simple Nested Loops Join

Sailors

Reserves

Key idea:

Take each record of S and match it with each record of R.

Steps:

1. Get tuple of S.
2. Iterate through each tuple in R.

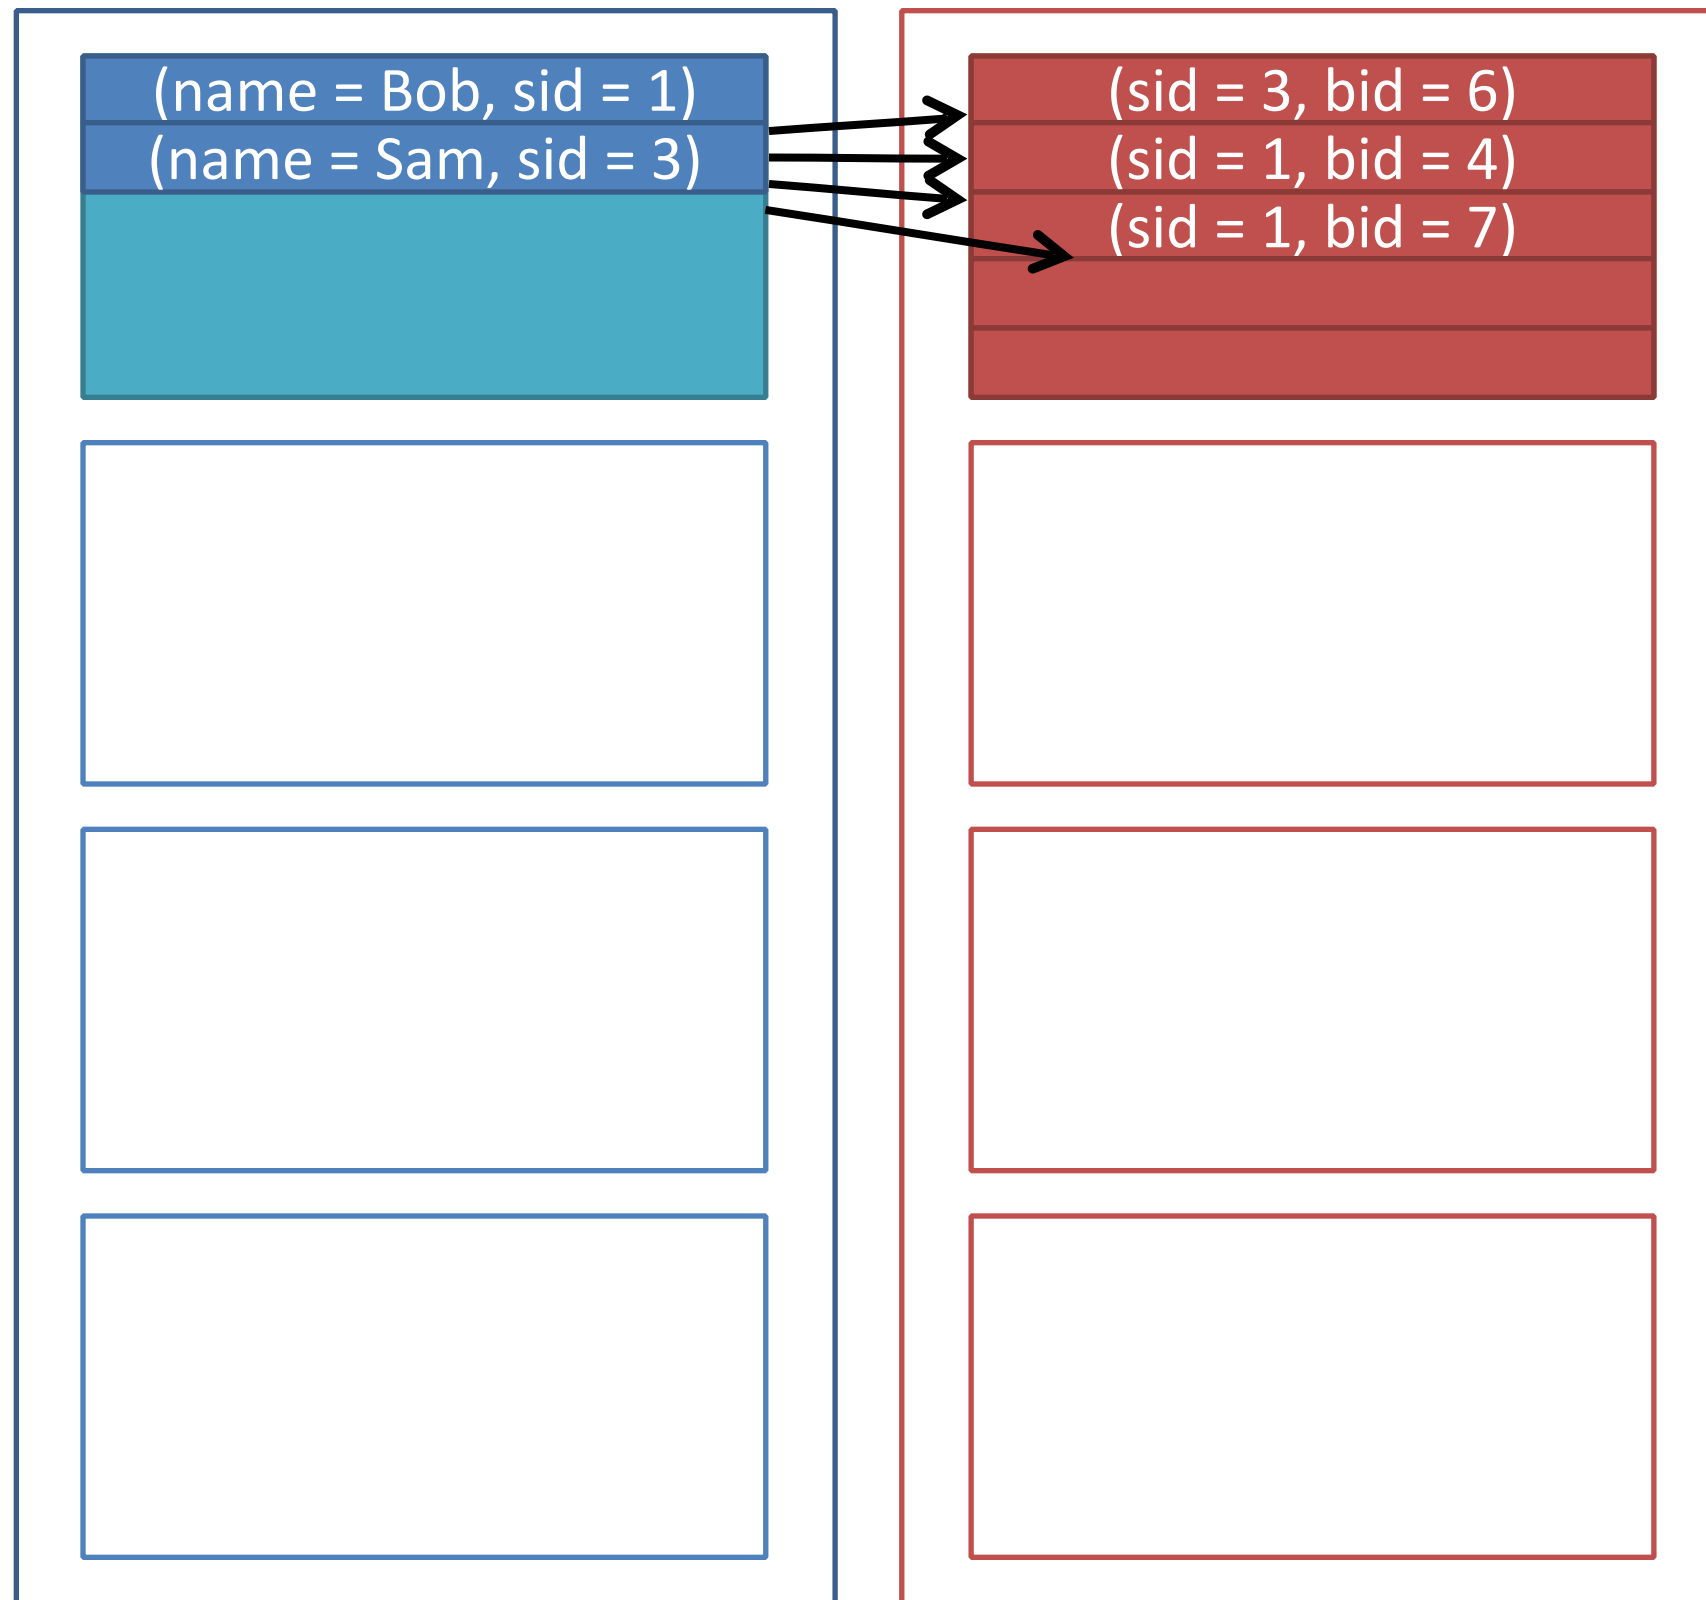
Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)

Simple Nested Loops Join

Sailors

Reserves



Key idea:

Take each record of S and match it with each record of R.

Steps:

1. Get tuple of S.
2. Iterate through each tuple in R.

Output:

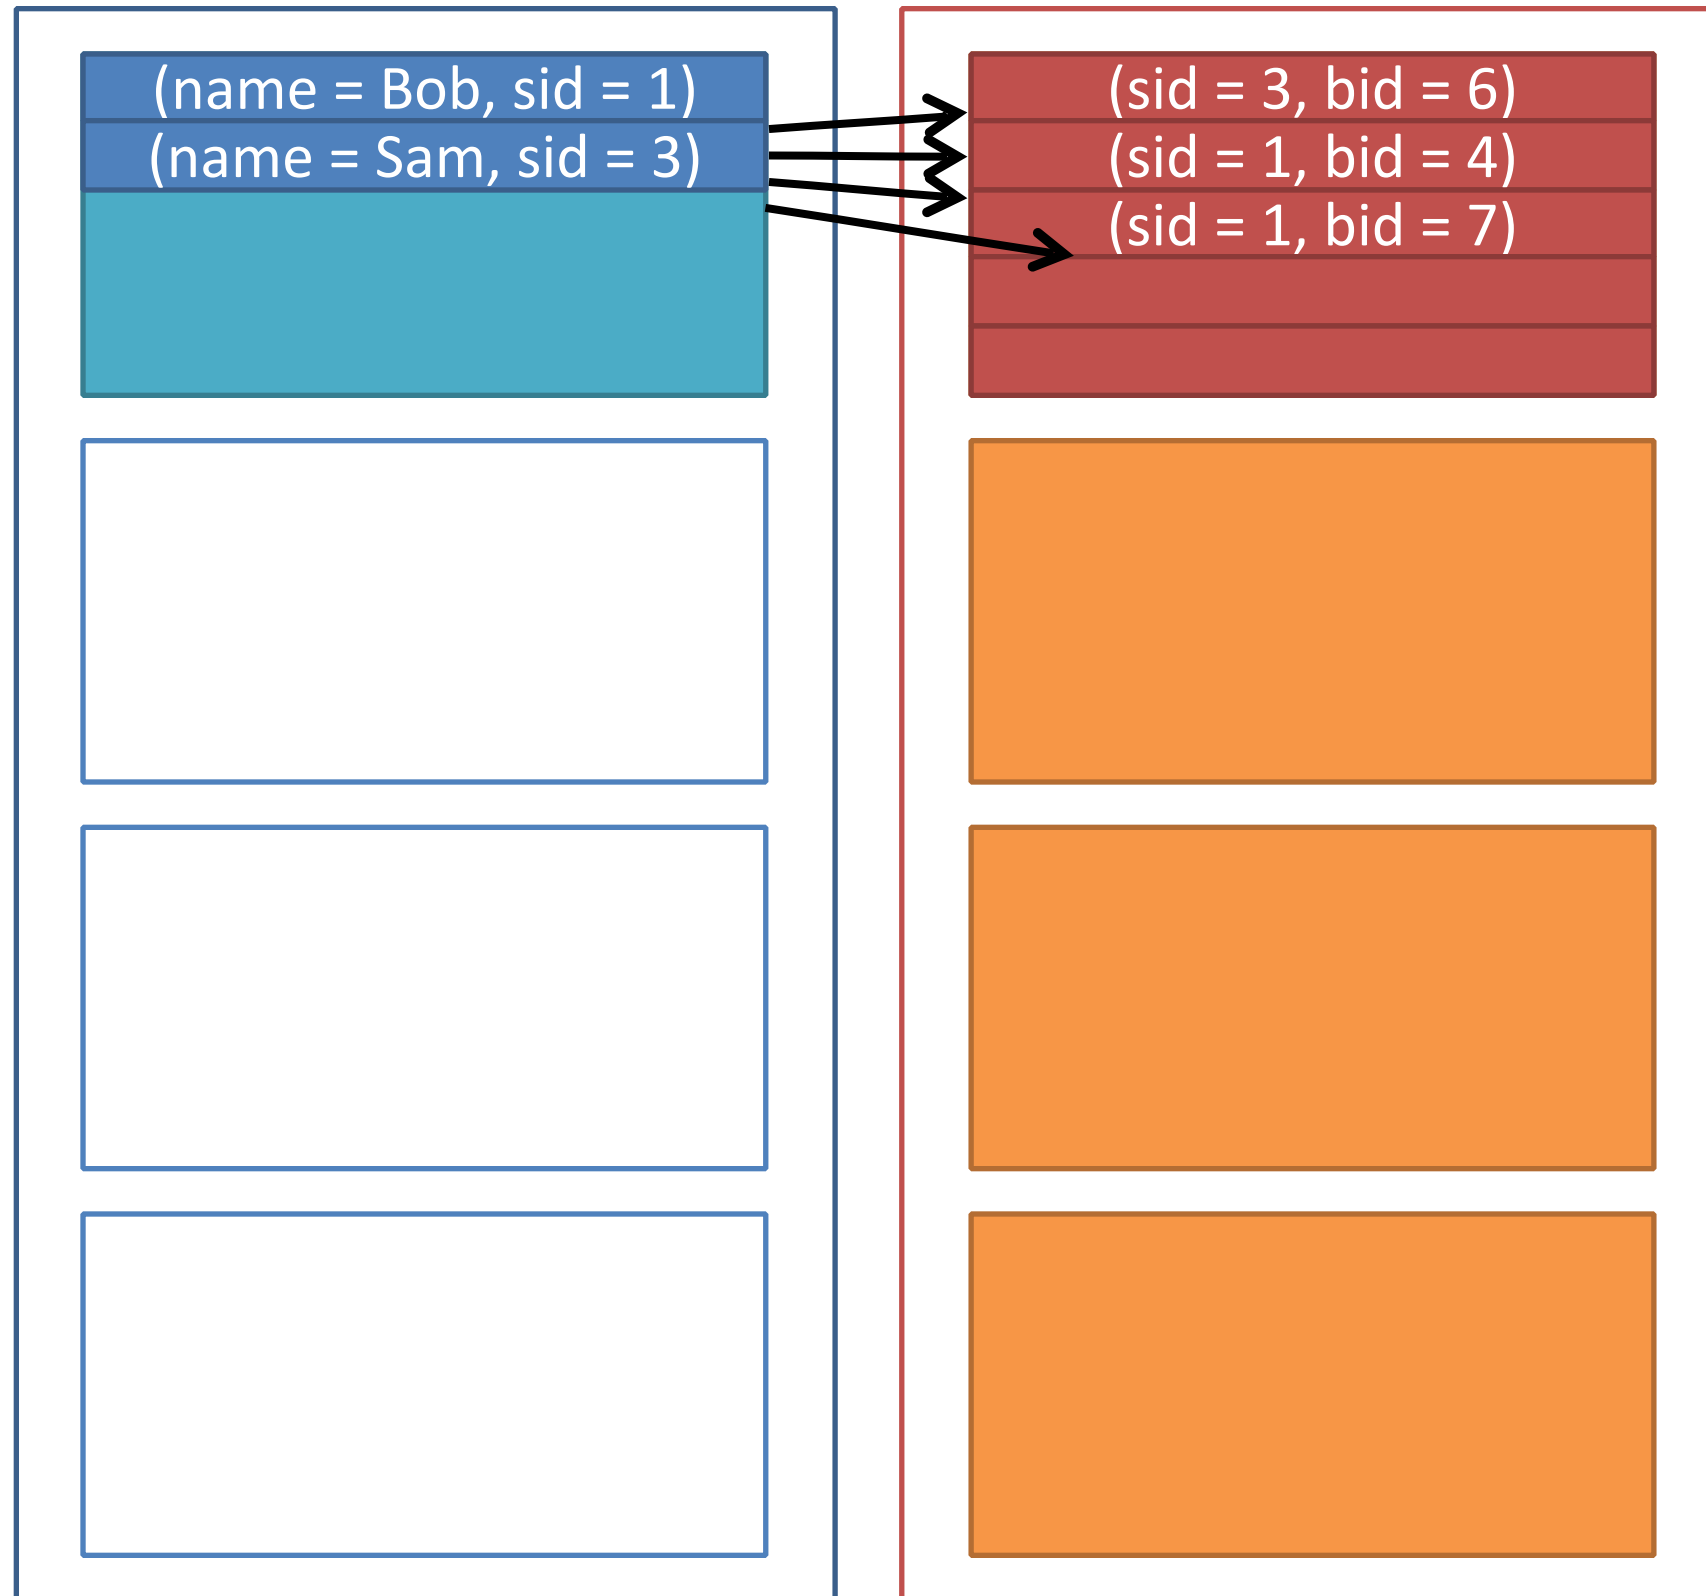
(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)
(name = Sam, sid = 3, bid = 6)

Notation: $[S]$ == “# pages in S ” ; $|S|$ == “# tuples in S ”

Simple Nested Loops Join

Sailors

Reserves



Key idea:

Take each record of S and match it with each record of R .

Steps:

1. Get tuple of S .
2. Iterate through each tuple in R .



I/Os:

$$[S] + |S| * [R]$$

Page-Oriented Nested Loops Join

Sailors

Reserves



Key idea:

Take each page of S and match with each page of R.



Steps:

1. Get page of S.
2. Iterate through each page in R.
3. Compare tuples in each.

Page-Oriented Nested Loops Join

Sailors

Reserves



Key idea:

Take each page of S and match with each page of R.

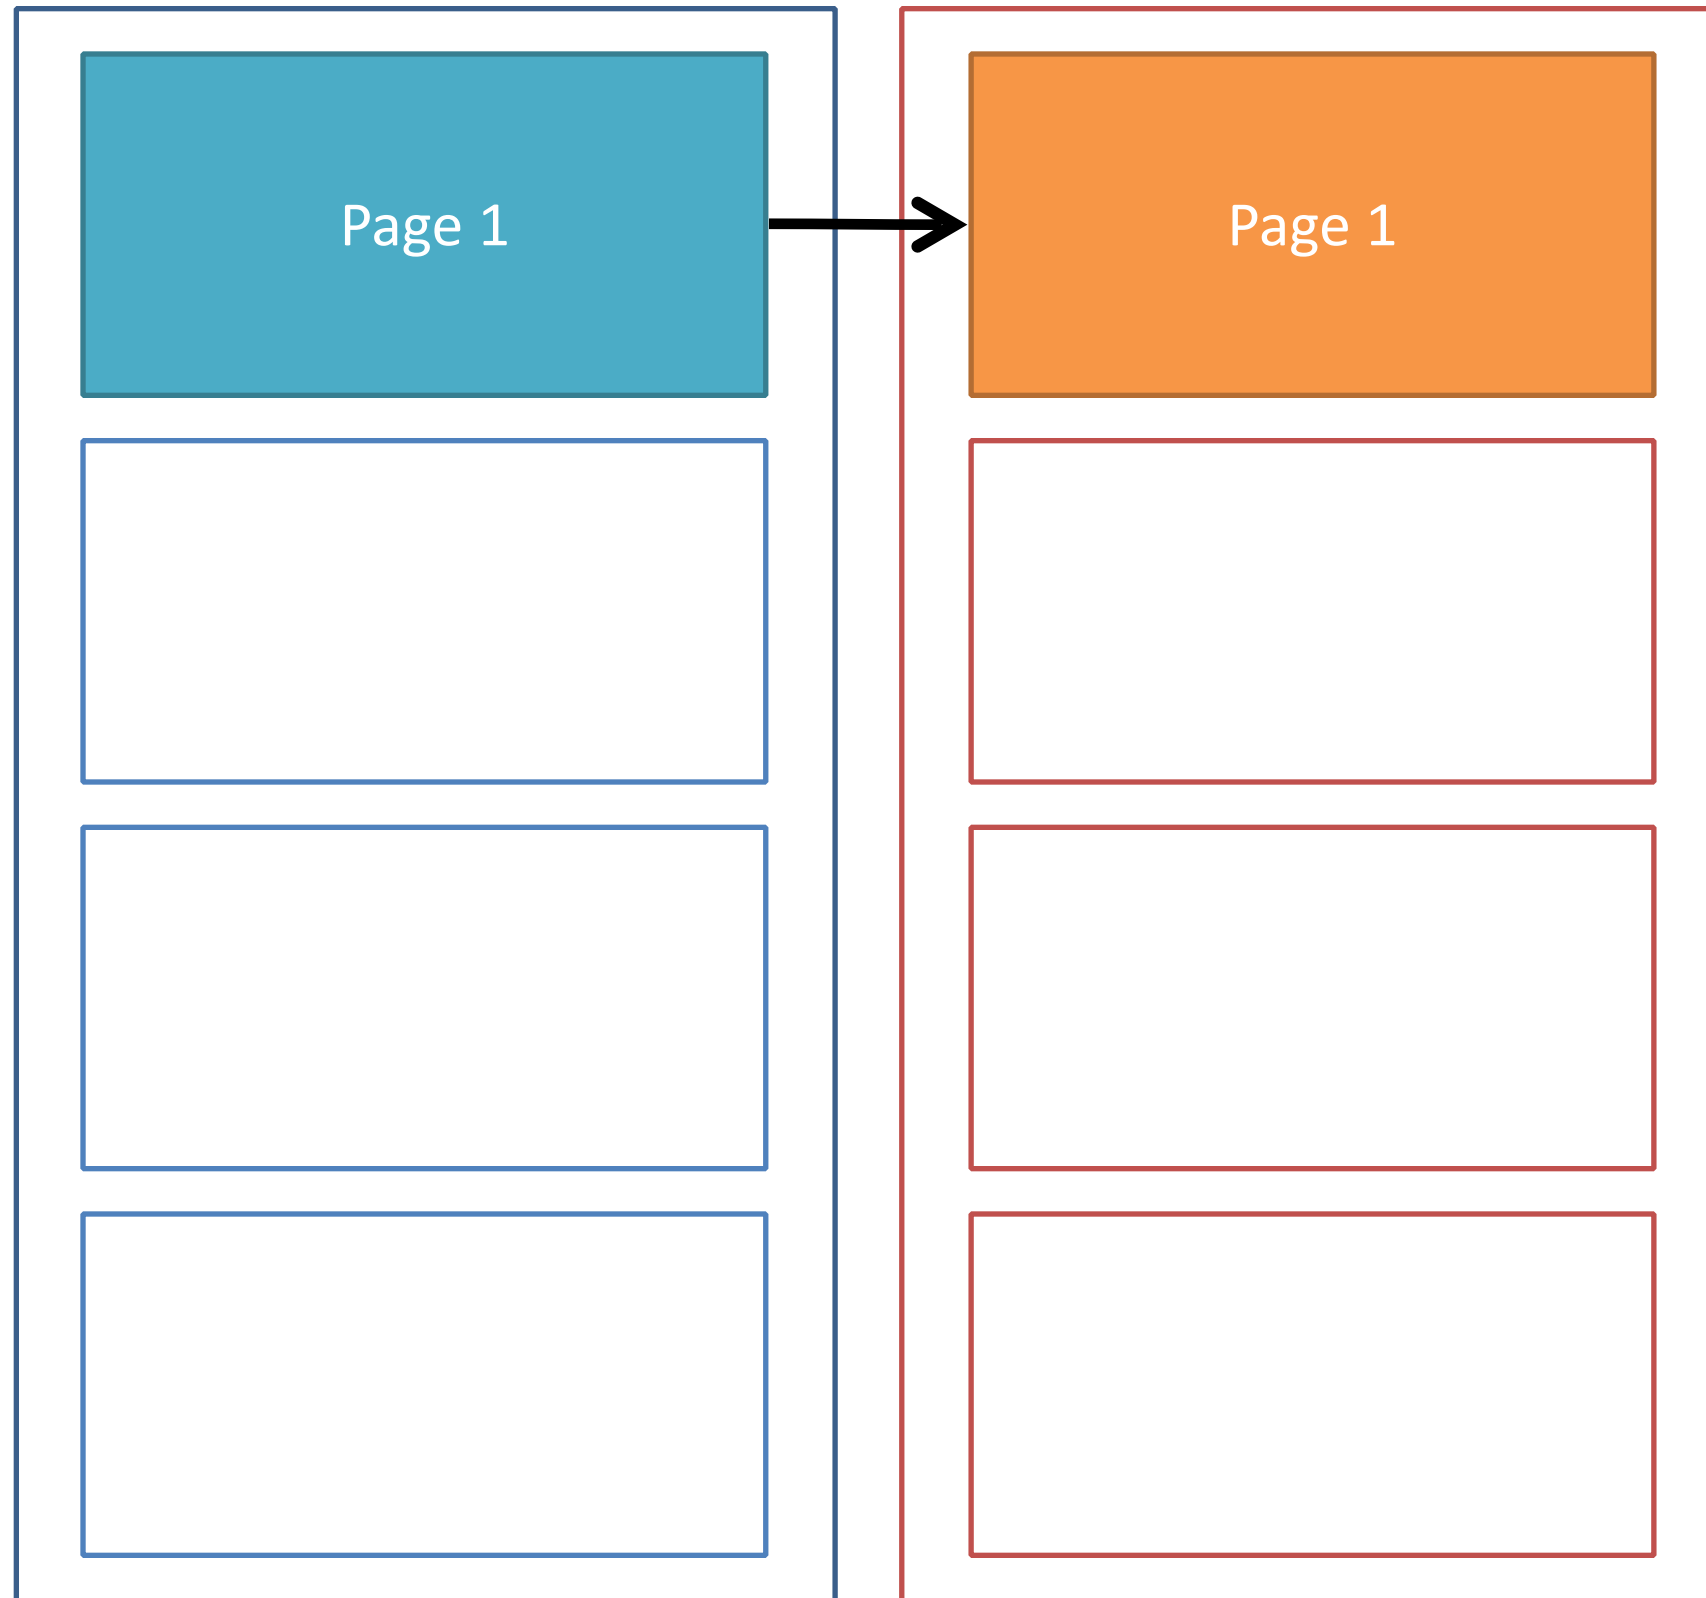
Steps:

1. Get page of S.
2. Iterate through each page in R.
3. Compare tuples in each.

Page-Oriented Nested Loops Join

Sailors

Reserves



Key idea:

Take each page of S and match with each page of R.

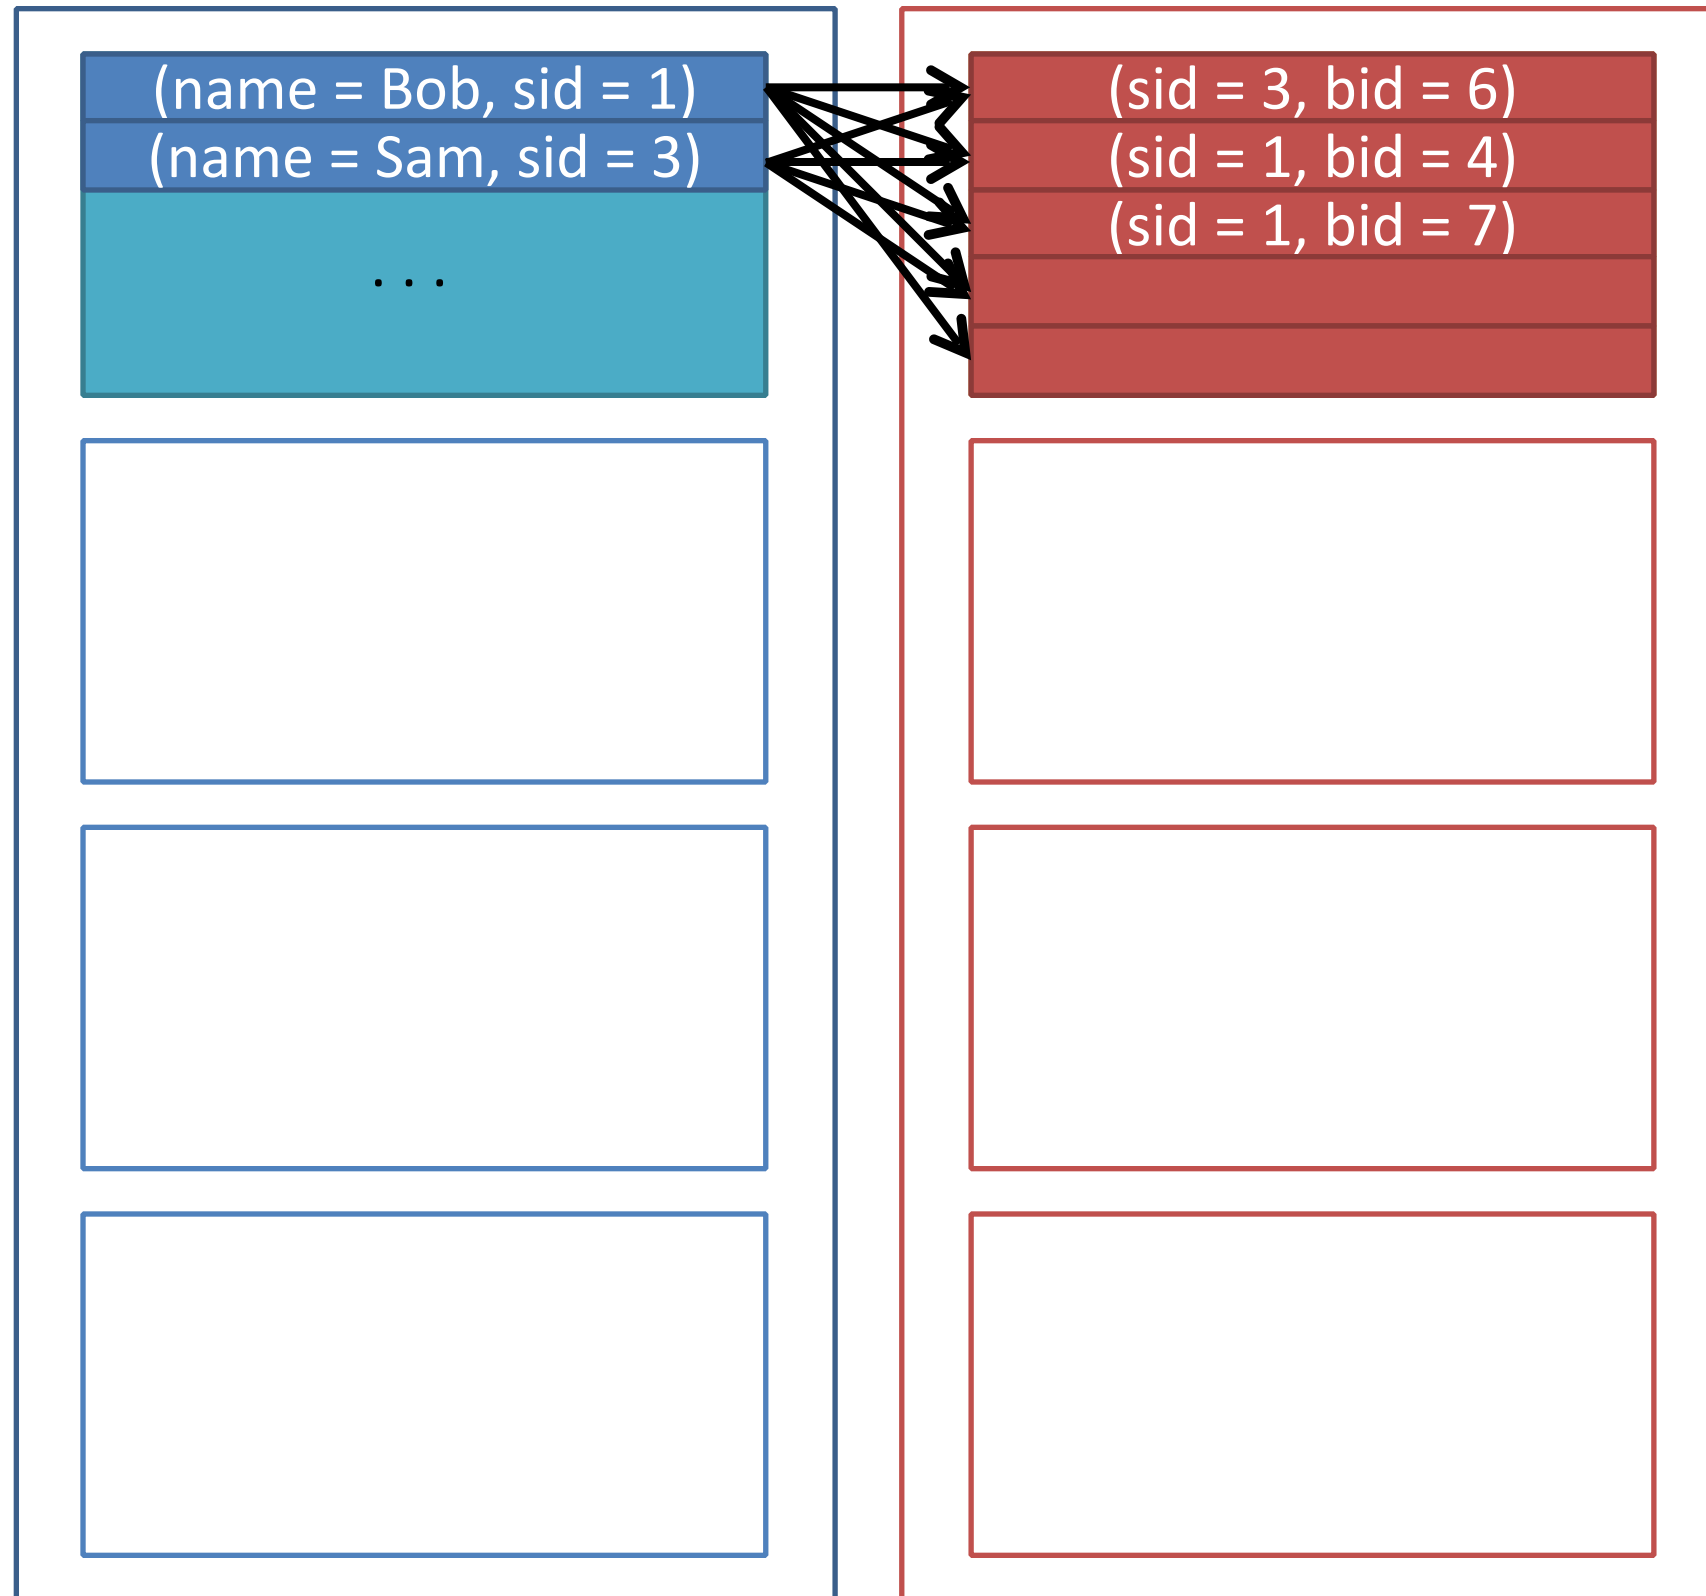
Steps:

1. Get page of S.
2. Iterate through each page in R.
3. Compare tuples in each.

Page-Oriented Nested Loops Join

Sailors

Reserves



Key idea:

Take each page of S and match with each page of R.

Steps:

1. Get page of S.
2. Iterate through each page in R.
3. Compare tuples in each.

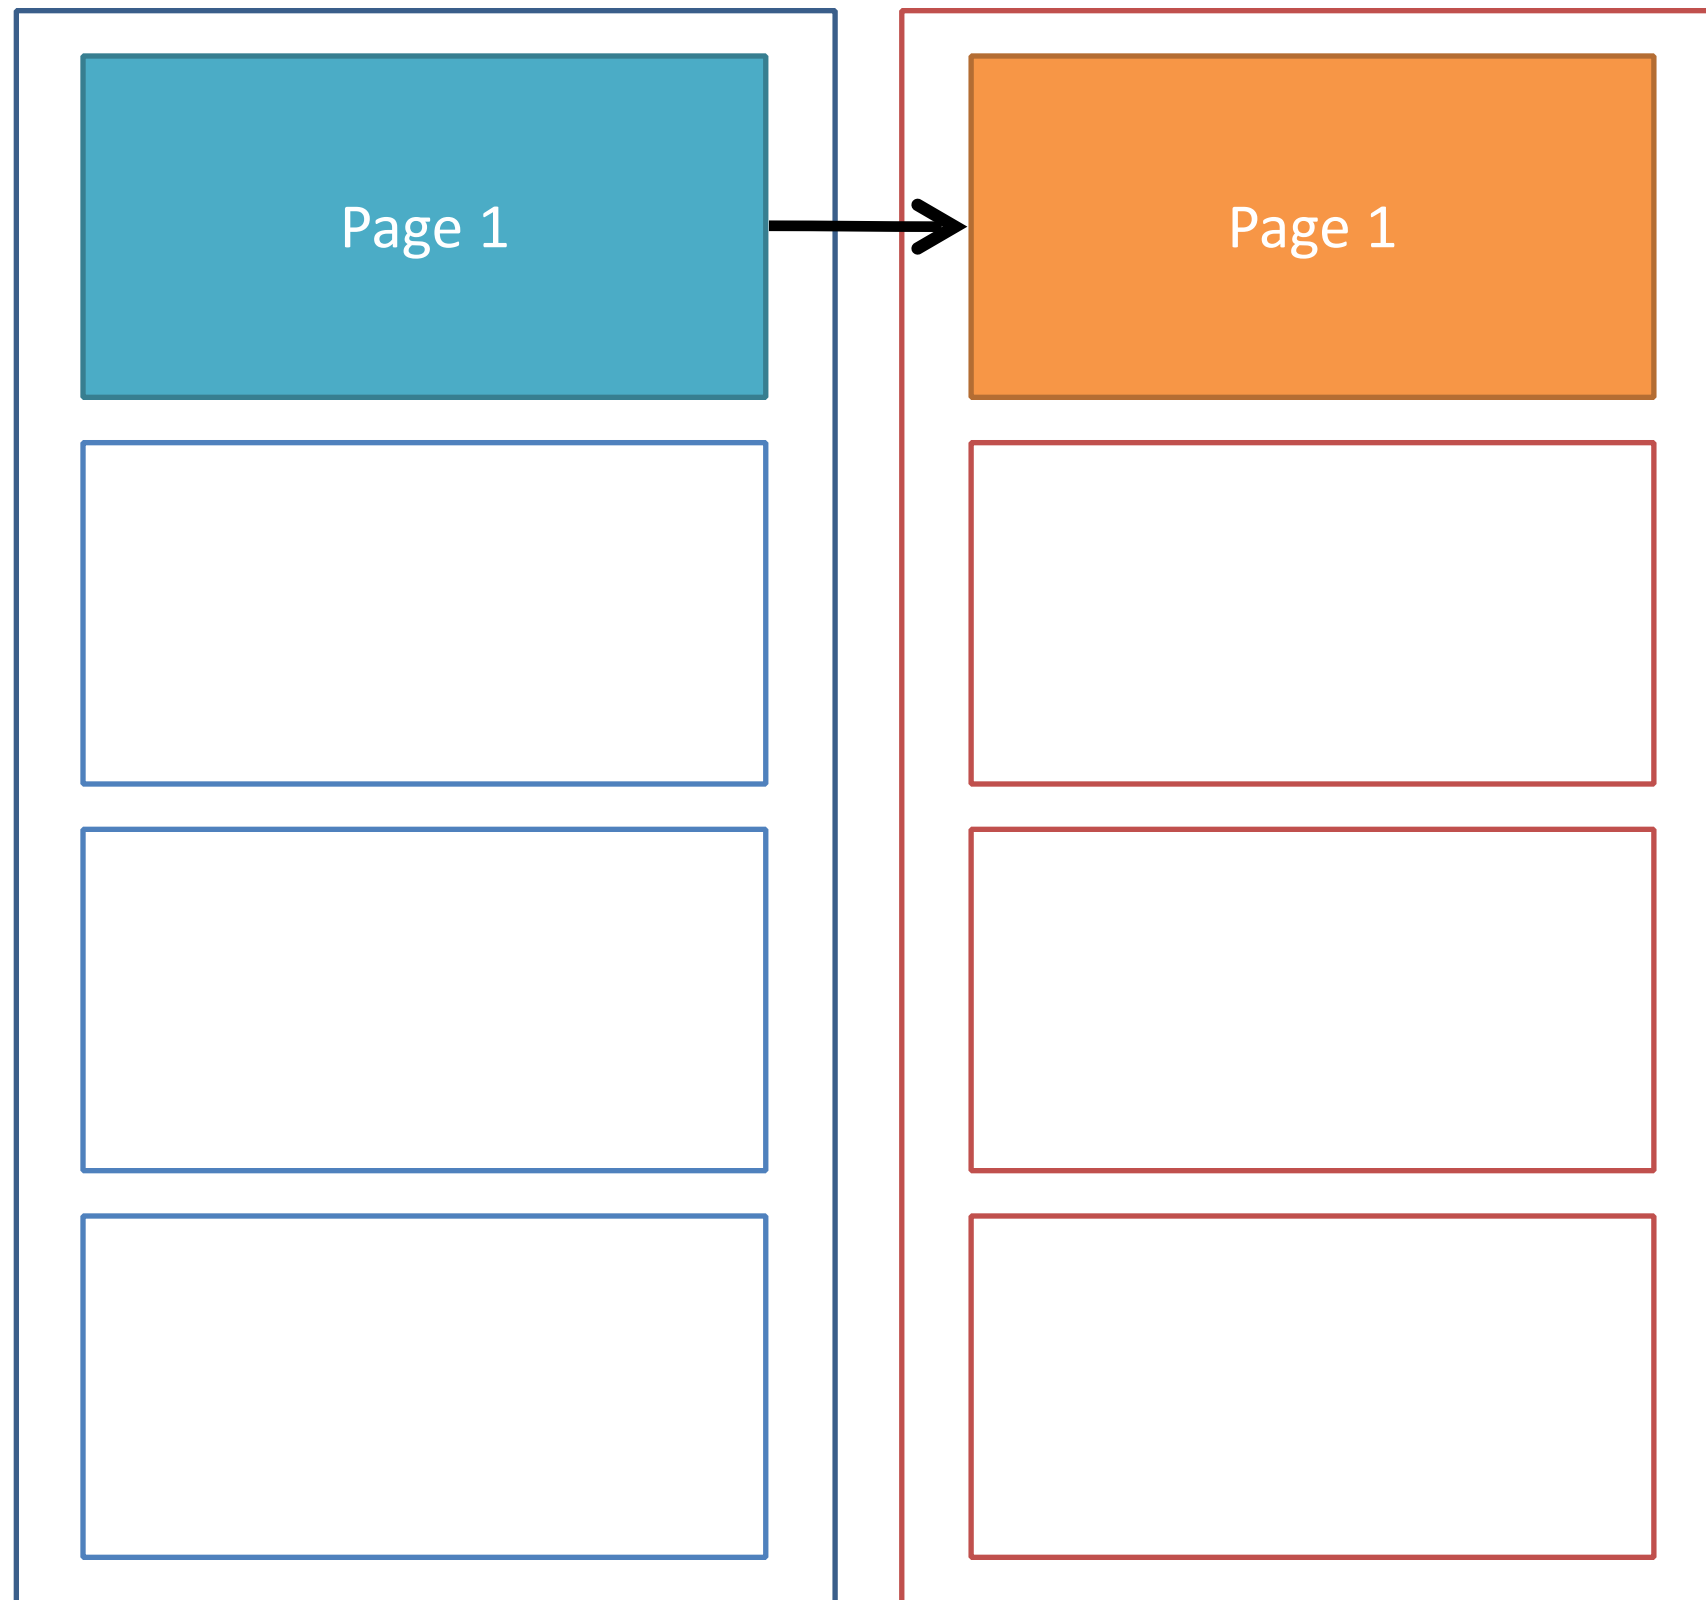
Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)
(name = Sam, sid = 3, bid = 6)

Page-Oriented Nested Loops Join

Sailors

Reserves



Key idea:

Take each page of S and match with each page of R.

Steps:

1. Get page of S.
2. Iterate through each page in R.
3. Compare tuples in each.


Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)
(name = Sam, sid = 3, bid = 6)

Page-Oriented Nested Loops Join

Sailors

Reserves



Key idea:

Take each page of *S* and match with each page of *R*.

Steps:

1. Get page of *S*.
2. Iterate through each page in *R*.
3. Compare tuples in each.


Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)
(name = Sam, sid = 3, bid = 6)

Page-Oriented Nested Loops Join

Sailors

Reserves



Key idea:

Take each page of S and match with each page of R.

Steps:

1. Get page of S.
2. Iterate through each page in R.
3. Compare tuples in each.

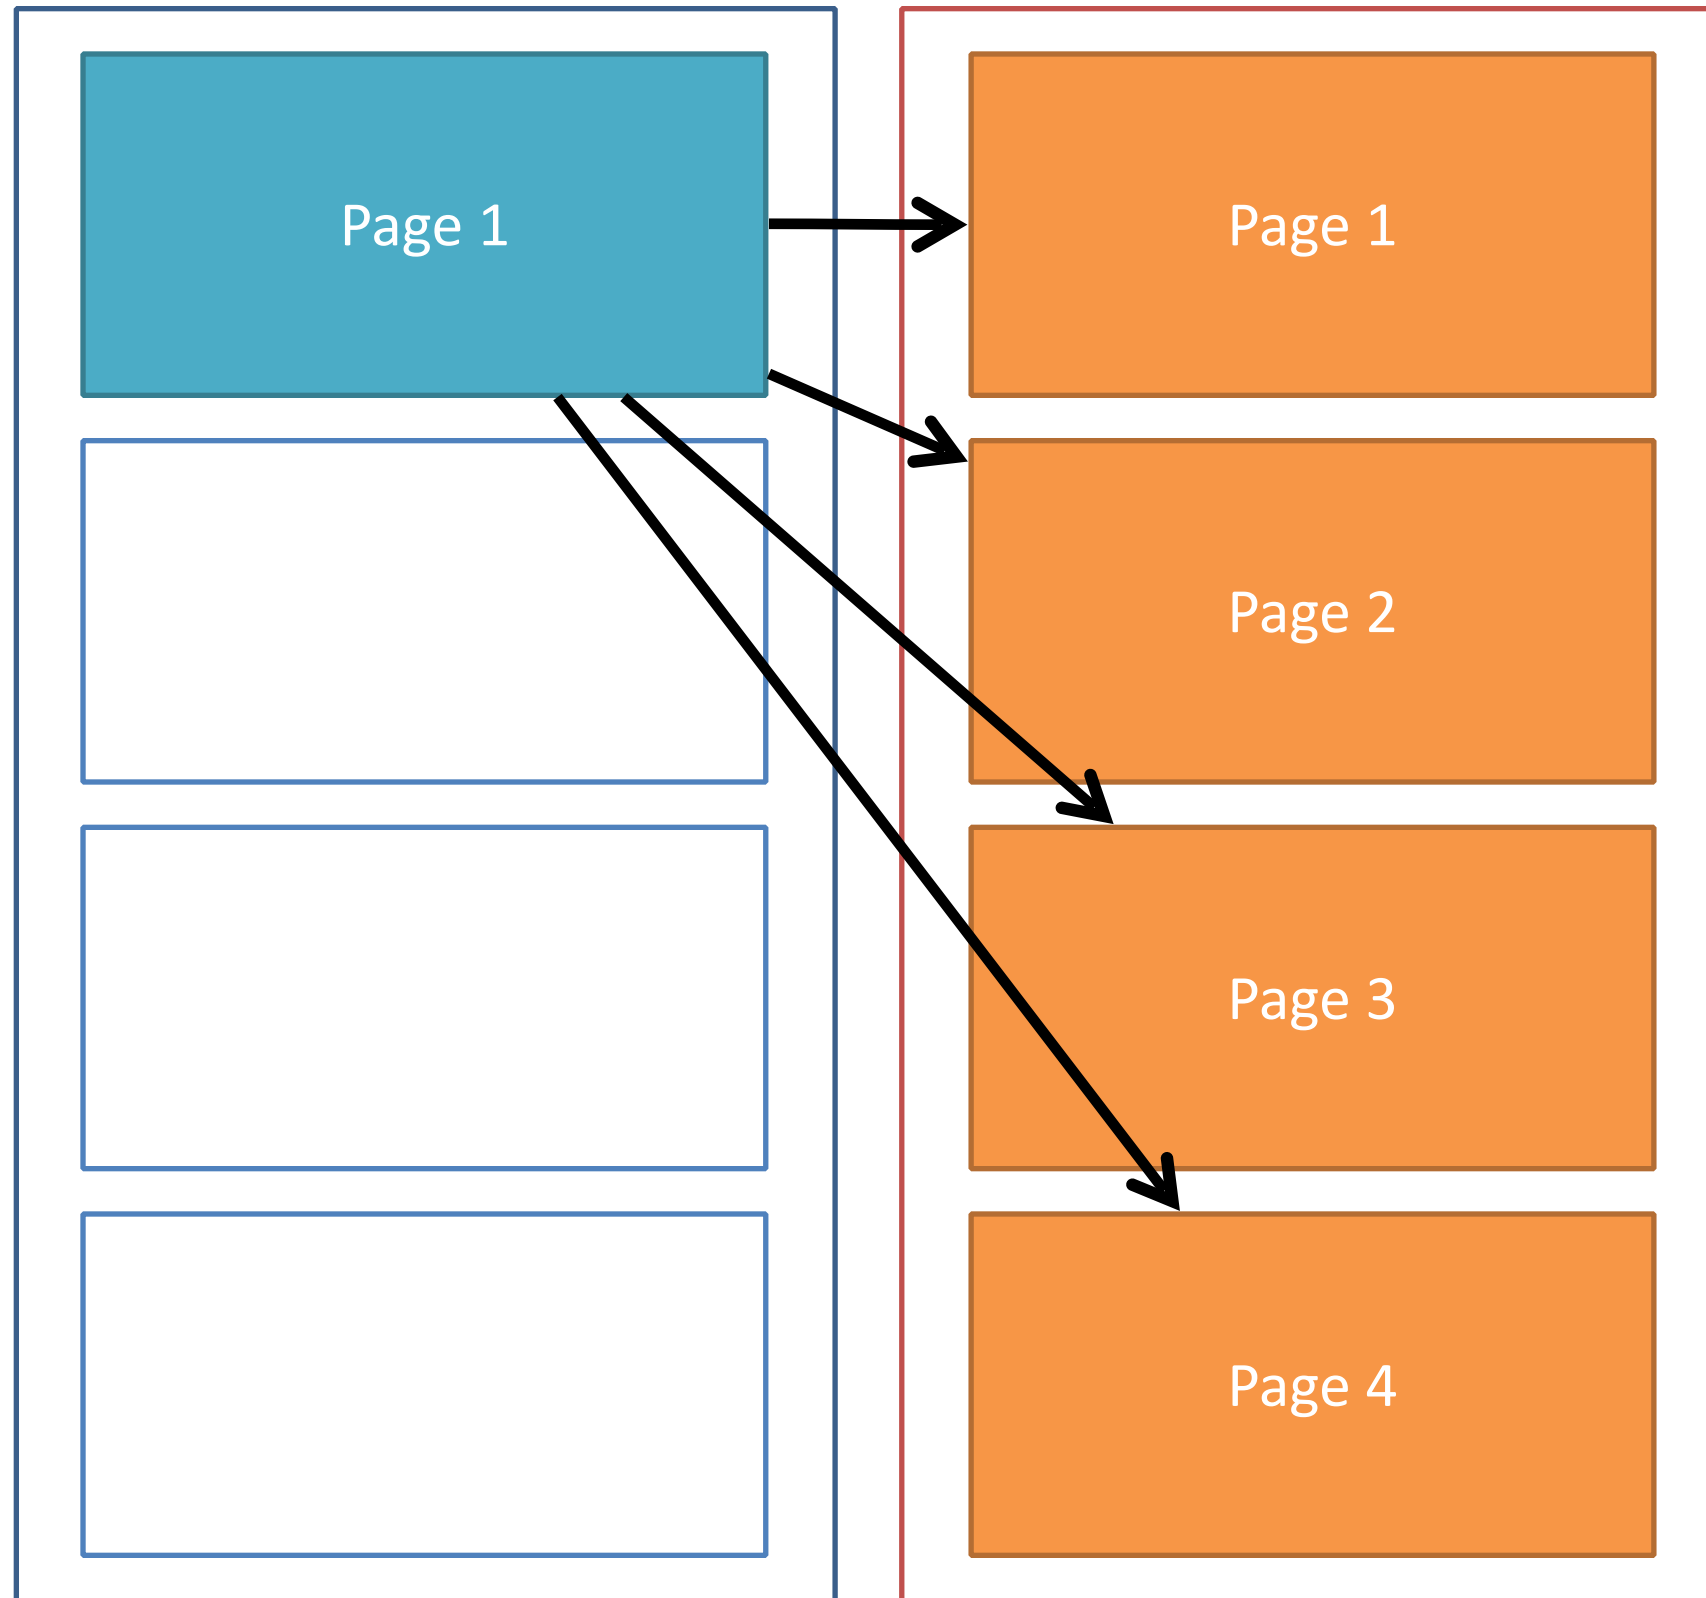
Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)
(name = Sam, sid = 3, bid = 6)

Page-Oriented Nested Loops Join

Sailors

Reserves



Key idea:

Take each page of S and match with each page of R.

Steps:

1. Get page of S.
2. Iterate through each page in R.
3. Compare tuples in each.

Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)
(name = Sam, sid = 3, bid = 6)

Page-Oriented Nested Loops Join

Sailors

Reserves

Page 1

Page 2

Key idea:

Take each page of S and match with each page of R.

Steps:

1. Get page of S.
2. Iterate through each page in R.
3. Compare tuples in each.

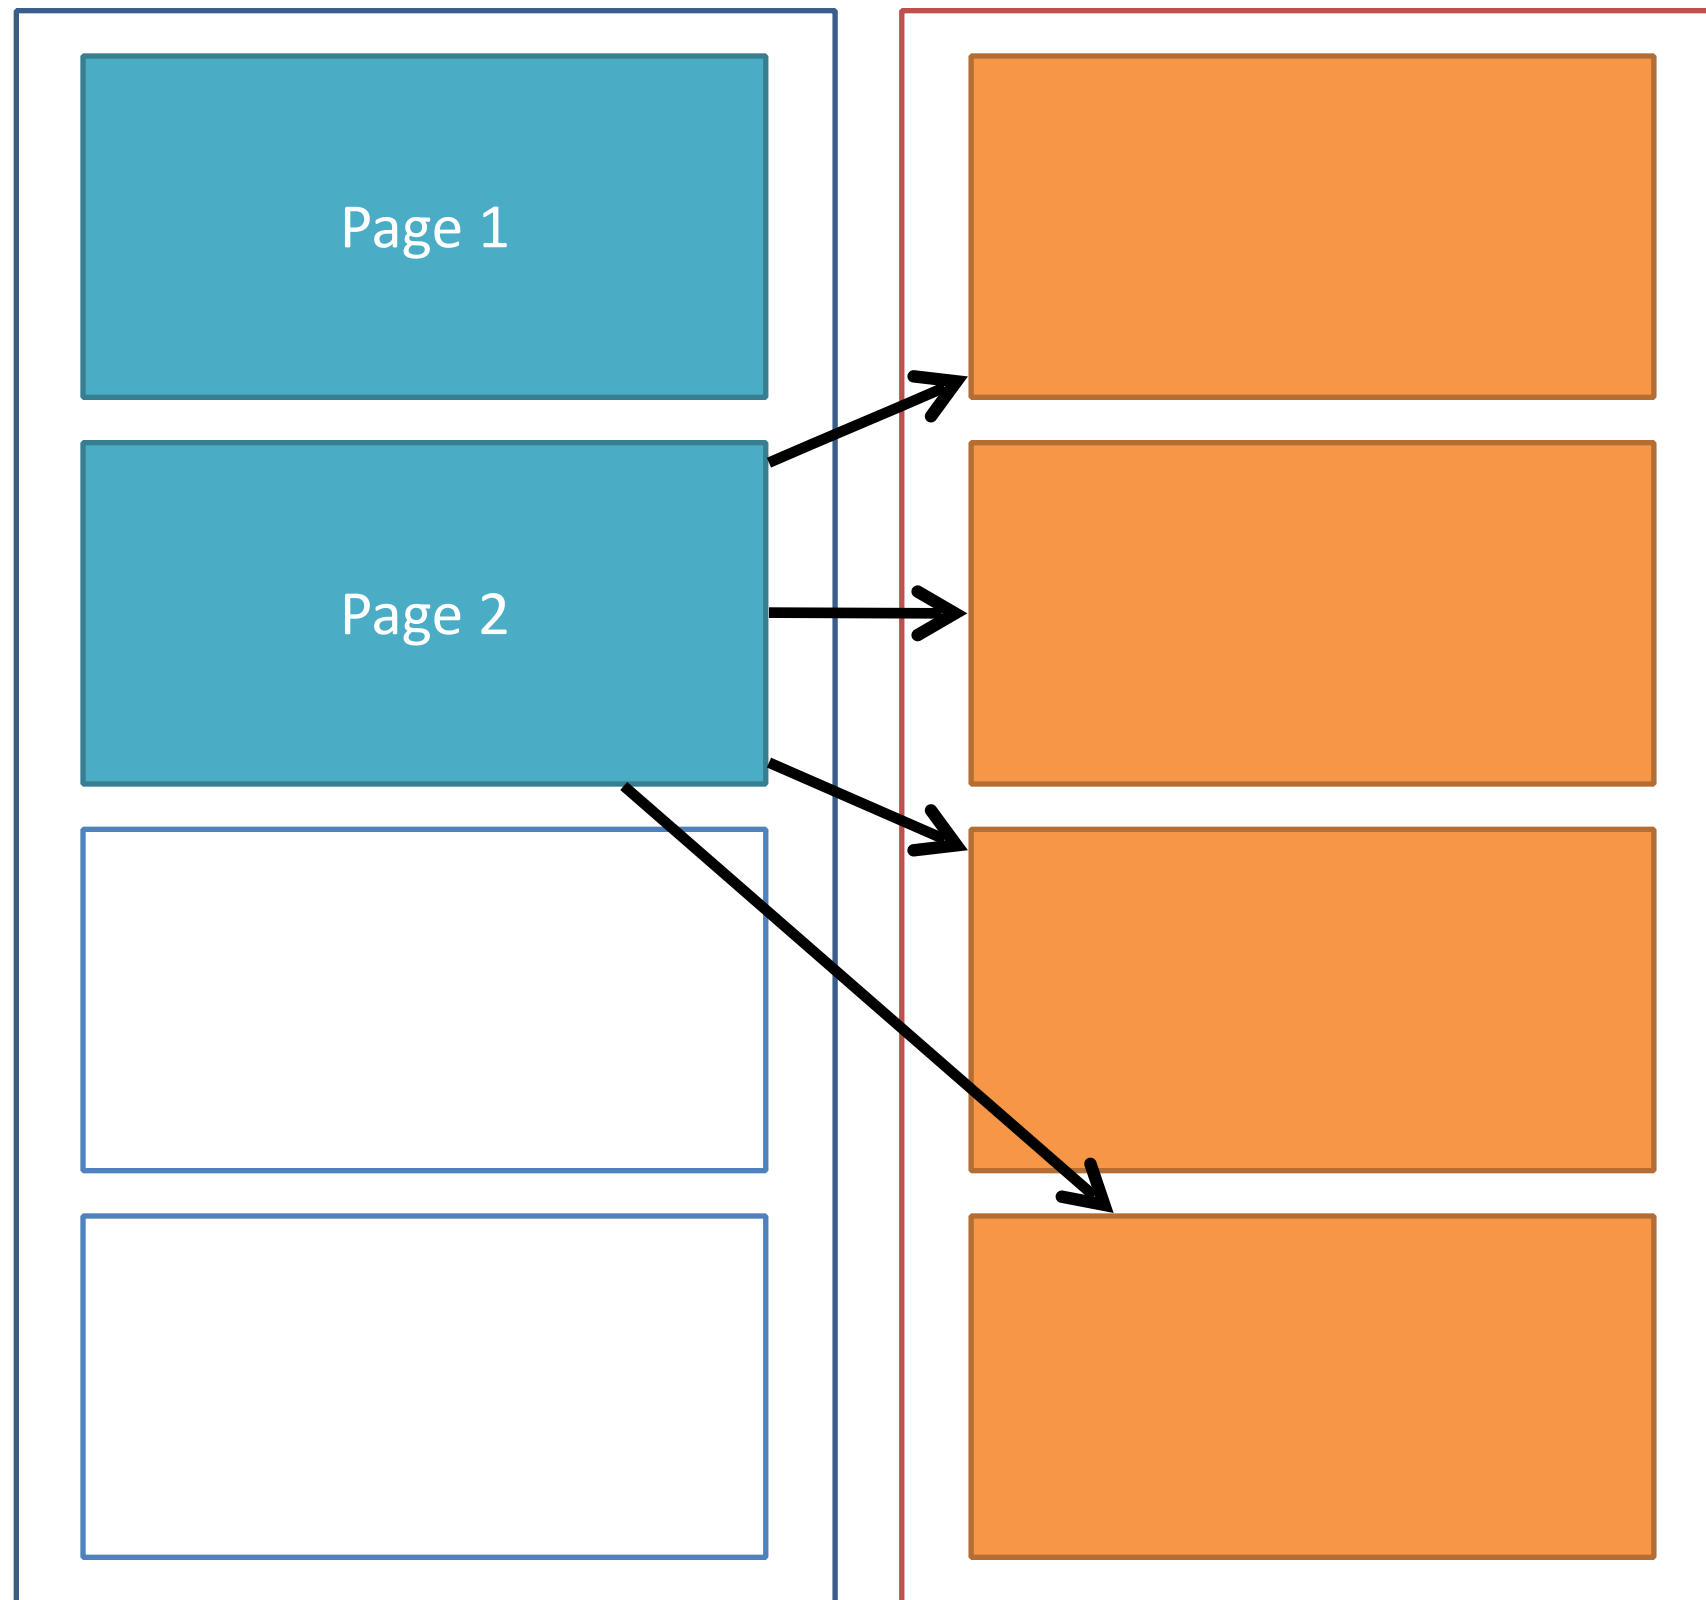
Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)
(name = Sam, sid = 3, bid = 6)

Page-Oriented Nested Loops Join

Sailors

Reserves



Key idea:

Take each page of S and match with each page of R.

Steps:

1. Get page of S.
2. Iterate through each page in R.
3. Compare tuples in each.

Output:

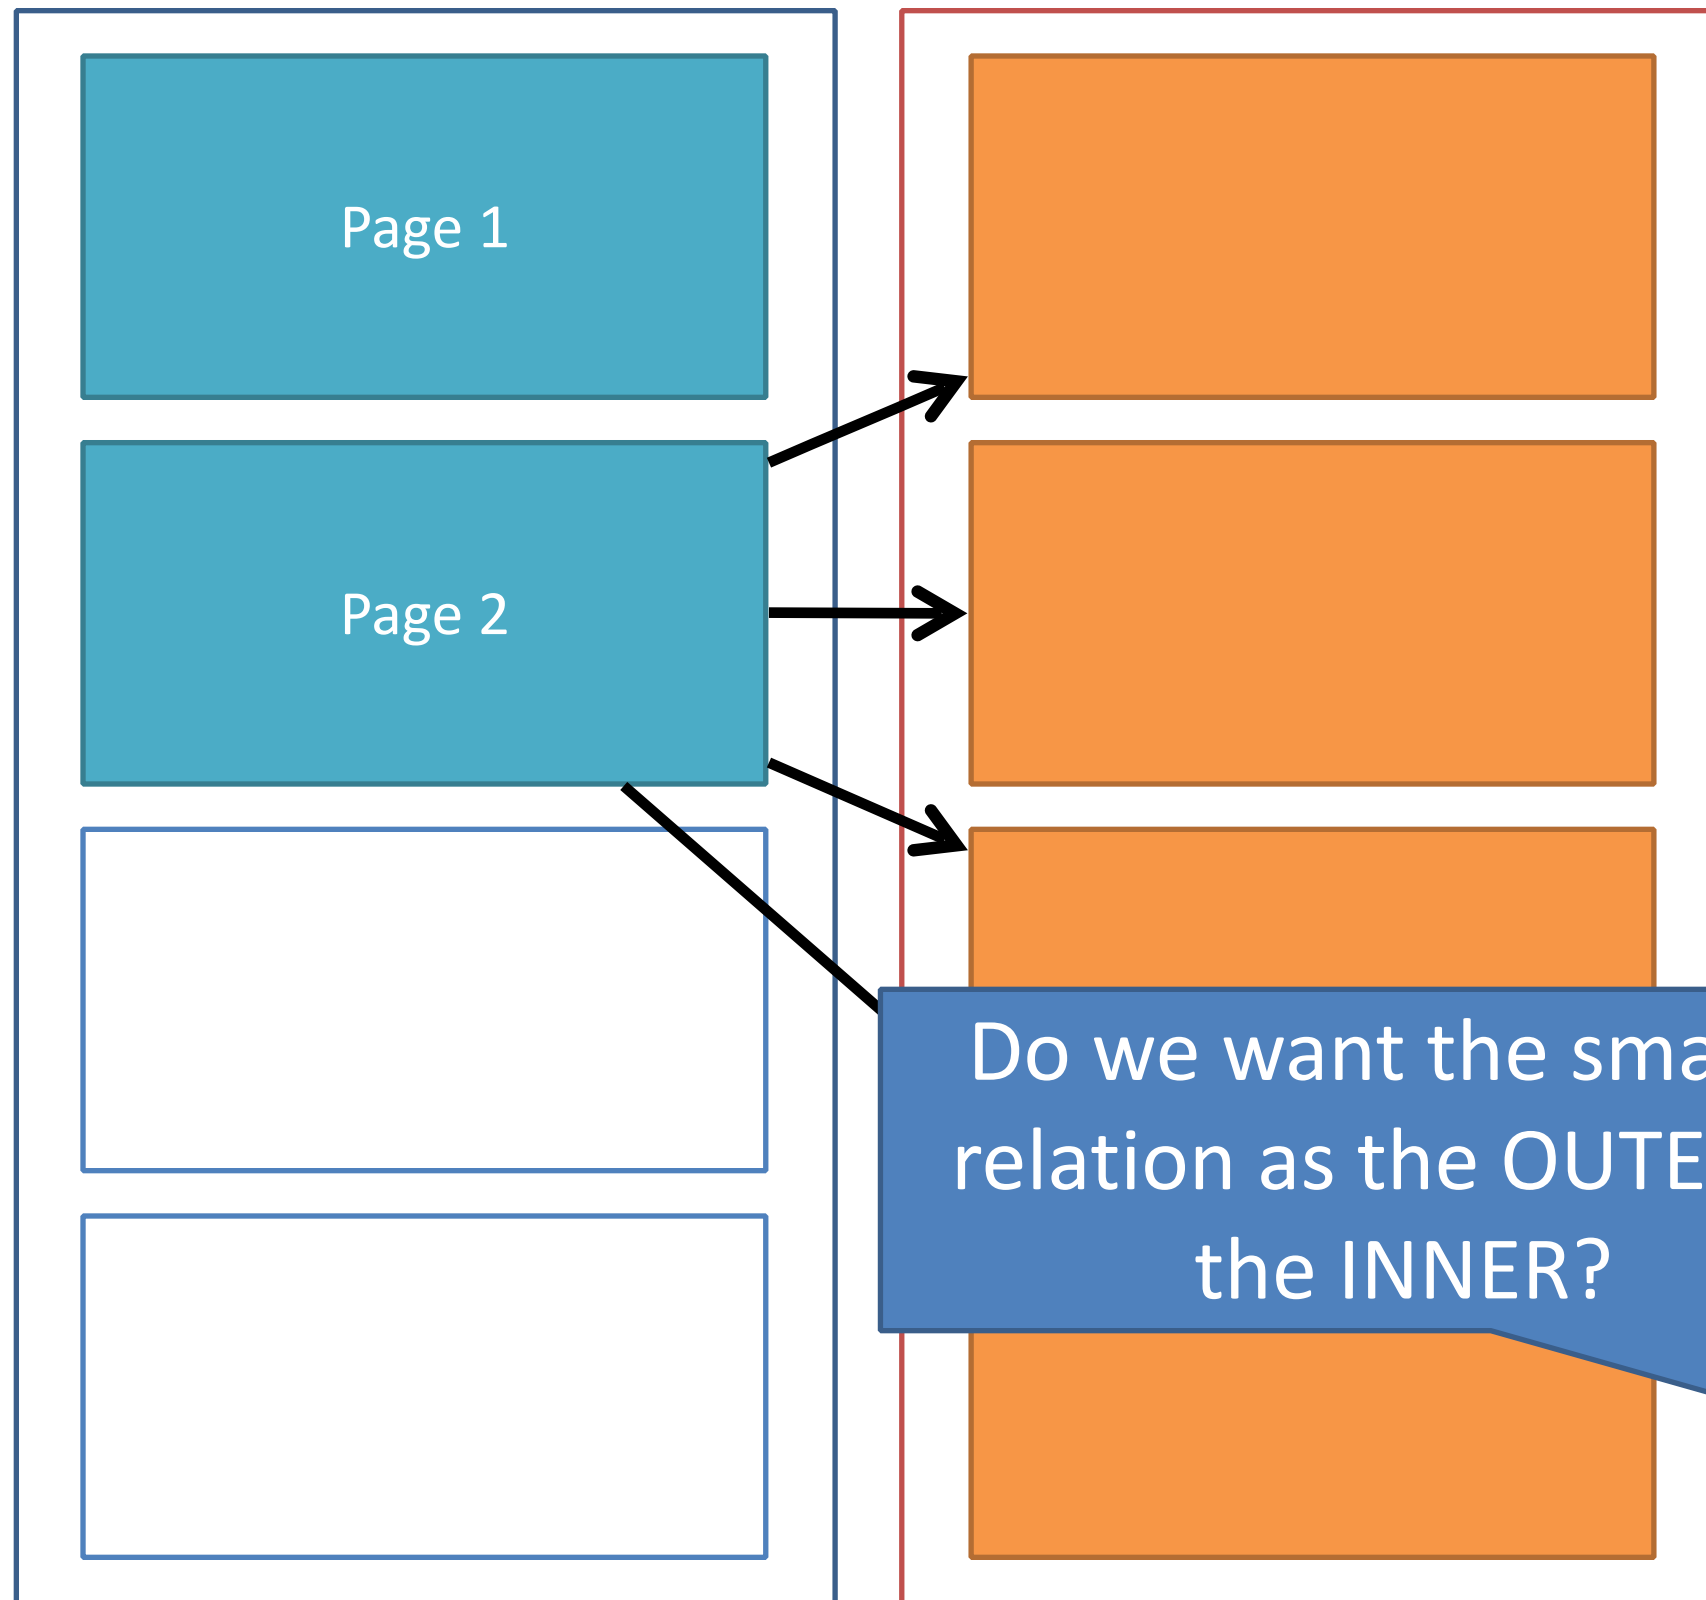
(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)
(name = Sam, sid = 3, bid = 6)

Notation: $[S]$ == “# pages in S” ; $|S|$ == “# tuples in S”

Page-Oriented Nested Loops Join

Sailors

Reserves



Key idea:

Take each page of S and match with each page of R.

Steps:

1. Get page of S.
2. Iterate through each page in R.



Compare tuples in each.

$$[S] + [S] * [R]$$

Chunk Nested Loops Join

Sailors

Reserves



Key idea:

Take **k pages** of S
and match with
each page of R.



Steps:

1. Get **k** pages of S.
2. Iterate through
each page in R.
3. Compare tuples in
each.

Chunk Nested Loops Join

Sailors

Reserves



Key idea:

Take **k pages** of S
and match with
each page of R.


Steps:

1. Get **k** pages of S.
2. Iterate through
each page in R.
3. Compare tuples in
each.

Chunk Nested Loops Join

Sailors

Reserves



Key idea:

Take **k pages** of S and match with each page of R.

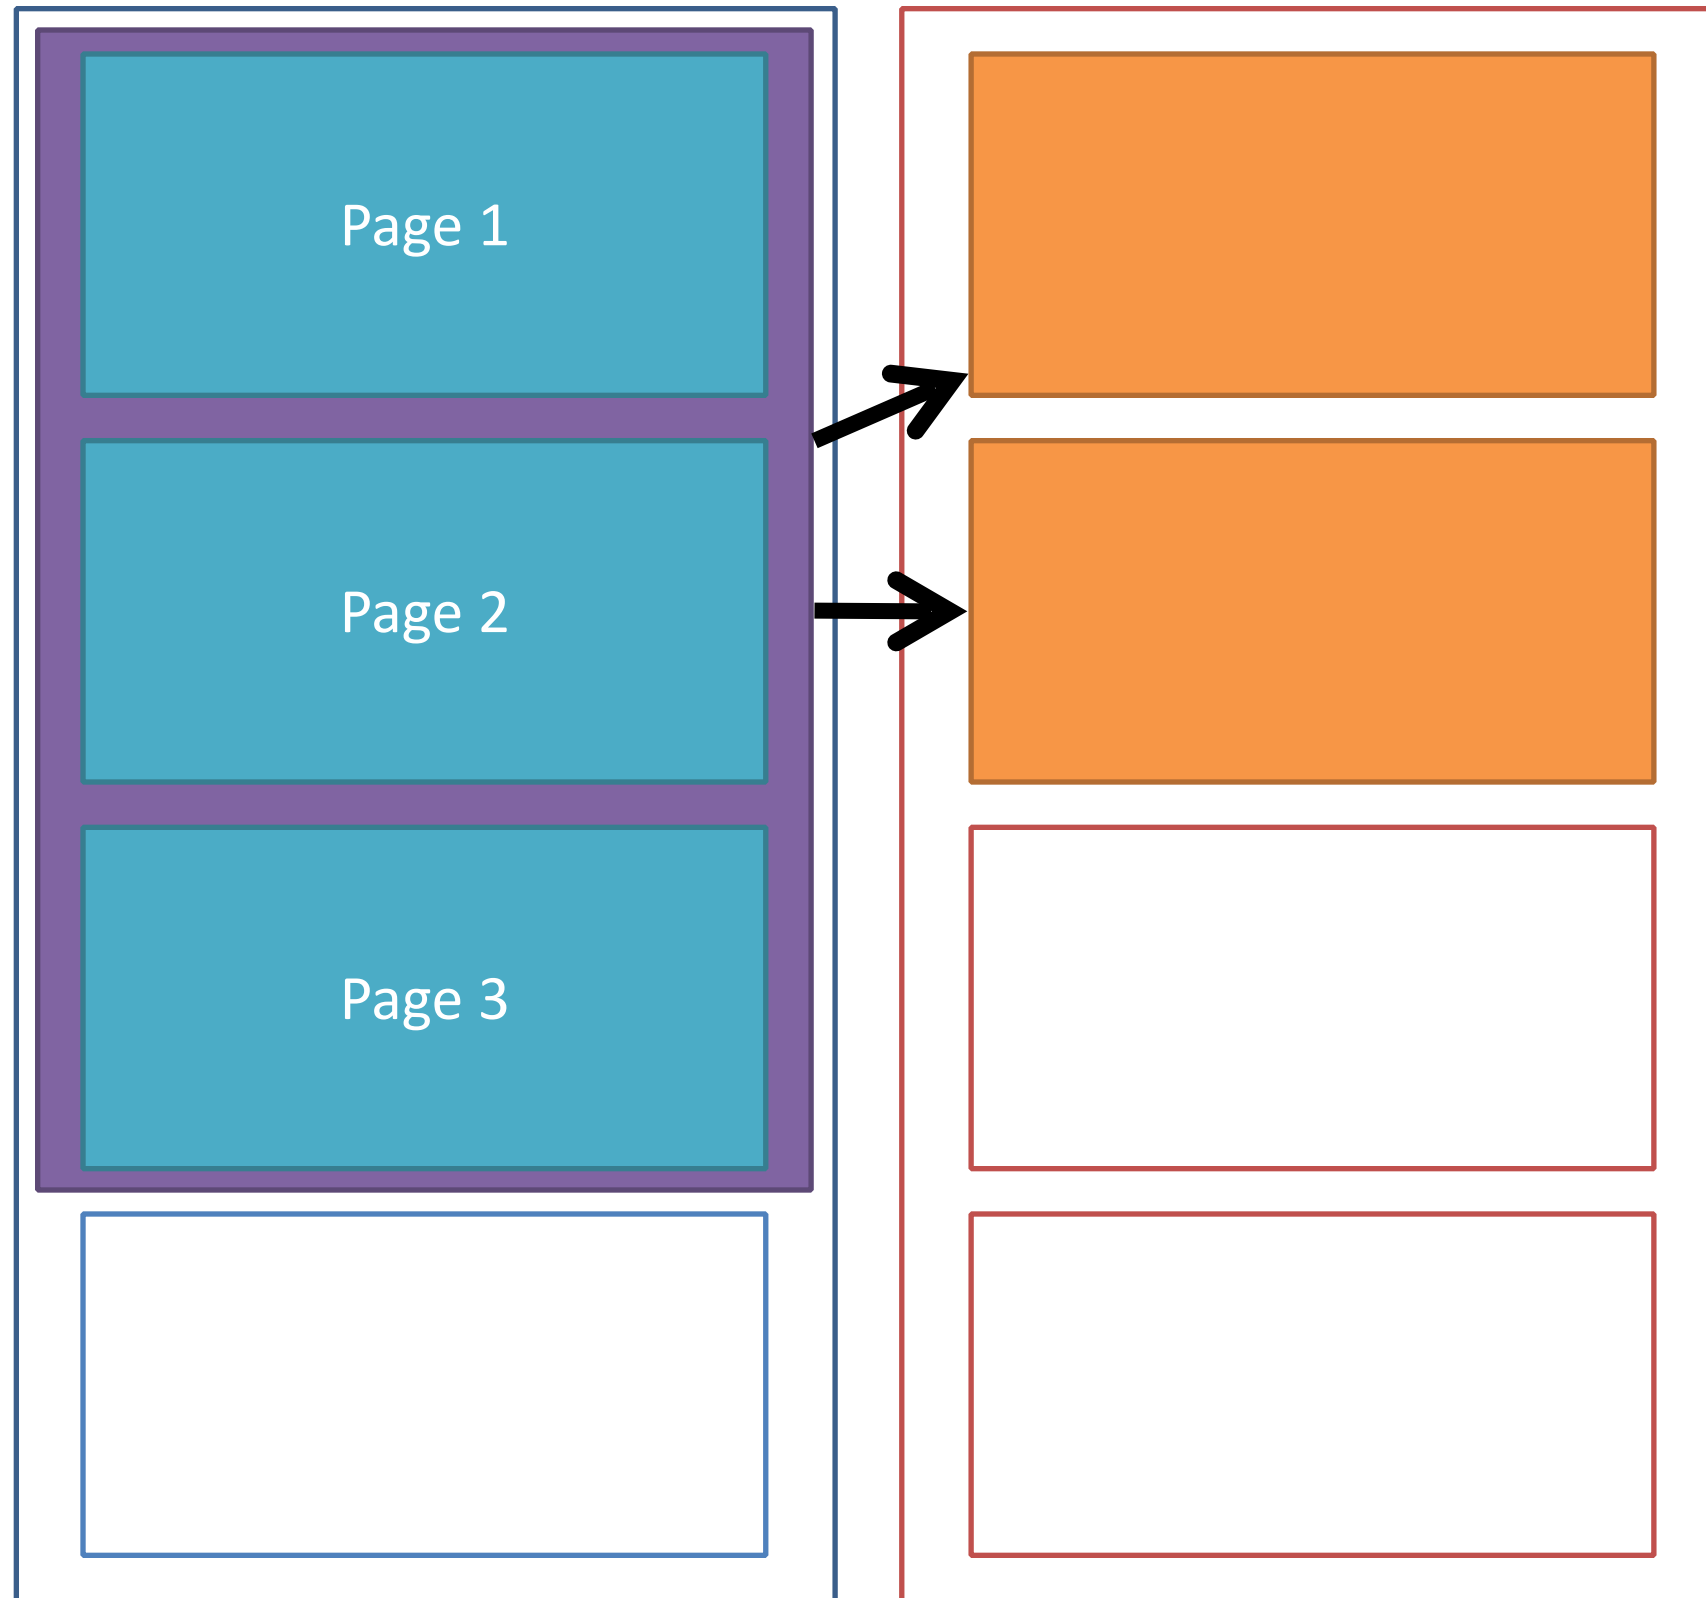
Steps:

1. Get **k** pages of S.
2. Iterate through each page in R.
3. Compare tuples in each.

Chunk Nested Loops Join

Sailors

Reserves



Key idea:

Take **k pages** of S and match with each page of R.

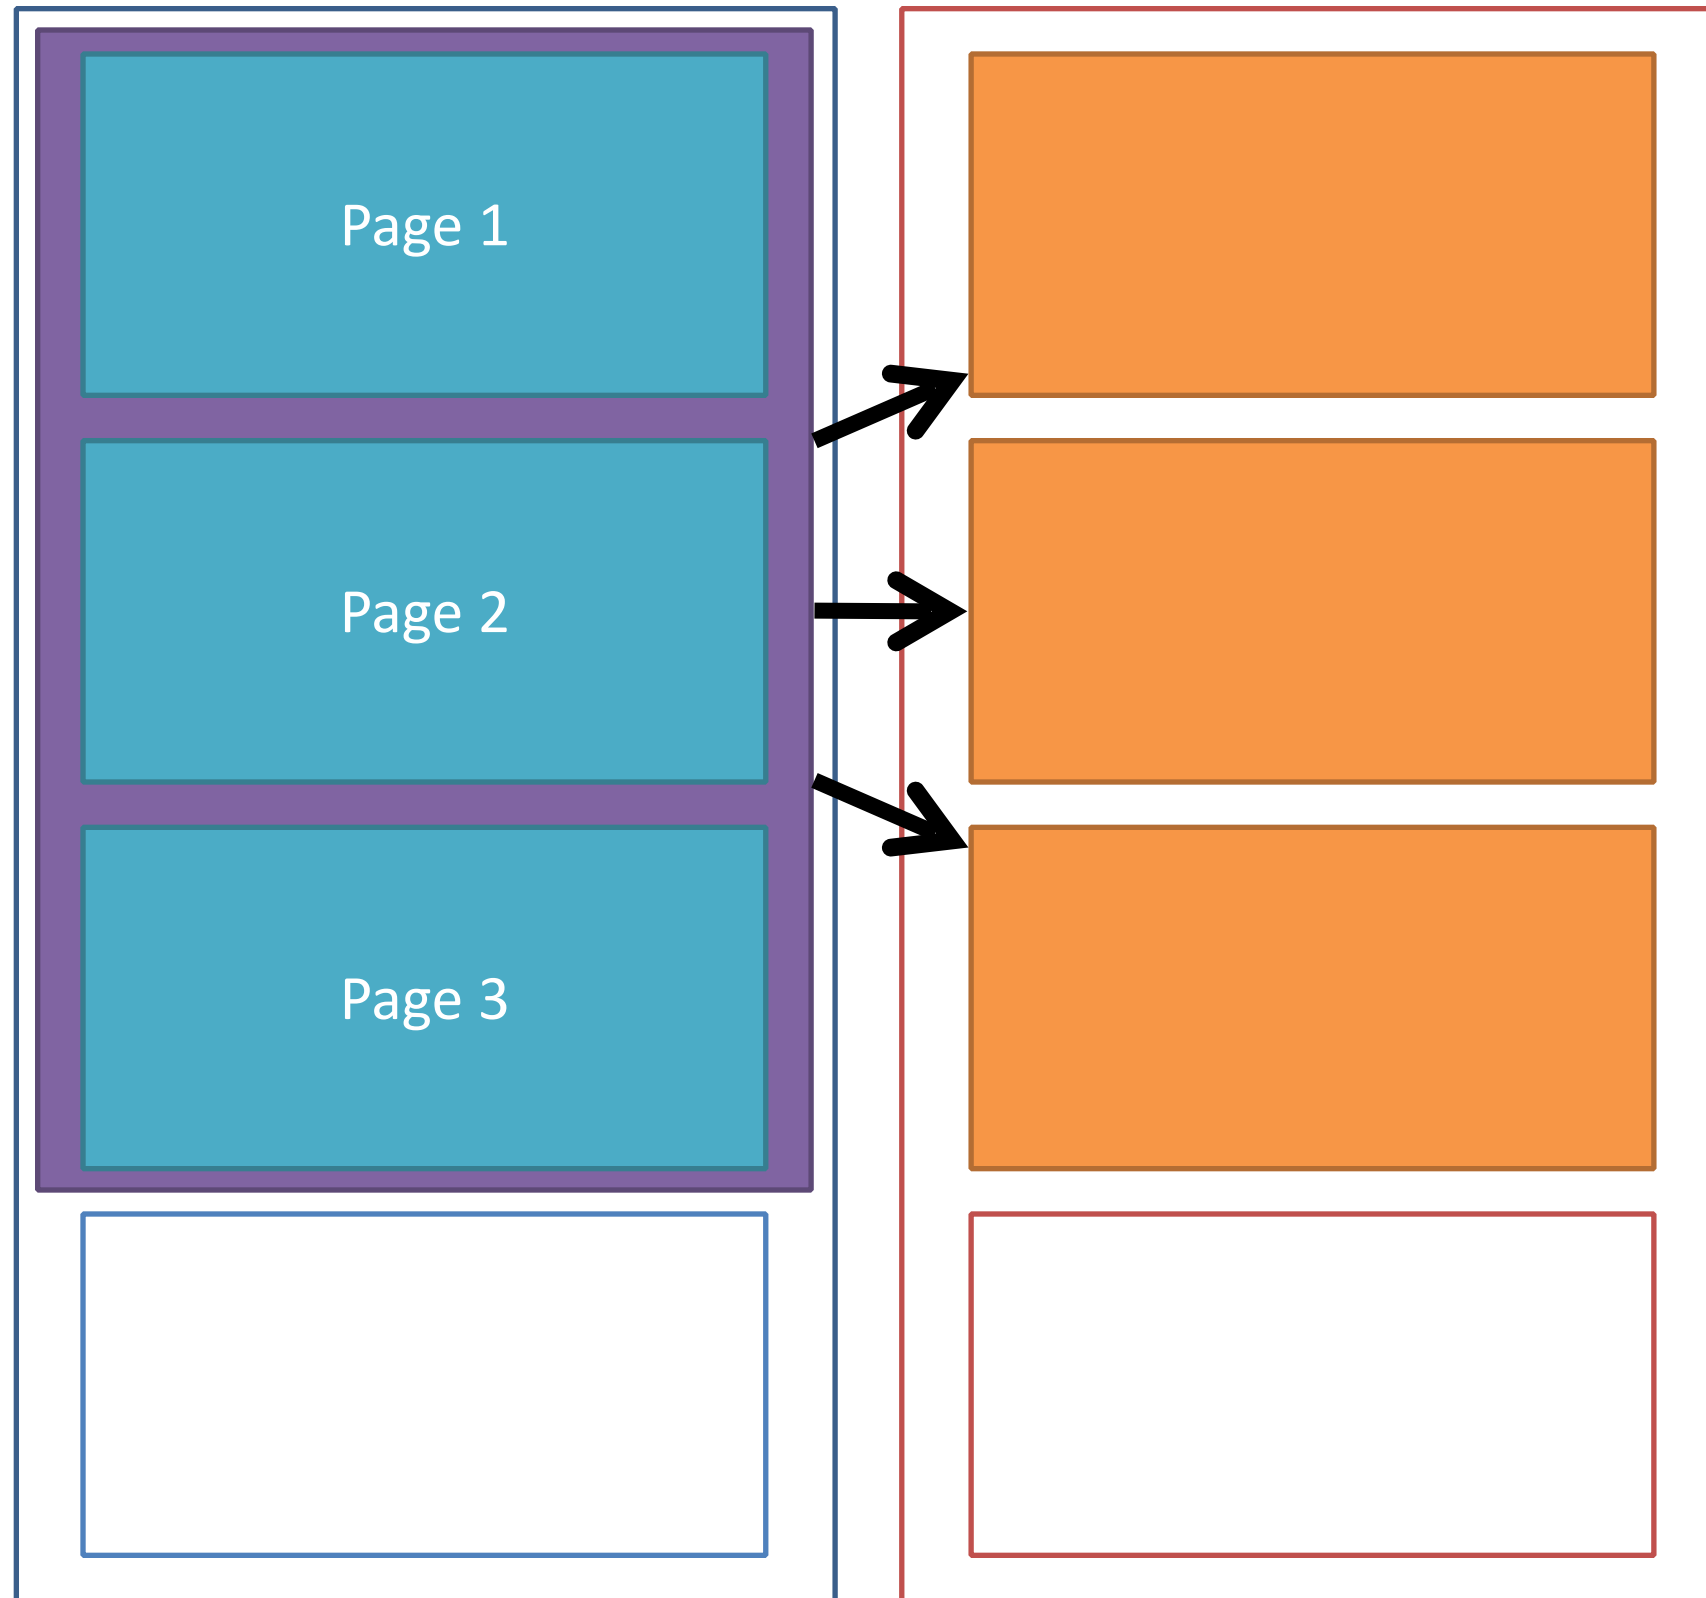
Steps:

1. Get **k** pages of S.
2. Iterate through each page in R.
3. Compare tuples in each.

Chunk Nested Loops Join

Sailors

Reserves



Key idea:

Take **k pages** of S and match with each page of R.

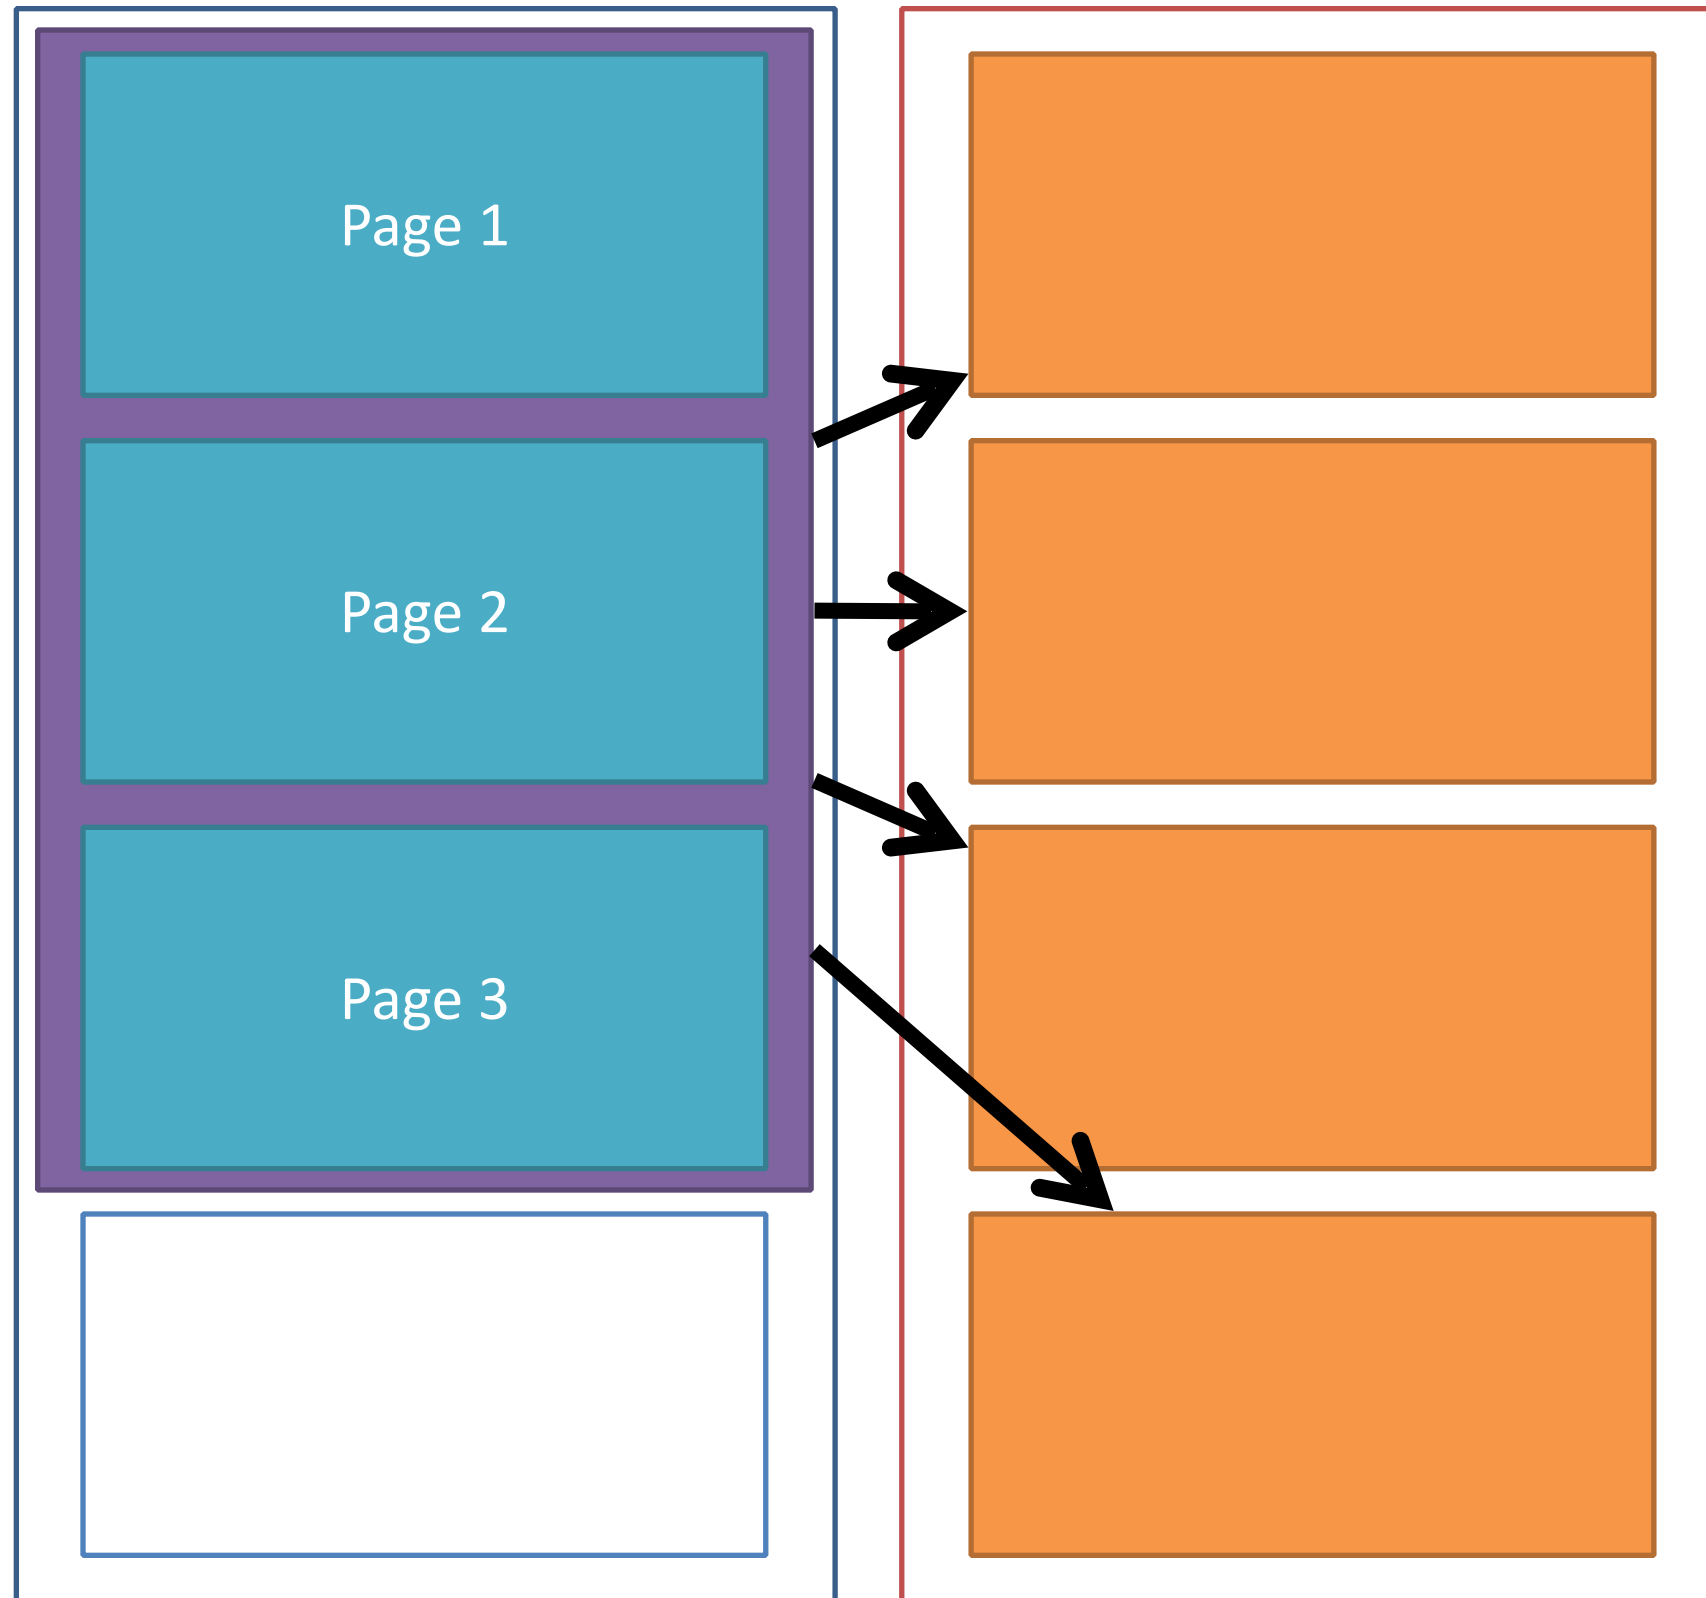
Steps:

1. Get **k** pages of S.
2. Iterate through each page in R.
3. Compare tuples in each.

Chunk Nested Loops Join

Sailors

Reserves



Key idea:

Take **k pages** of S and match with each page of R.

Steps:

1. Get **k** pages of S.
2. Iterate through each page in R.
3. Compare tuples in each.

Chunk Nested Loops Join

Sailors

Reserves

Page 1

Page 2

Page 3

Page 4

Key idea:

Take **k pages** of S
and match with
each page of R.

Steps:


1. Get **k** pages of S.
2. Iterate through
each page in R.
3. Compare tuples in
each.

Notation: $[S]$ == “# pages in S” ; $|S|$ == “# tuples in S”

Chunk Nested Loops Join

Sailors

Reserves



Key idea:

Take **k pages** of S and match with each page of R.

Steps:

1. Get k pages of S.
2. Iterate through each page in R.

Compare tuples in each.

:

$$[S] + ([S] / k) * [R]$$

Sort-Merge Join

Sailors

Reserves

Key idea:

Sort S and R, then
merge them!

Steps:

1. Sort S and R.
2. “Zip” or merge.

Sort-Merge Join

Sailors

Reserves

Key idea:

Sort S and R **on join column**, then merge them!

Steps:

1. Sort S and R.
2. “Zip” or merge.

(name = Bob, sid = 1)

(name = Sam, sid = 3)

(name = Sue, sid = 7)

(name = Jill, sid = 2)

(name = Joe, sid = 12)

(name = Sue, sid = 8)

(name = Yue, sid = 4)

Sort-Merge Join

Sailors

Reserves

Key idea:

Sort S and R **on join column**, then merge them!

Steps:

1. Sort S and R.
2. “Zip” or merge.

(name = Bob, sid = 1)

(name = Jill, sid = 2)

(name = Sam, sid = 3)

(name = Yue, sid = 4)

(name = Sue, sid = 7)

(name = Sue, sid = 8)

(name = Joe, sid = 12)

...

Sort-Merge Join

Sailors

Reserves

Key idea:

Sort S and R **on join column**, then merge them!

Steps:

1. Sort S and R.
2. “Zip” or merge.

(name = Bob, sid = 1)

(name = Jill, sid = 2)

(name = Sam, sid = 3)

(name = Yue, sid = 4)

(name = Sue, sid = 7)

(name = Sue, sid = 8)

(name = Joe, sid = 12)

...

(sid = 1, bid = 4)

(sid = 1, bid = 7)

(sid = 3, bid = 6)

(sid = 4, bid = 3)

(sid = 8, bid = 1)

(sid = 8, bid = 13)

(sid = 8, bid = 15)


(sid = 12, bid = 1)

...

Sort-Merge Join

Sailors

Reserves



Key idea:

Sort S and R **on join column**, then merge them!

Steps:

1. Sort S and R.
2. “Zip” or merge.

Output:

(name = Bob, sid = 1, bid = 4)

Sort-Merge Join

Sailors

Reserves



Key idea:

Sort S and R **on join column**, then merge them!

Steps:

1. Sort S and R.
2. “Zip” or merge.

Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)
(name = Sam, sid = 3, bid = 6)

(name = Bob, sid = 1)
(name = Jill, sid = 2)
(name = Sam, sid = 3)
(name = Yue, sid = 4)
(name = Sue, sid = 7)

(name = Sue, sid = 8)
(name = Joe, sid = 12)
...

(sid = 1, bid = 4)
(sid = 1, bid = 7)
(sid = 3, bid = 6)
(sid = 4, bid = 3)
(sid = 8, bid = 1)

(sid = 8, bid = 13)
(sid = 8, bid = 15)
(sid = 12, bid = 1)
...

Sort-Merge Join

Sailors

Reserves

Key idea:

Sort S and R **on join column**, then merge them!

Steps:

1. Sort S and R.
2. “Zip” or merge.

Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)
(name = Sam, sid = 3, bid = 6)

...

(name = Bob, sid = 1)
(name = Jill, sid = 2)
(name = Sam, sid = 3)
(name = Yue, sid = 4)
(name = Sue, sid = 7)

(name = Sue, sid = 8)
(name = Joe, sid = 12)
...

(sid = 1, bid = 4)
(sid = 1, bid = 7)
(sid = 3, bid = 6)
(sid = 4, bid = 3)
(sid = 8, bid = 1)

(sid = 8, bid = 13)
(sid = 8, bid = 15)
(sid = 12, bid = 1)
...

Sort-Merge Join

Sailors

Reserves

Key idea:

Sort S and R **on join column**, then merge them!

Steps:

1. Sort S and R.
2. “Zip” or merge.

Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)
(name = Sam, sid = 3, bid = 6)

...

(name = Bob, sid = 1)
(name = Jill, sid = 2)
(name = Sam, sid = 3)
(name = Yue, sid = 4)
(name = Sue, sid = 7)

(name = Sue, sid = 8)
(name = Joe, sid = 12)
...

(sid = 1, bid = 4)
(sid = 1, bid = 7)
(sid = 3, bid = 6)
(sid = 4, bid = 3)
(sid = 8, bid = 1)

(sid = 8, bid = 13)
(sid = 8, bid = 15)
(sid = 12, bid = 1)
...

Sort-Merge Join

Sailors

Reserves

Key idea:

Sort S and R **on join column**, then merge them!

Steps:

1. Sort S and R.
2. “Zip” or merge.

Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)
(name = Sam, sid = 3, bid = 6)

...

(name = Bob, sid = 1)
(name = Jill, sid = 2)
(name = Sam, sid = 3)
(name = Yue, sid = 4)
(name = Sue, sid = 7)

(name = Sue, sid = 8)
(name = Joe, sid = 12)
...

(sid = 1, bid = 4)
(sid = 1, bid = 7)
(sid = 3, bid = 6)
(sid = 4, bid = 3)
(sid = 8, bid = 1)

(sid = 8, bid = 13)
(sid = 8, bid = 15)
(sid = 12, bid = 1)
...

Sort-Merge Join

Sailors

Reserves

Key idea:

Sort S and R **on join column**, then merge them!

Steps:

1. Sort S and R.
2. “Zip” or merge.

Output:

(name = Bob, sid = 1, bid = 4)
(name = Bob, sid = 1, bid = 7)
(name = Sam, sid = 3, bid = 6)

...

(name = Bob, sid = 1)
(name = Jill, sid = 2)
(name = Sam, sid = 3)
(name = Yue, sid = 4)
(name = Sue, sid = 7)

(name = Sue, sid = 8)
(name = Joe, sid = 12)
...

(sid = 1, bid = 4)
(sid = 1, bid = 7)
(sid = 3, bid = 6)
(sid = 4, bid = 3)
(sid = 8, bid = 1)

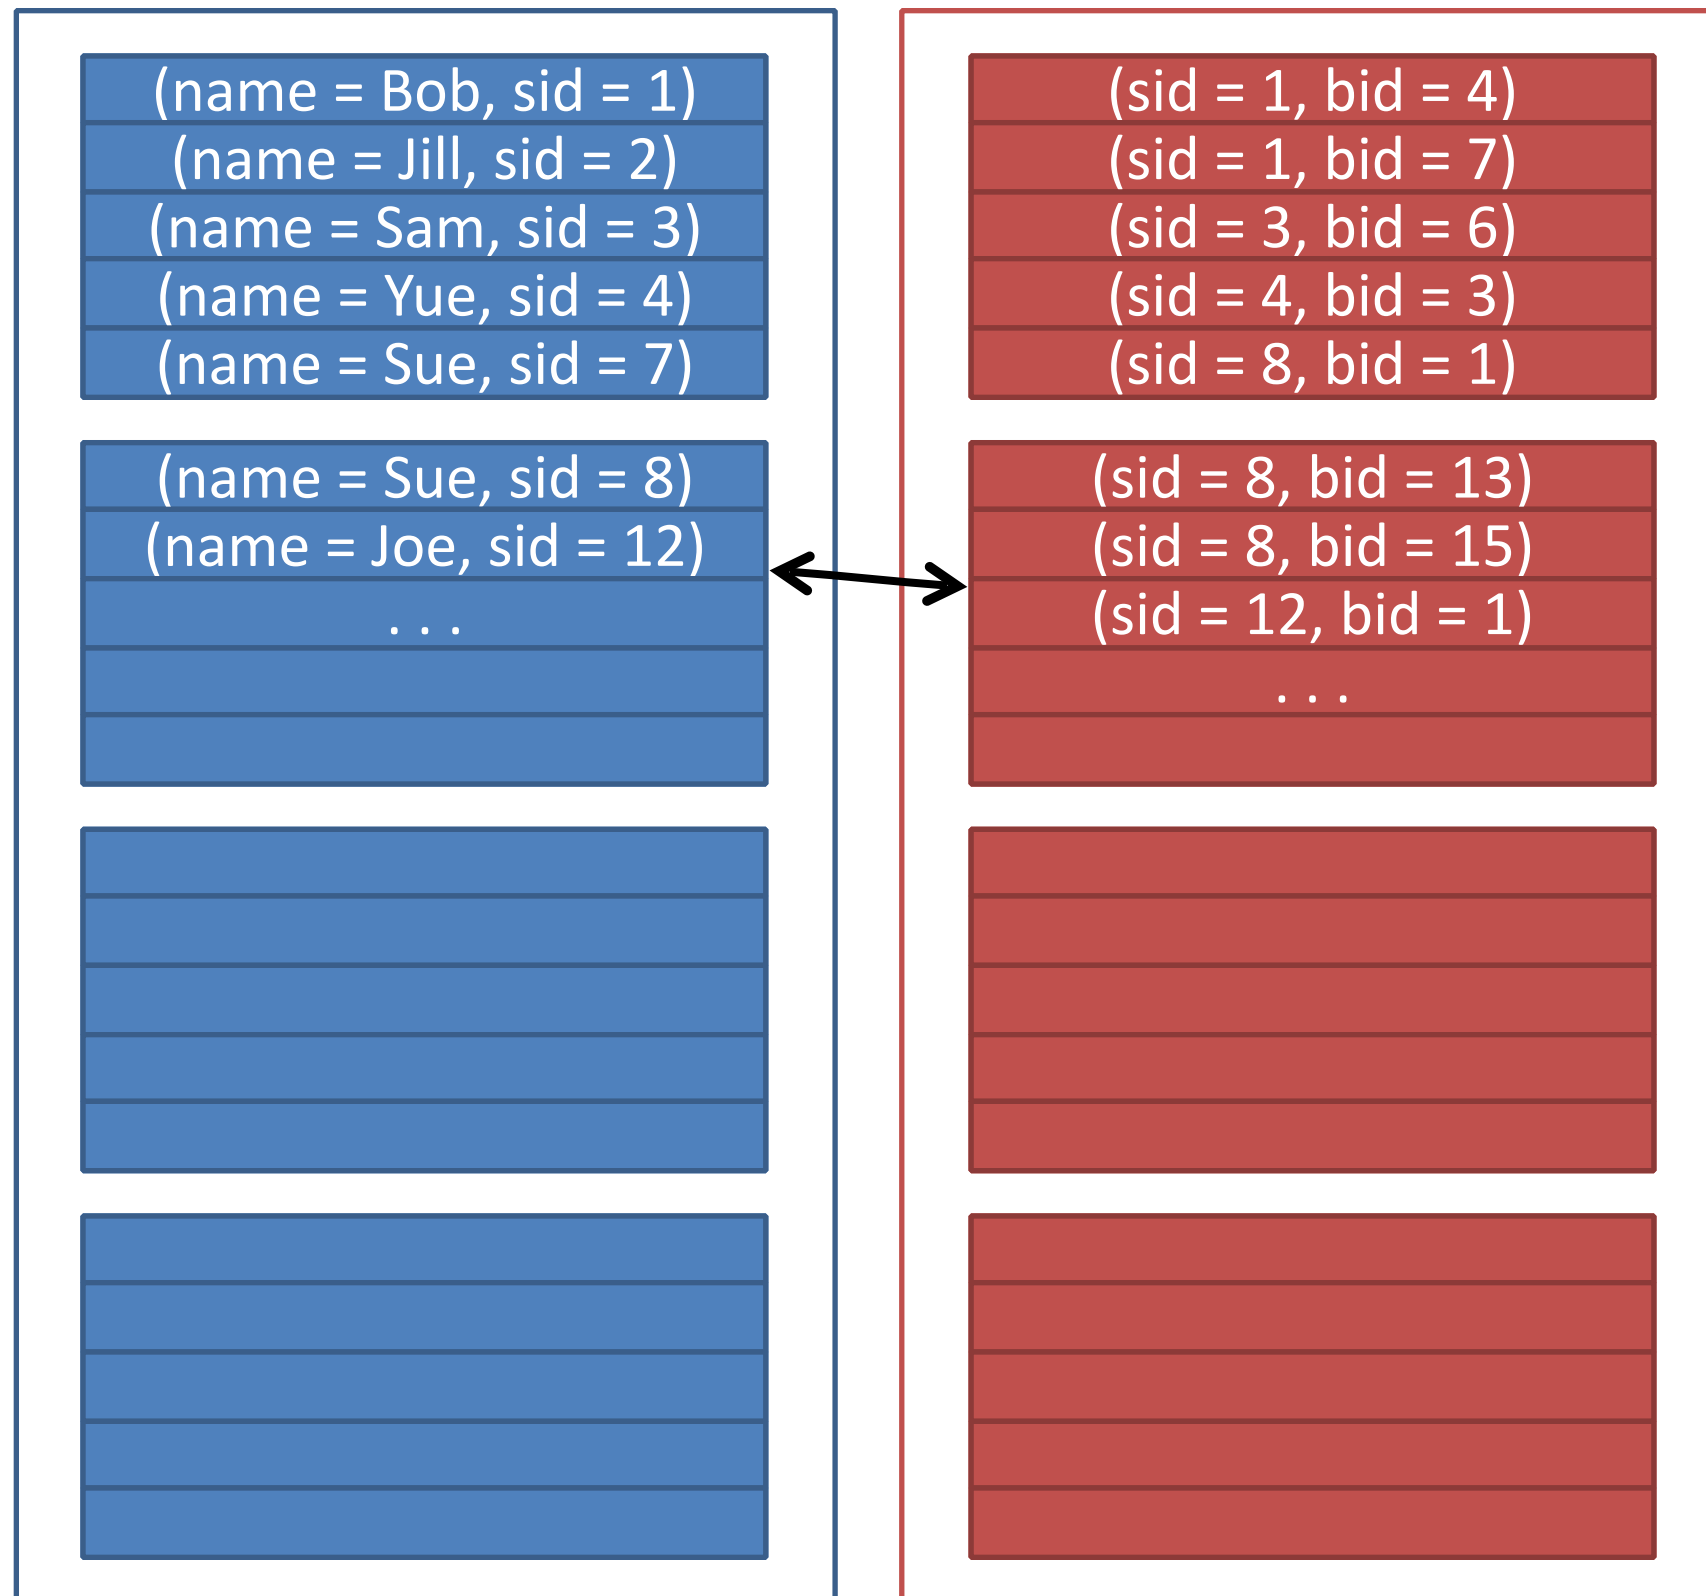
(sid = 8, bid = 13)
(sid = 8, bid = 15)
(sid = 12, bid = 1)
...

Notation: $[S] == \text{"\# pages in S"} ; |S| == \text{"\# tuples in S"}$

Sort-Merge Join

Sailors

Reserves



Key idea:

Sort S and R **on join column**, then merge them!

Steps:

1. Sort S and R.
2. “Zip” or merge.

I/Os:

$$\sim 5([S] + [R])$$

$$\text{Sorting: } 4([S] + [R])$$

$$\text{Merging: } [S] + [R]$$

Optimizing Sort-Merge Join

Sailors

Reserves

(name = Bob, sid = 1)
(name = Jill, sid = 2)
(name = Sam, sid = 3)
(name = Yue, sid = 4)
(name = Sue, sid = 7)

(name = Sue, sid = 8)
(name = Joe, sid = 12)
...

(sid = 1, bid = 4)
(sid = 1, bid = 7)
(sid = 3, bid = 6)
(sid = 4, bid = 3)
(sid = 8, bid = 1)

(sid = 8, bid = 13)
(sid = 8, bid = 15)
(sid = 12, bid = 1)
...

Key idea:

Internal Sort on
both. Perform
merge on all runs!

Steps:

1. Internal sort S and R. (Pass 0)
2. Merge all runs.

Optimizing Sort-Merge Join

Sailors

Reserves

Key idea:

Internal Sort on
both. Perform
merge on all runs!

Steps:

1. Internal sort S and R. (Pass 0)
2. Merge all runs.

(name = Bob, sid = 1)
(name = Jill, sid = 2)
(name = Yue, sid = 4)
(name = Sue, sid = 8)
(name = Jack, sid = 18)

(name = Cat, sid = 22)

...

(name = Sam, sid = 3)
(name = Sue, sid = 7)
(name = Joe, sid = 12)

...

(sid = 1, bid = 4)

(sid = 1, bid = 7)

(sid = 4, bid = 3)

(sid = 8, bid = 1)

(sid = 8, bid = 13)

(sid = 12, bid = 1)

...

(sid = 3, bid = 6)

(sid = 8, bid = 15)

...

Optimizing Sort-Merge Join

Sailors

Reserves

Key idea:

Internal Sort on both.
Perform merge on all runs!

Steps:

1. Internal sort S and R.
(Pass 0)

merge all runs.

NOTE: What does this
assume about the number
of runs?

$\sim 3([S] + [R])$

Pass 0: $2([S] + [R])$

Merging: $[S] + [R]$

(name = Bob, sid = 1)
(name = Jill, sid = 2)
(name = Yue, sid = 4)
(name = Sue, sid = 8)
(name = Jack, sid = 18)

(name = Cat, sid = 22)
...


(name = Sam, sid = 3)
(name = Sue, sid = 7)
(name = Joe, sid = 12)
...

(sid = 1, bid = 4)
(sid = 1, bid = 7)
(sid = 4, bid = 3)
(sid = 8, bid = 1)
(sid = 8, bid = 13)


(sid = 12, bid = 1)
...

Hash-Join

Sailors



Reserves



Key idea:

Partition S and R
using same hash fn,
then collect same
partitions

Steps:

1. Partition S and R
2. Re-Hash, collect

Hash-Join

Sailors

Reserves

Key idea:

Partition S and R
using same hash fn,
then collect same
partitions

Steps:

1. Partition S and R
2. Re-Hash, collect

(name = Bob, sid = 1)

(name = Sam, sid = 3)

(name = Sue, sid = 7)

(name = Jill, sid = 2)

(name = Joe, sid = 12)

(name = Sue, sid = 8)

(name = Yue, sid = 4)

Hash function: sid mod 4 Hash-Join

Sailors

Reserves

Key idea:

Partition S and R
using same hash fn,
then collect same
partitions

Steps:

1. Partition S and R
2. Re-Hash, collect

(name = Joe, sid = 12)
(name = Sue, sid = 8)
(name = Yue, sid = 4)
...

(name = Bob, sid = 1)
...

(name = Jill, sid = 2)
...

(name = Sue, sid = 7)
(name = Sam, sid = 3)
...

(sid = 12, bid = 1)
(sid = 8, bid = 13)
(sid = 8, bid = 15)
(sid = 4, bid = 3)
(sid = 8, bid = 1)

(sid = 1, bid = 4)
(sid = 1, bid = 7)
...

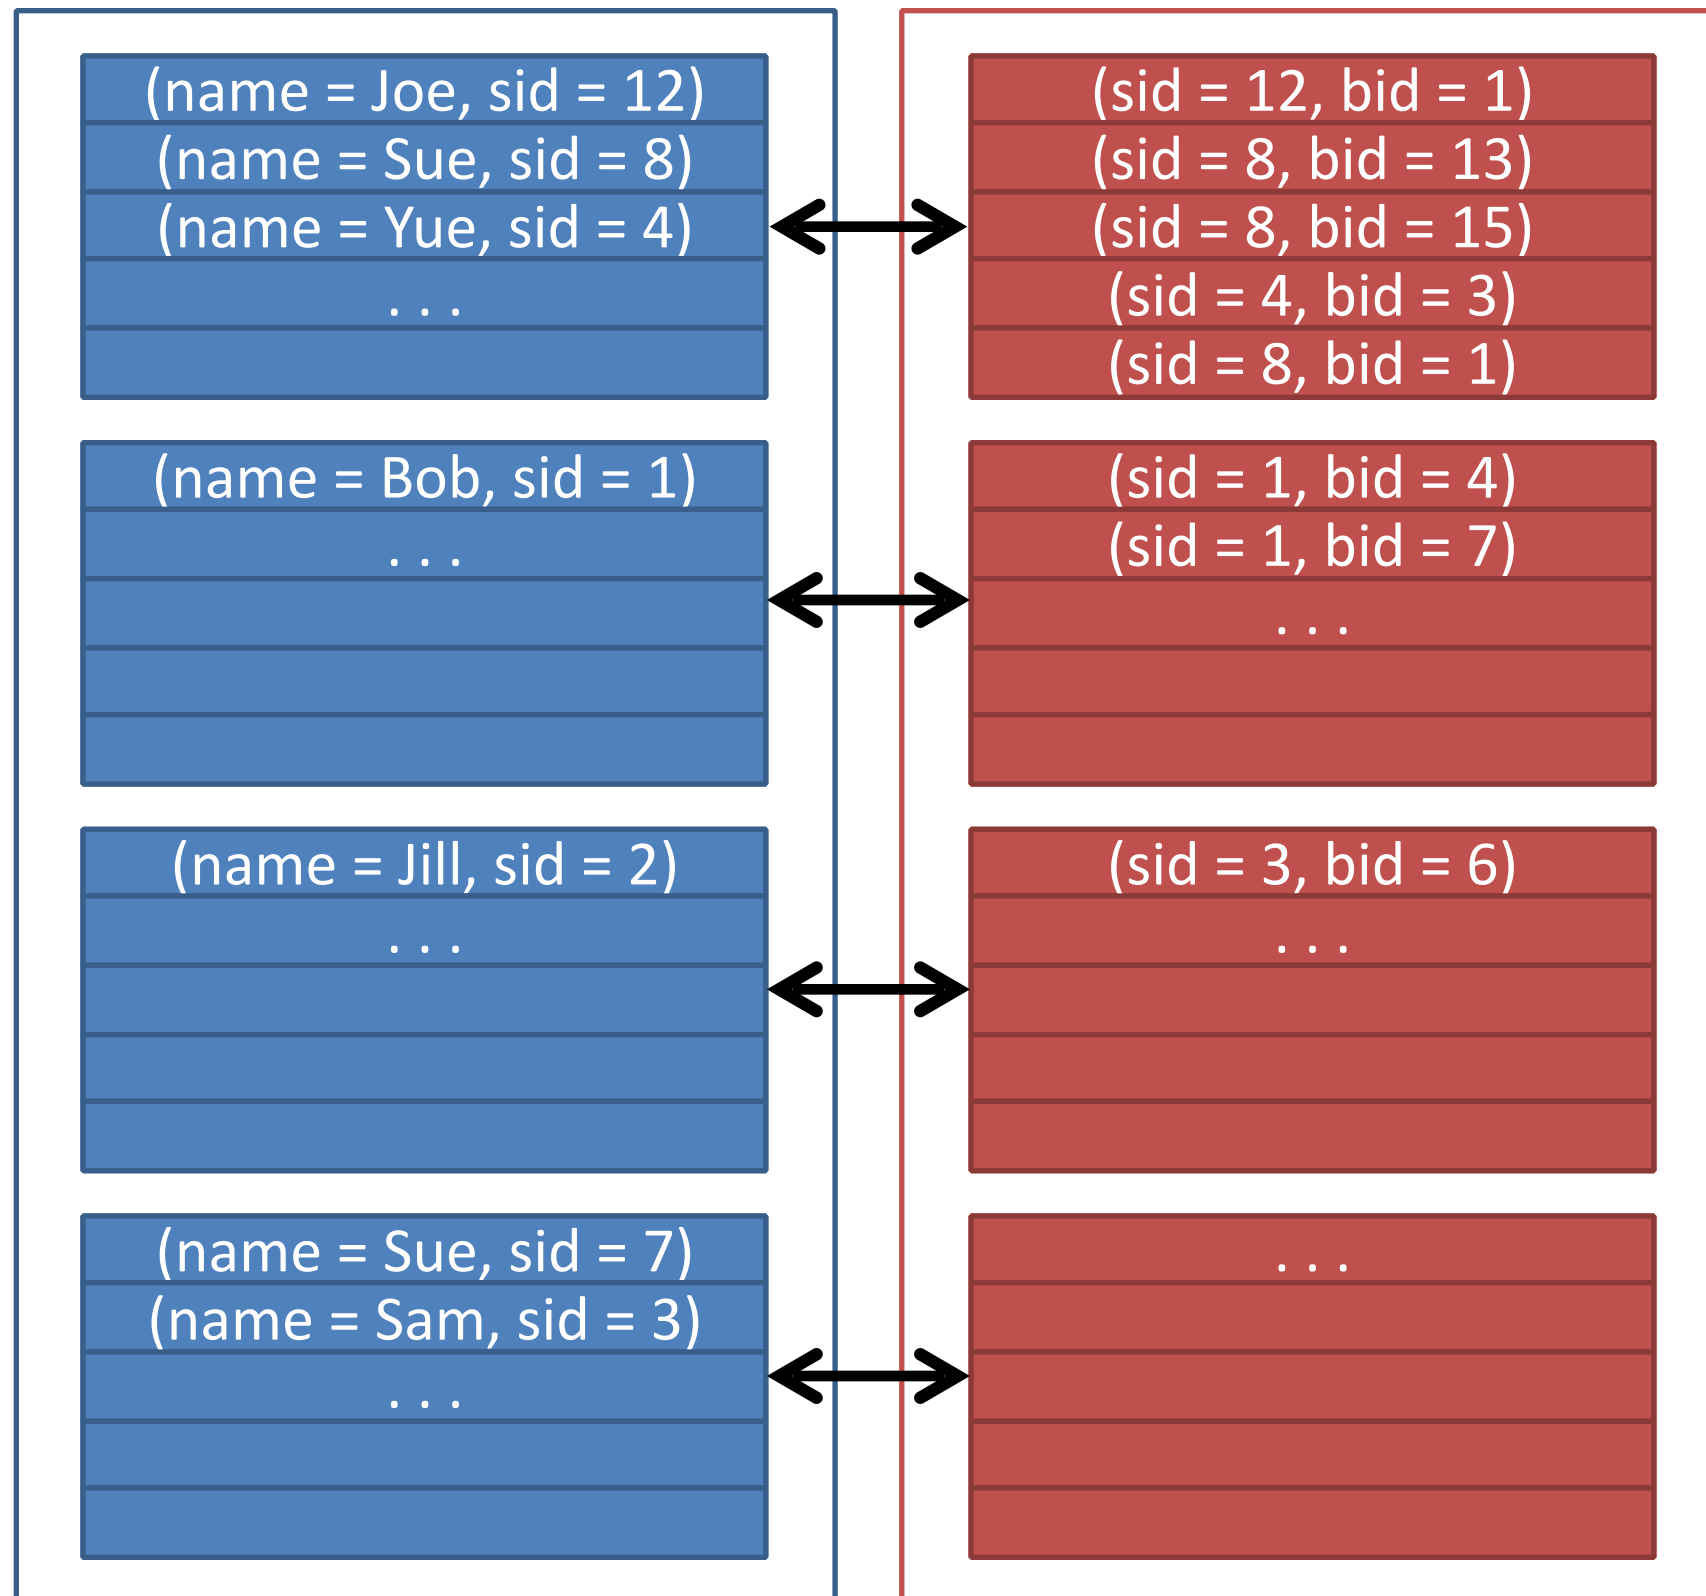
(sid = 3, bid = 6)
...

...

Hash-Join

Sailors

Reserves



Key idea:

Partition S and R
using same hash fn,
then collect same
partitions

Steps:

1. Partition S and R
2. Re-Hash, collect

Notation: $[S]$ == “# pages in S ” ; $|S|$ == “# tuples in S ”

Hash-Join

Sailors

Reserves

Key idea:

Partition S and R
using same hash fn,
then collect same
partitions

Steps:

1. Partition S and R
Hash, collect

2. $3([S] + [R])$

Partition: $2([S] + [R])$

Re-Hash: $[S] + [R]$

NOTE: This is no different
from what we previously
assumed about hashing.

NOTE: What are we
assuming about the size of
partitions?

(name = Joe, sid = 12)
(name = Sue, sid = 8)
(name = Yue, sid = 4)
...


(sid = 12, bid = 1)
(sid = 8, bid = 13)
(sid = 8, bid = 15)
(sid = 4, bid = 3)
(sid = 8, bid = 1)

(name = Bob,
...)

(name = Jill, sid = 2)
...

(name = Sue, sid = 7)
(name = Sam, sid = 3)
...

Hybrid Hashing




$N \leq B$: Load all data into memory and hash

Hybrid Hashing




$N \leq B$: Load all data into memory and hash

Hybrid Hashing



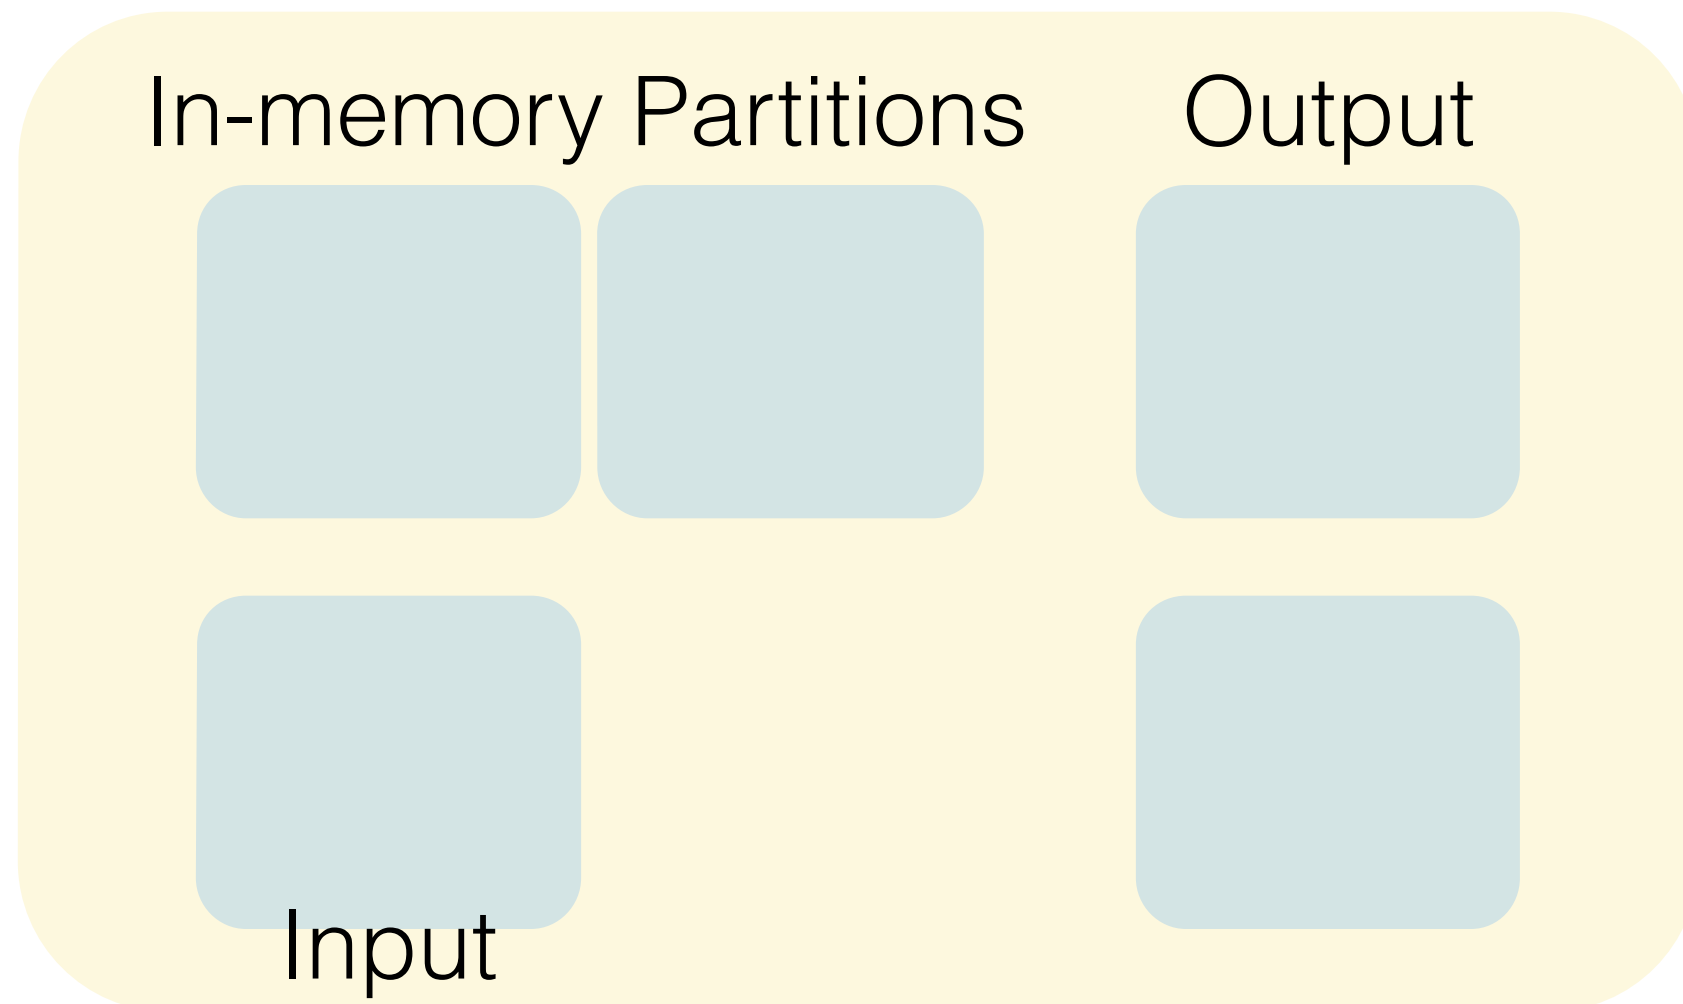

$N \leq B$: Load all data into memory and hash

Hybrid Hashing



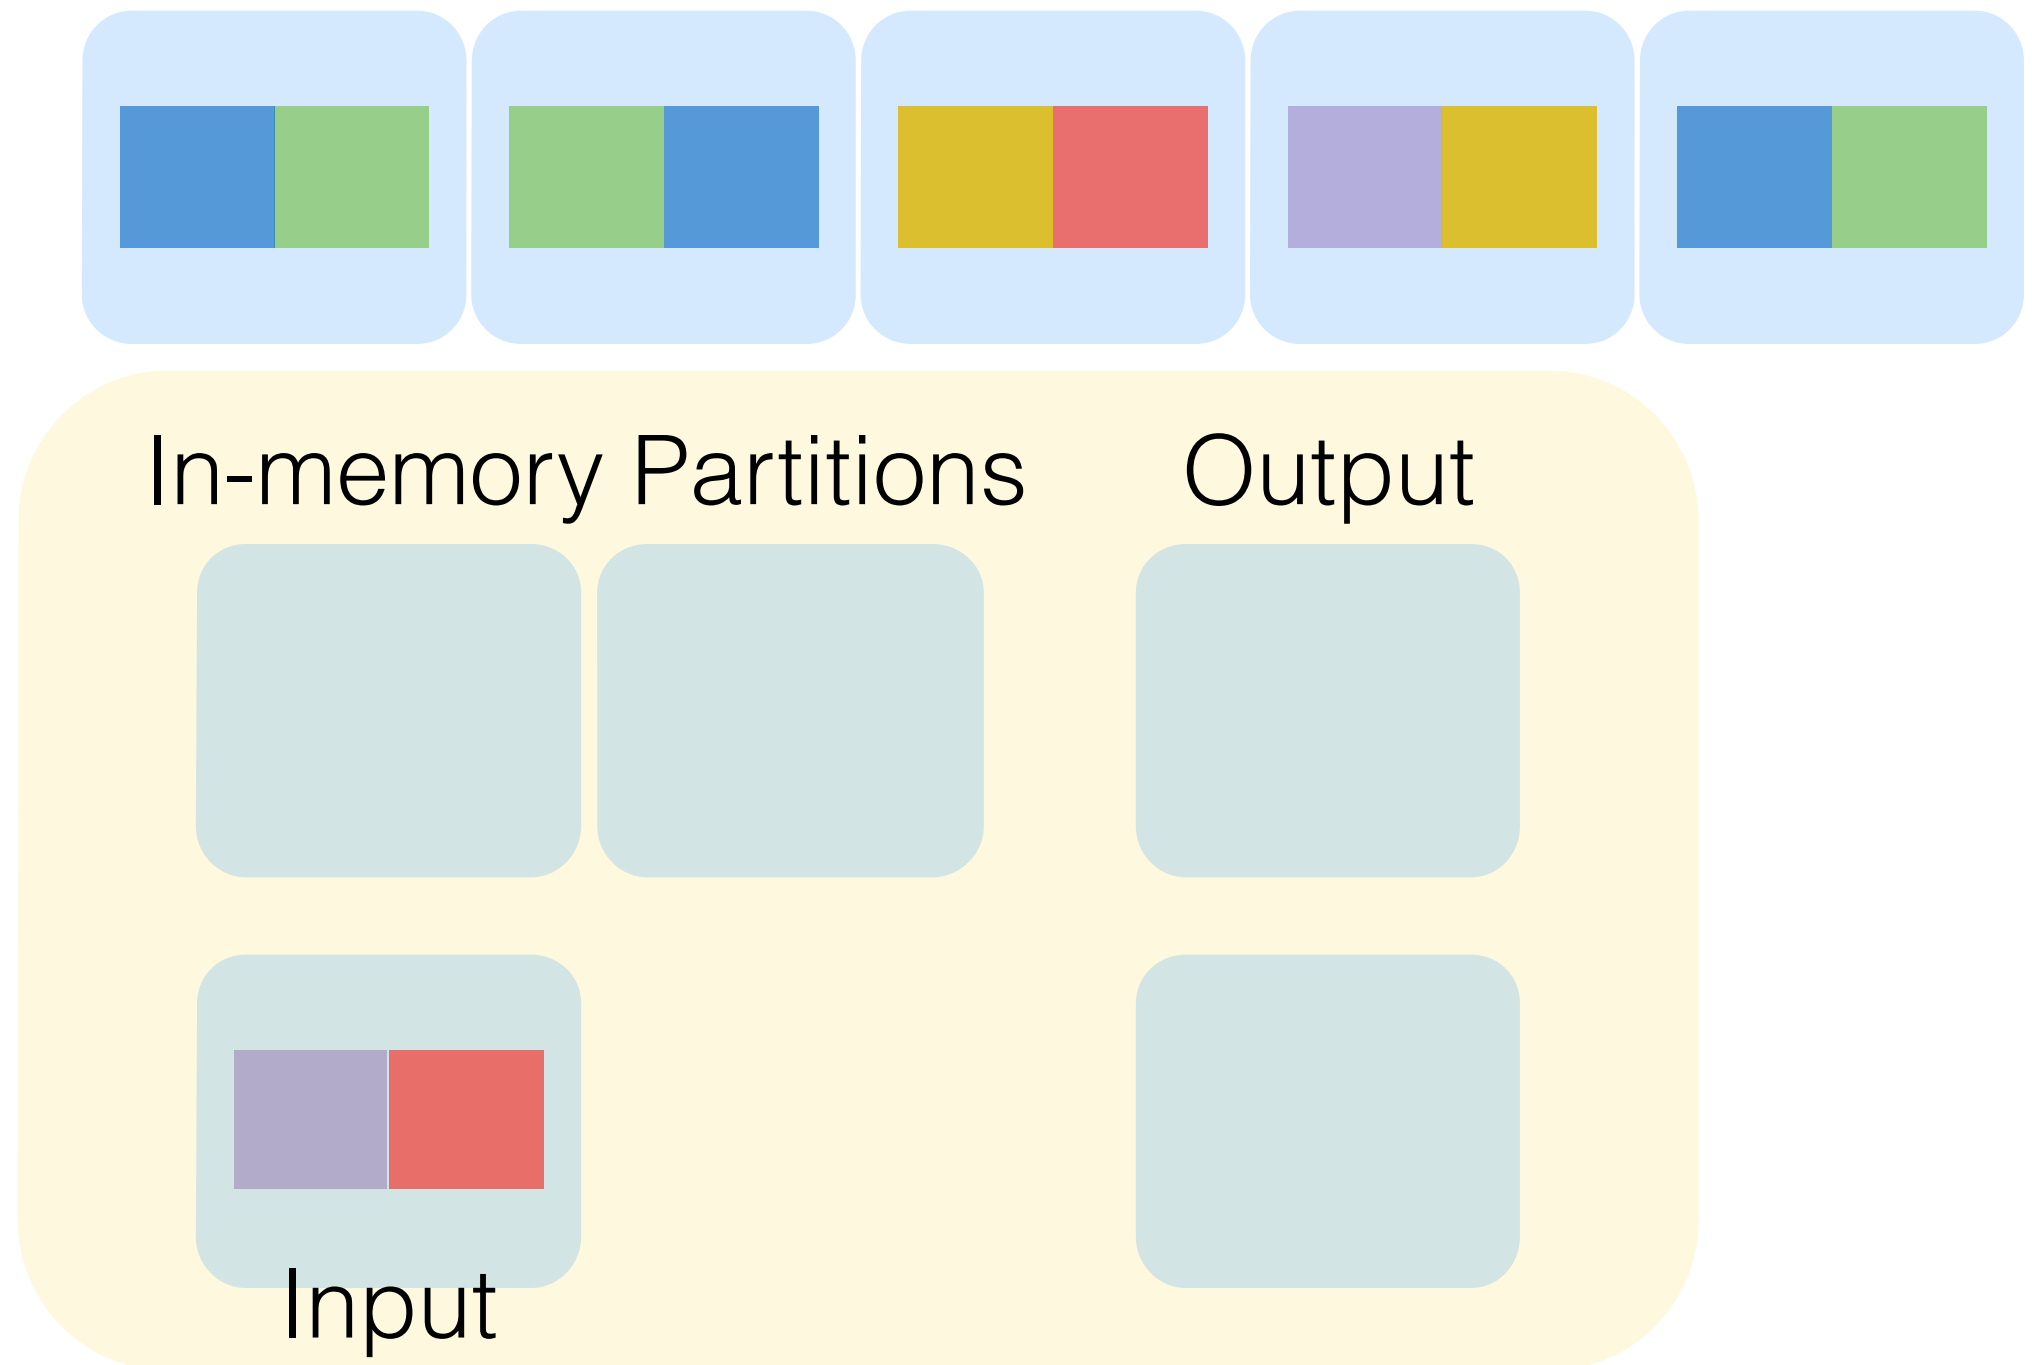
$N \geq B$: Must perform at least 2 passes

Hybrid Hashing



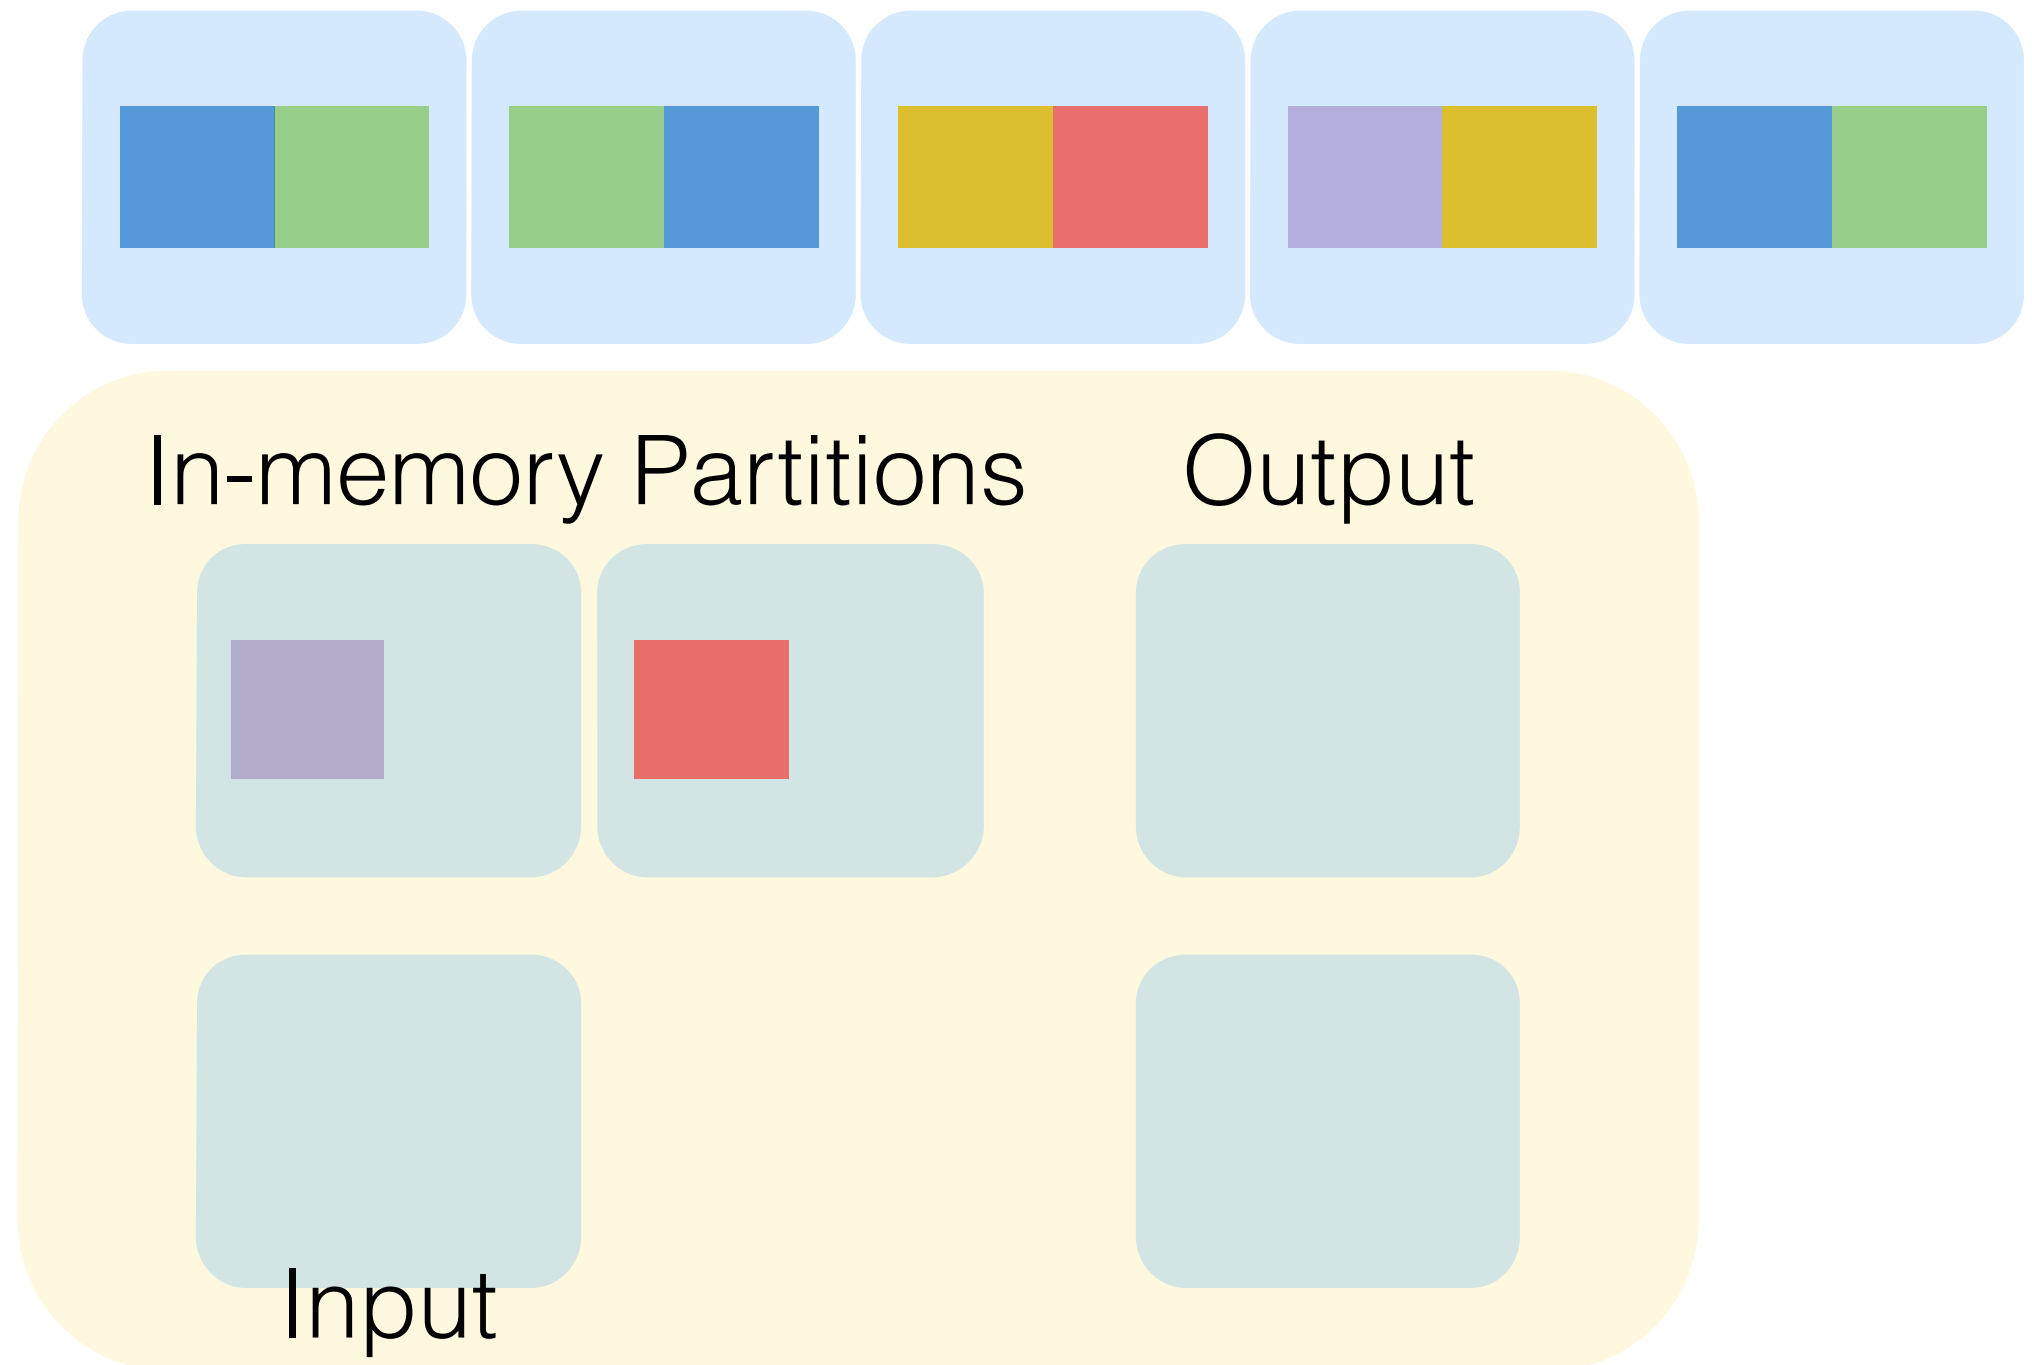
Solution: Keep some partitions in memory

Hybrid Hashing



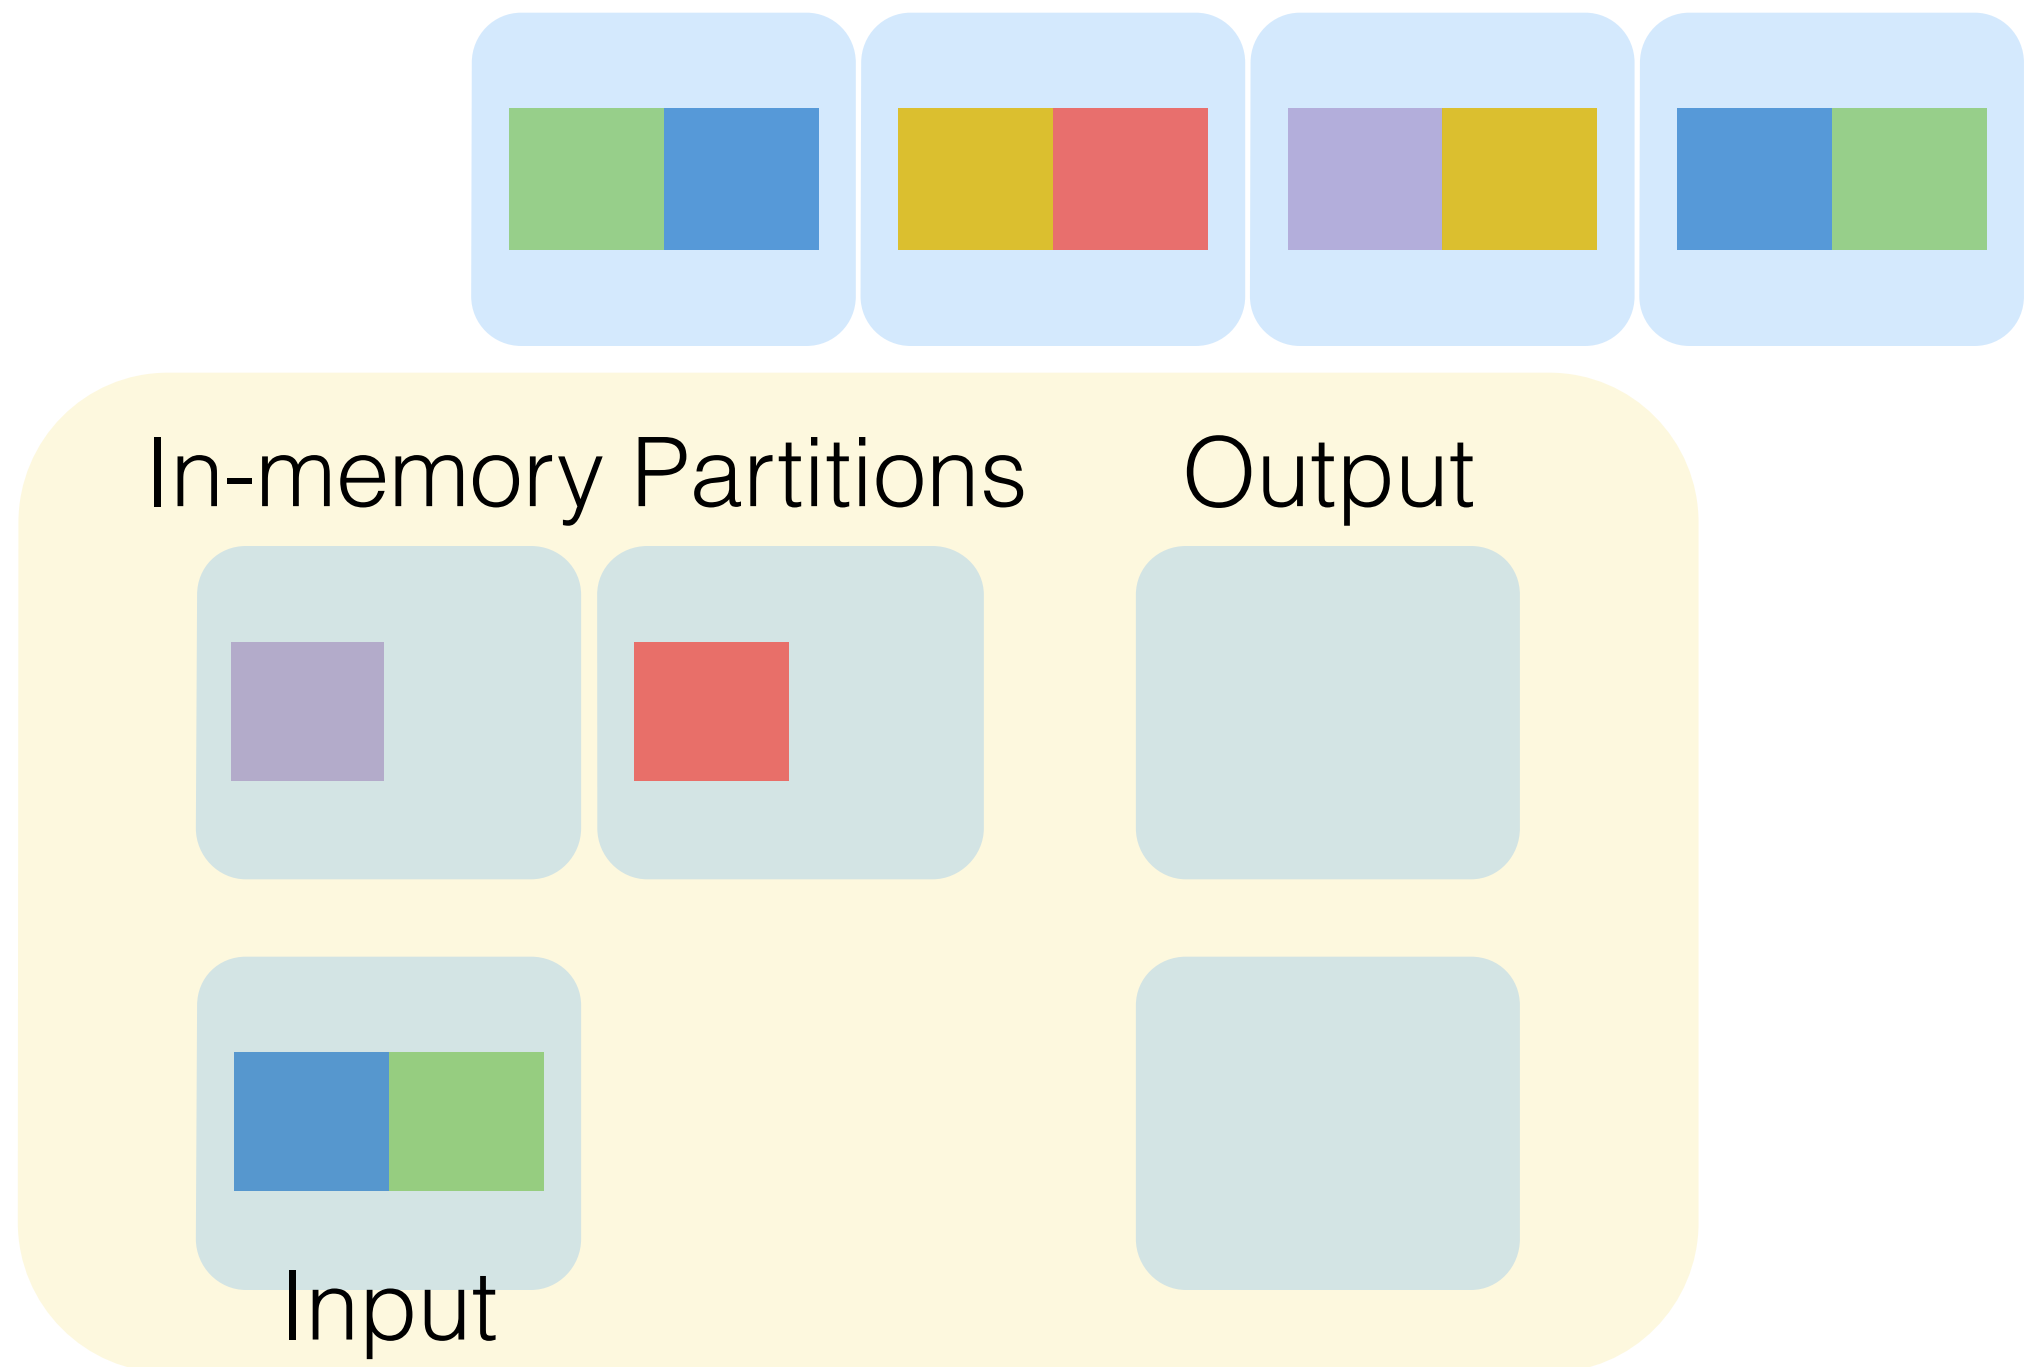
Solution: Keep some partitions in memory

Hybrid Hashing



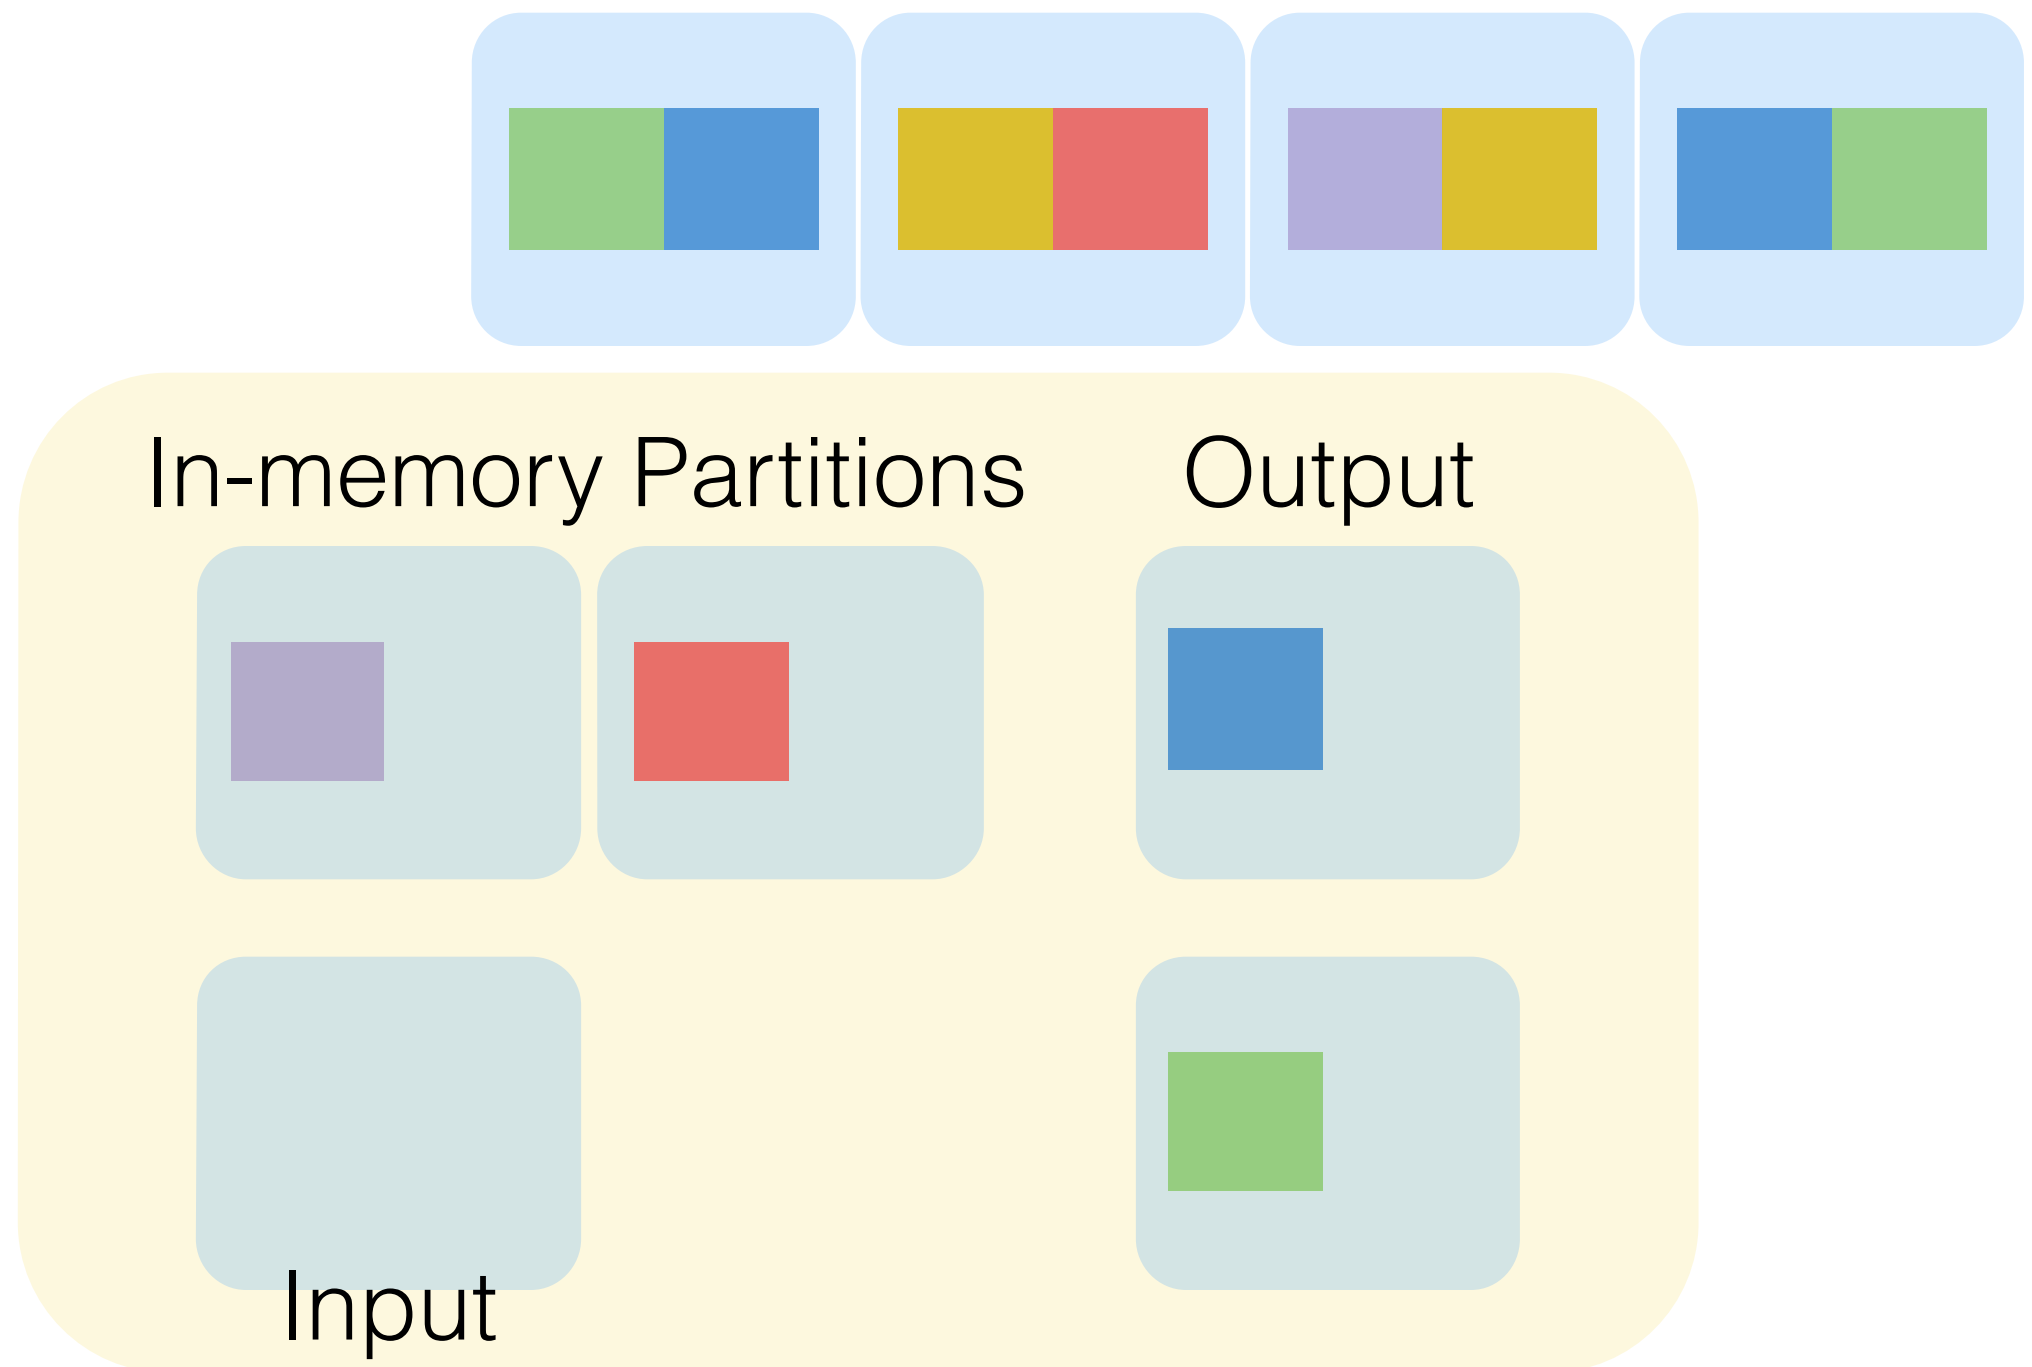
Solution: Keep some partitions in memory

Hybrid Hashing



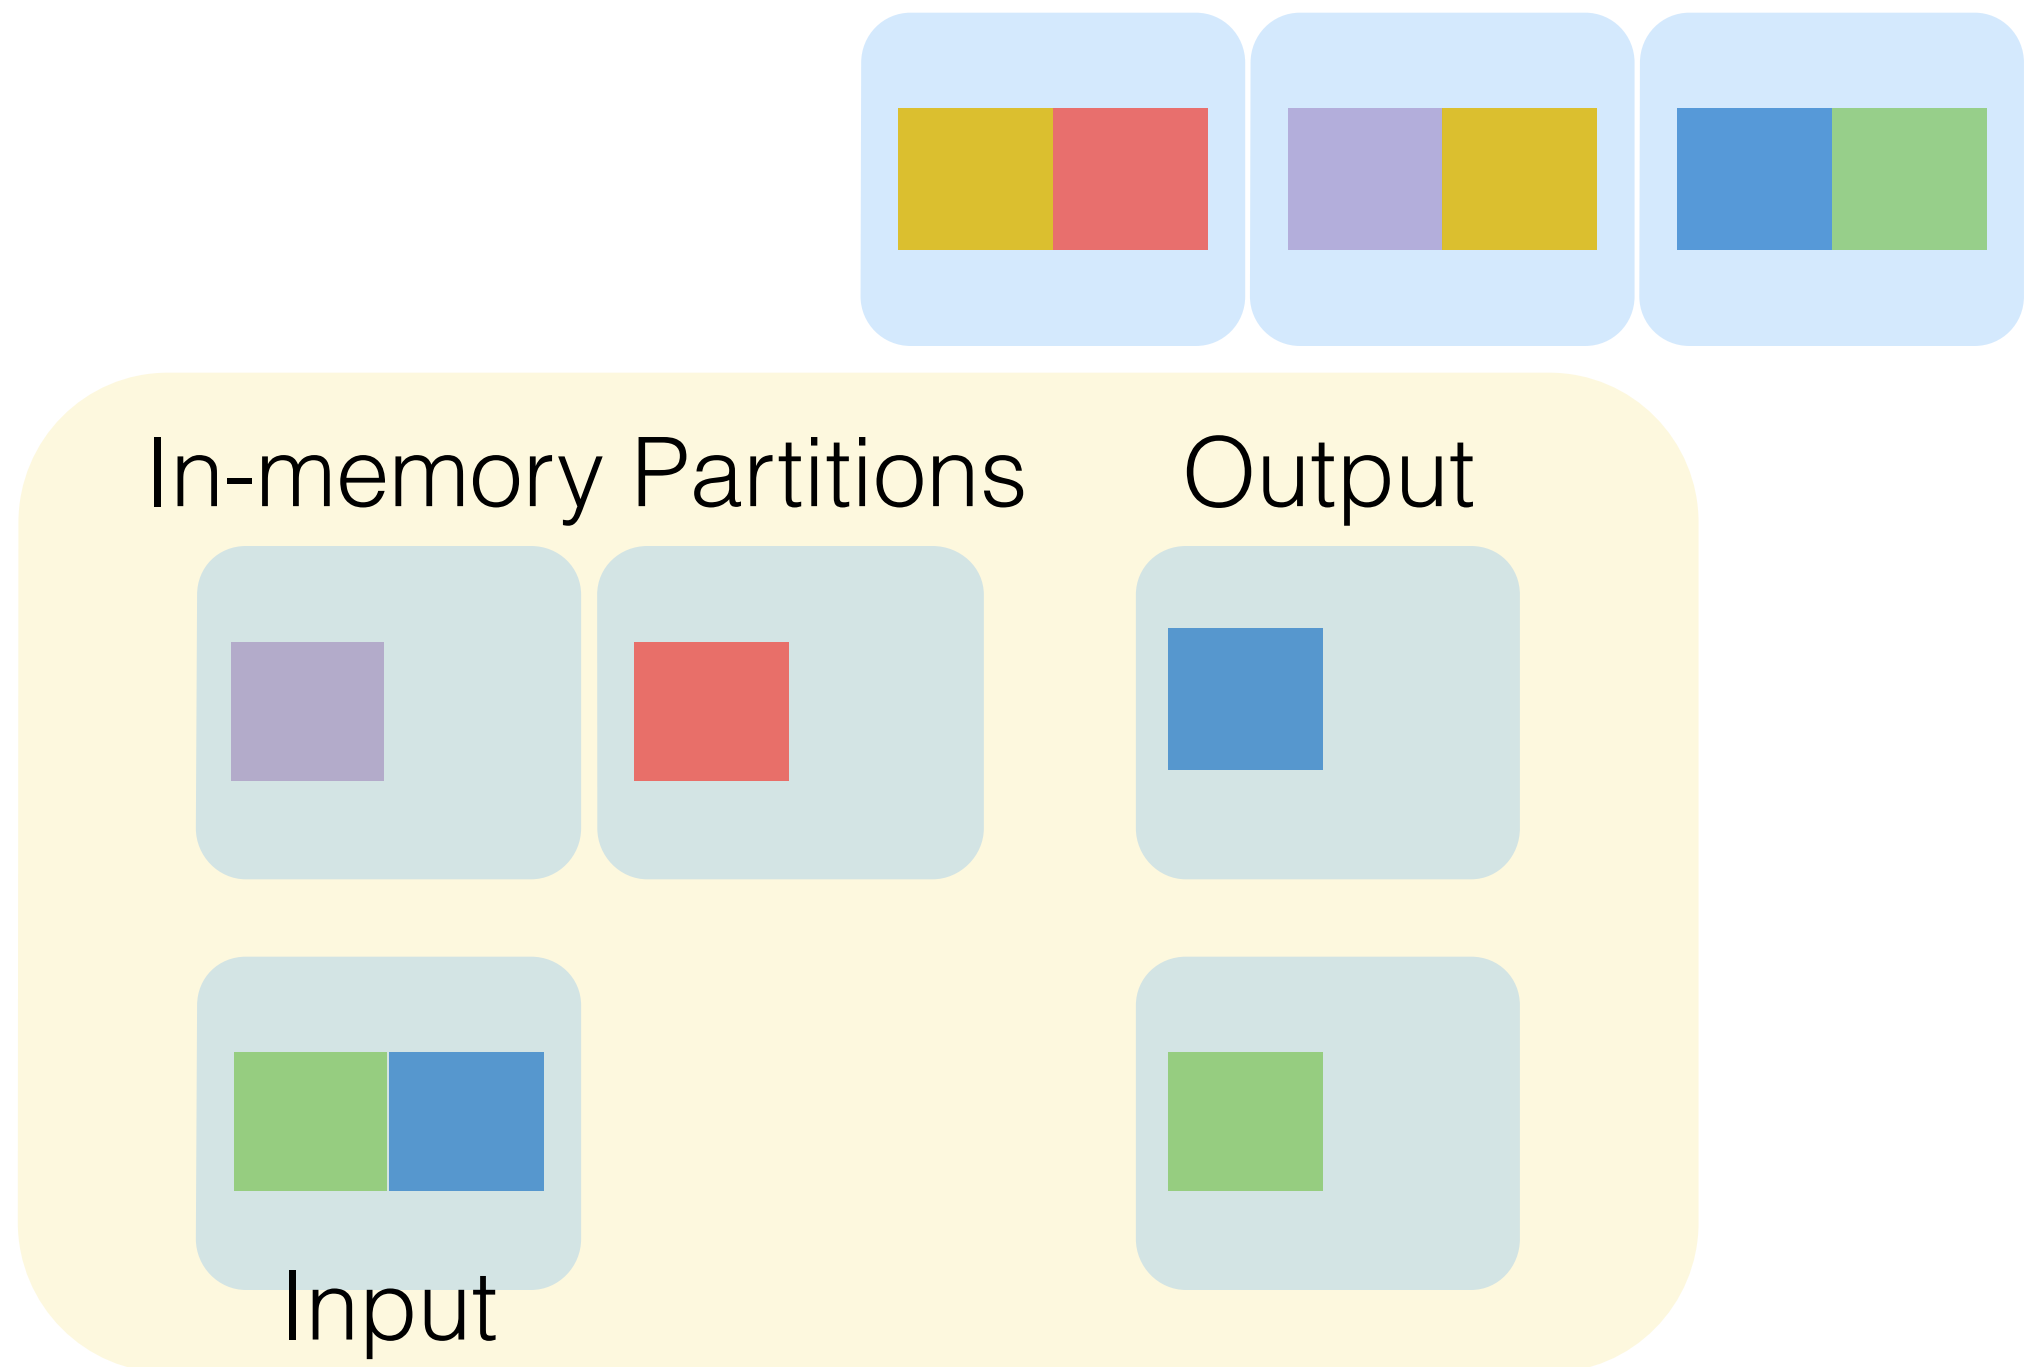
Solution: Keep some partitions in memory

Hybrid Hashing



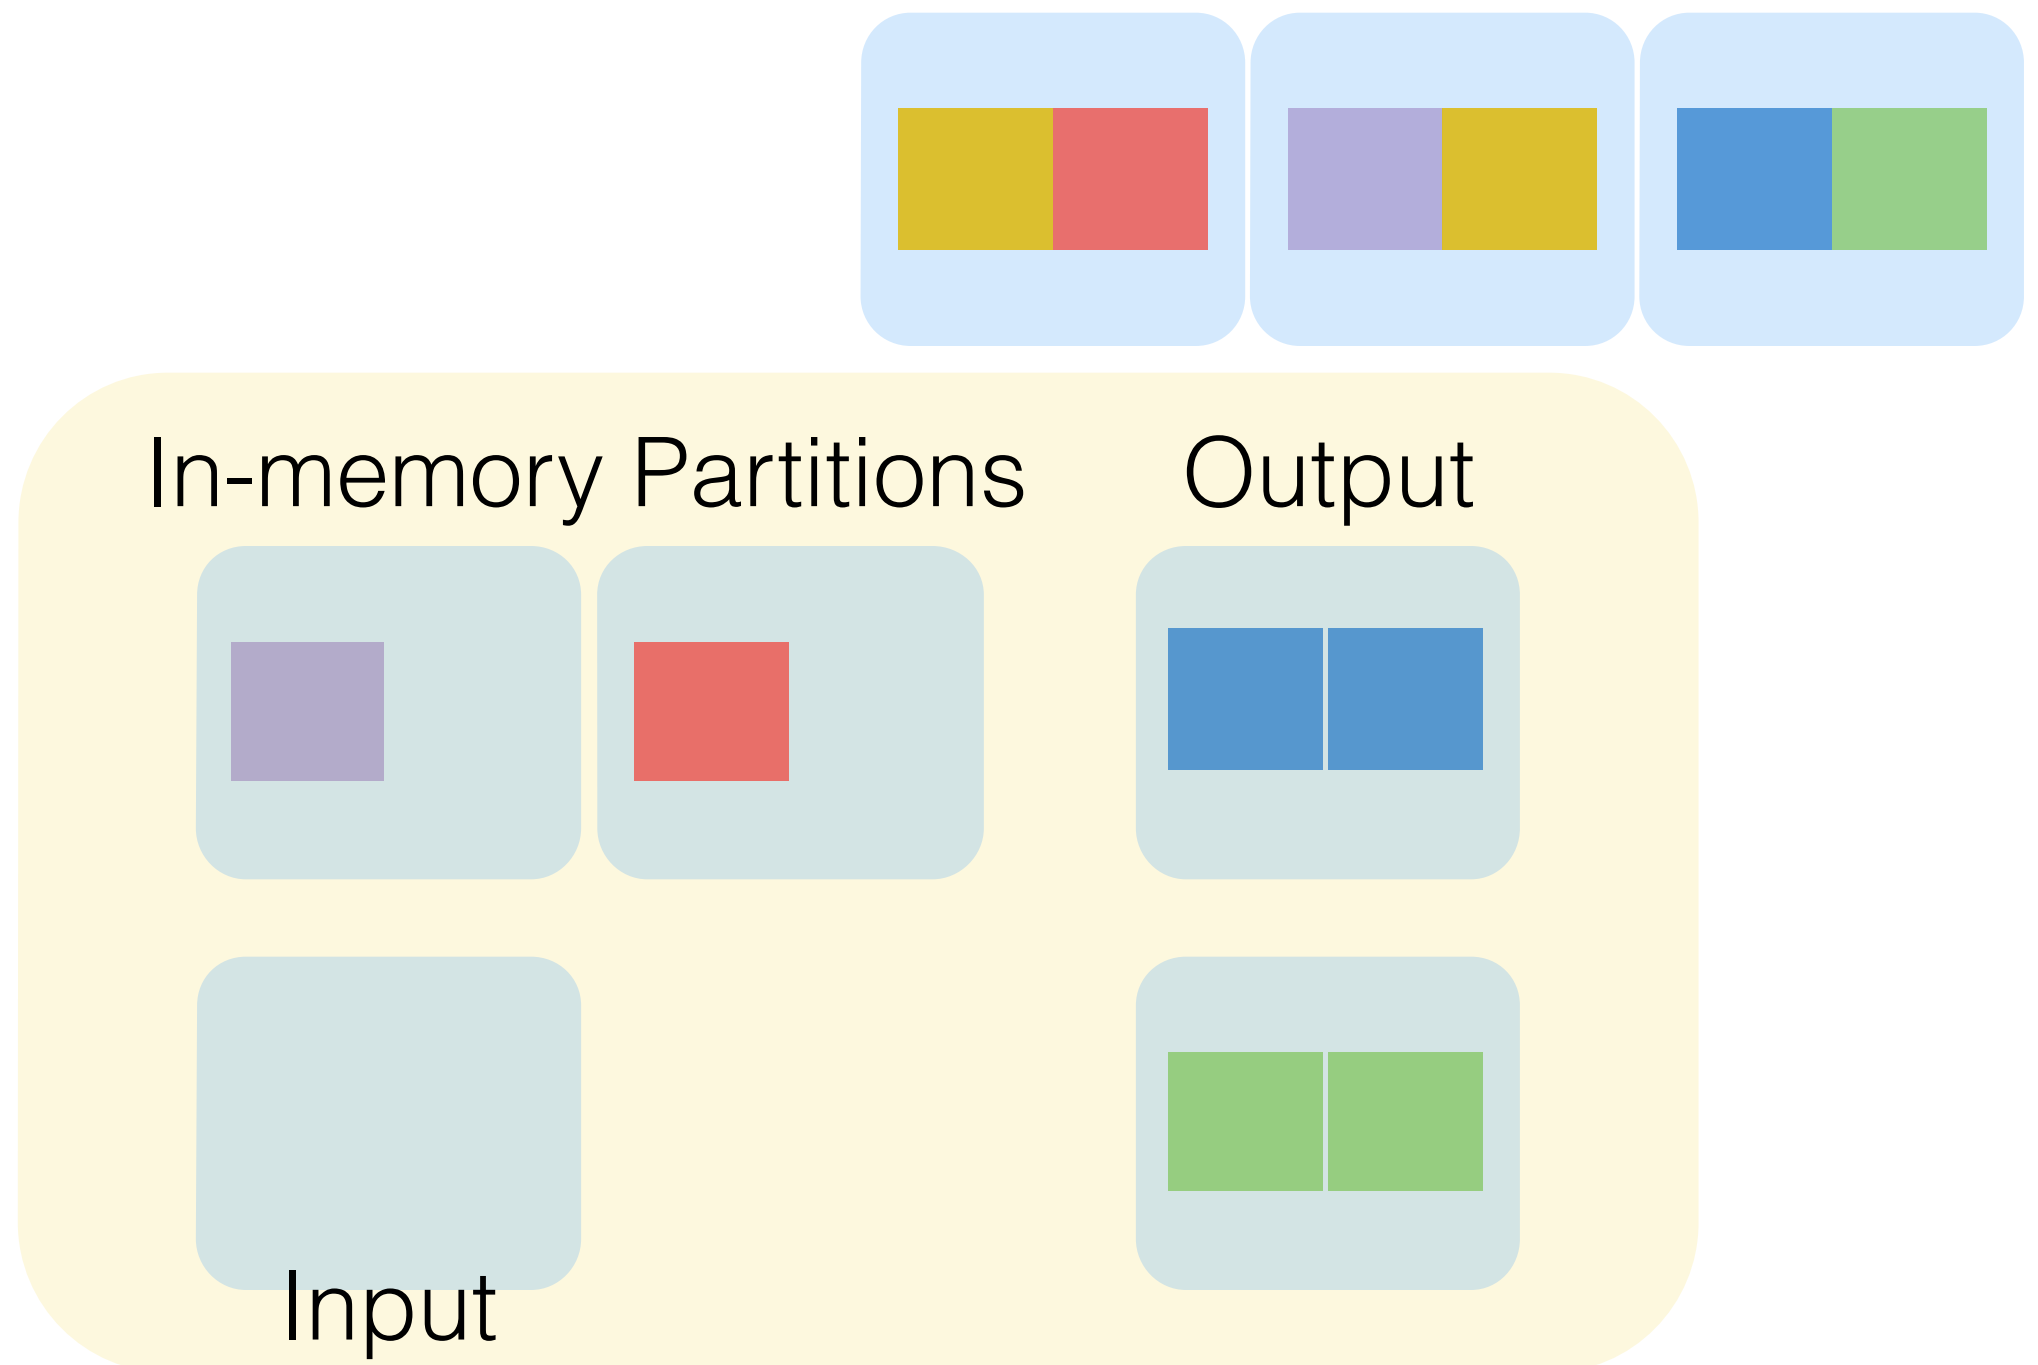
Solution: Keep some partitions in memory

Hybrid Hashing




Solution: Keep some partitions in memory

Hybrid Hashing



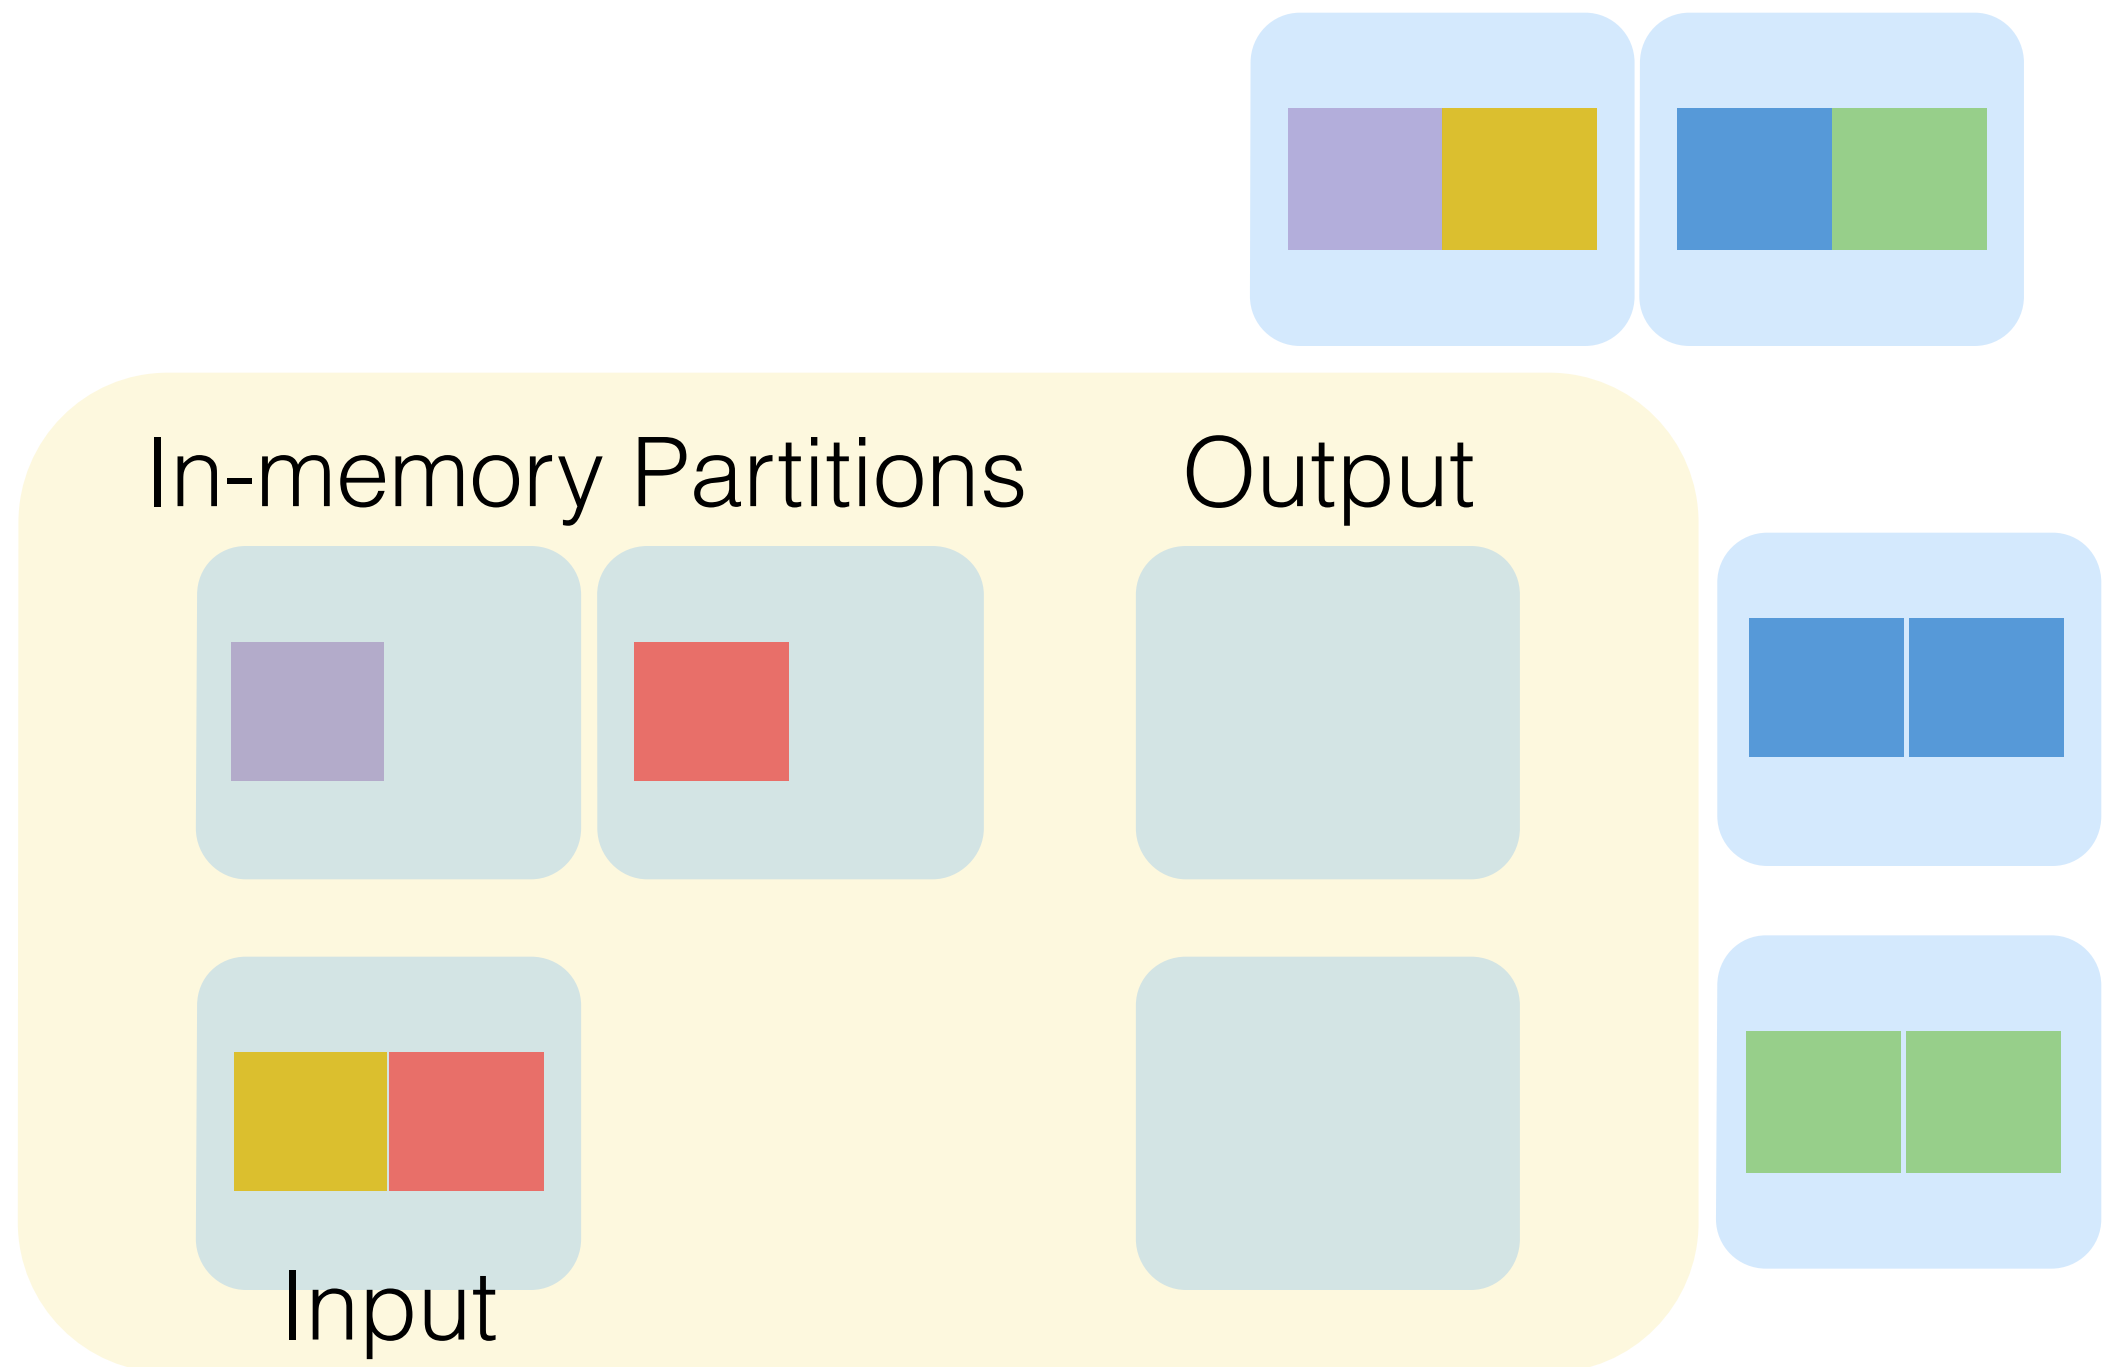
Solution: Keep some partitions in memory

Hybrid Hashing



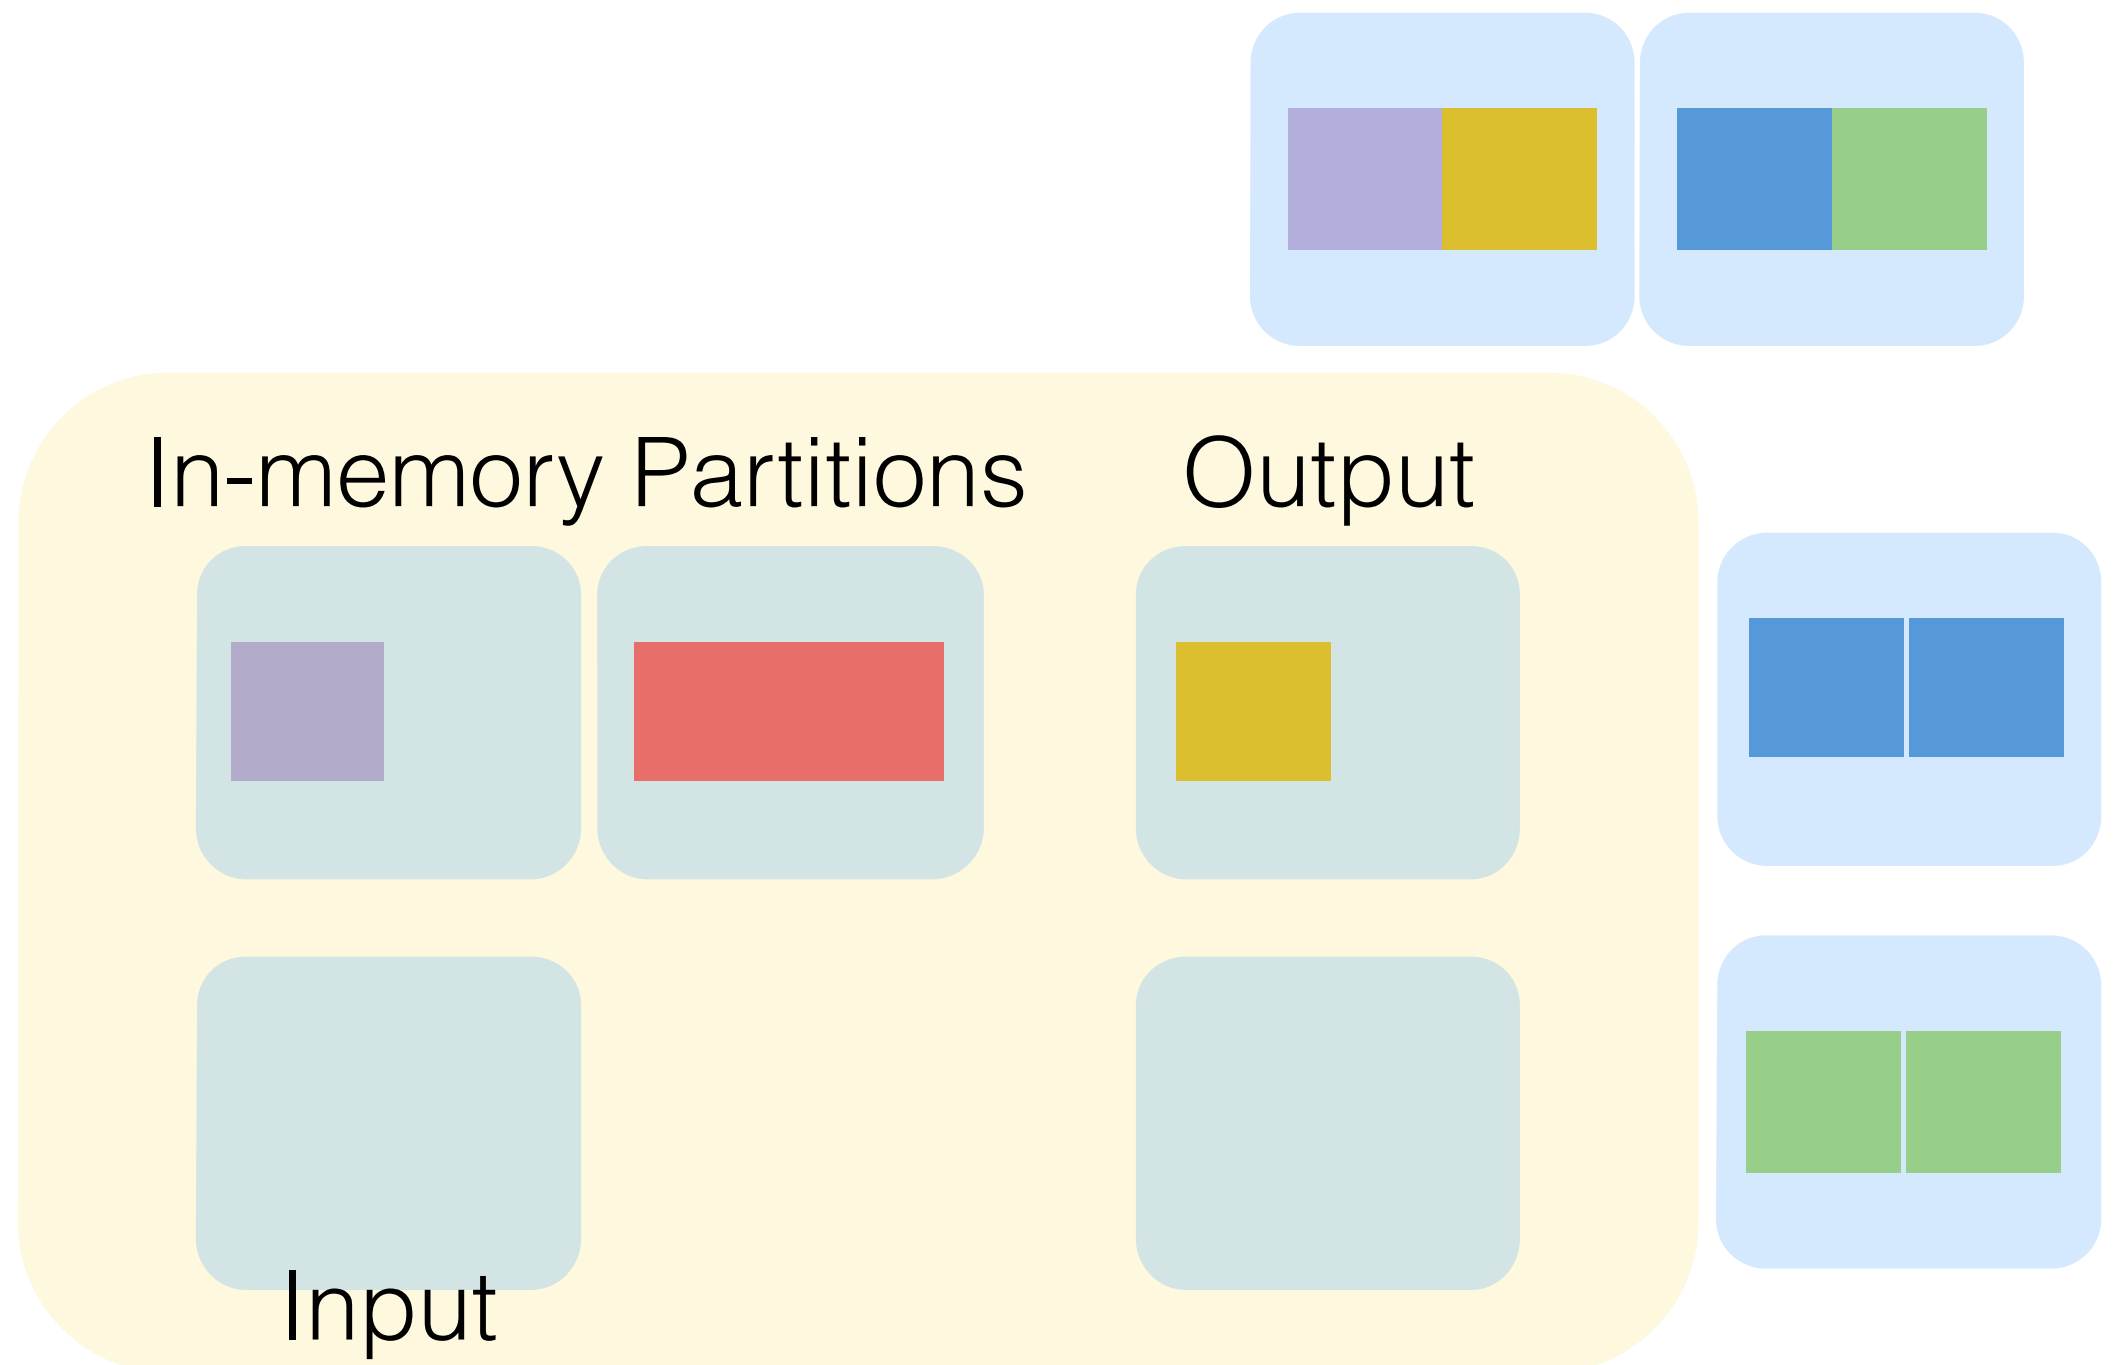
Solution: Keep some partitions in memory

Hybrid Hashing



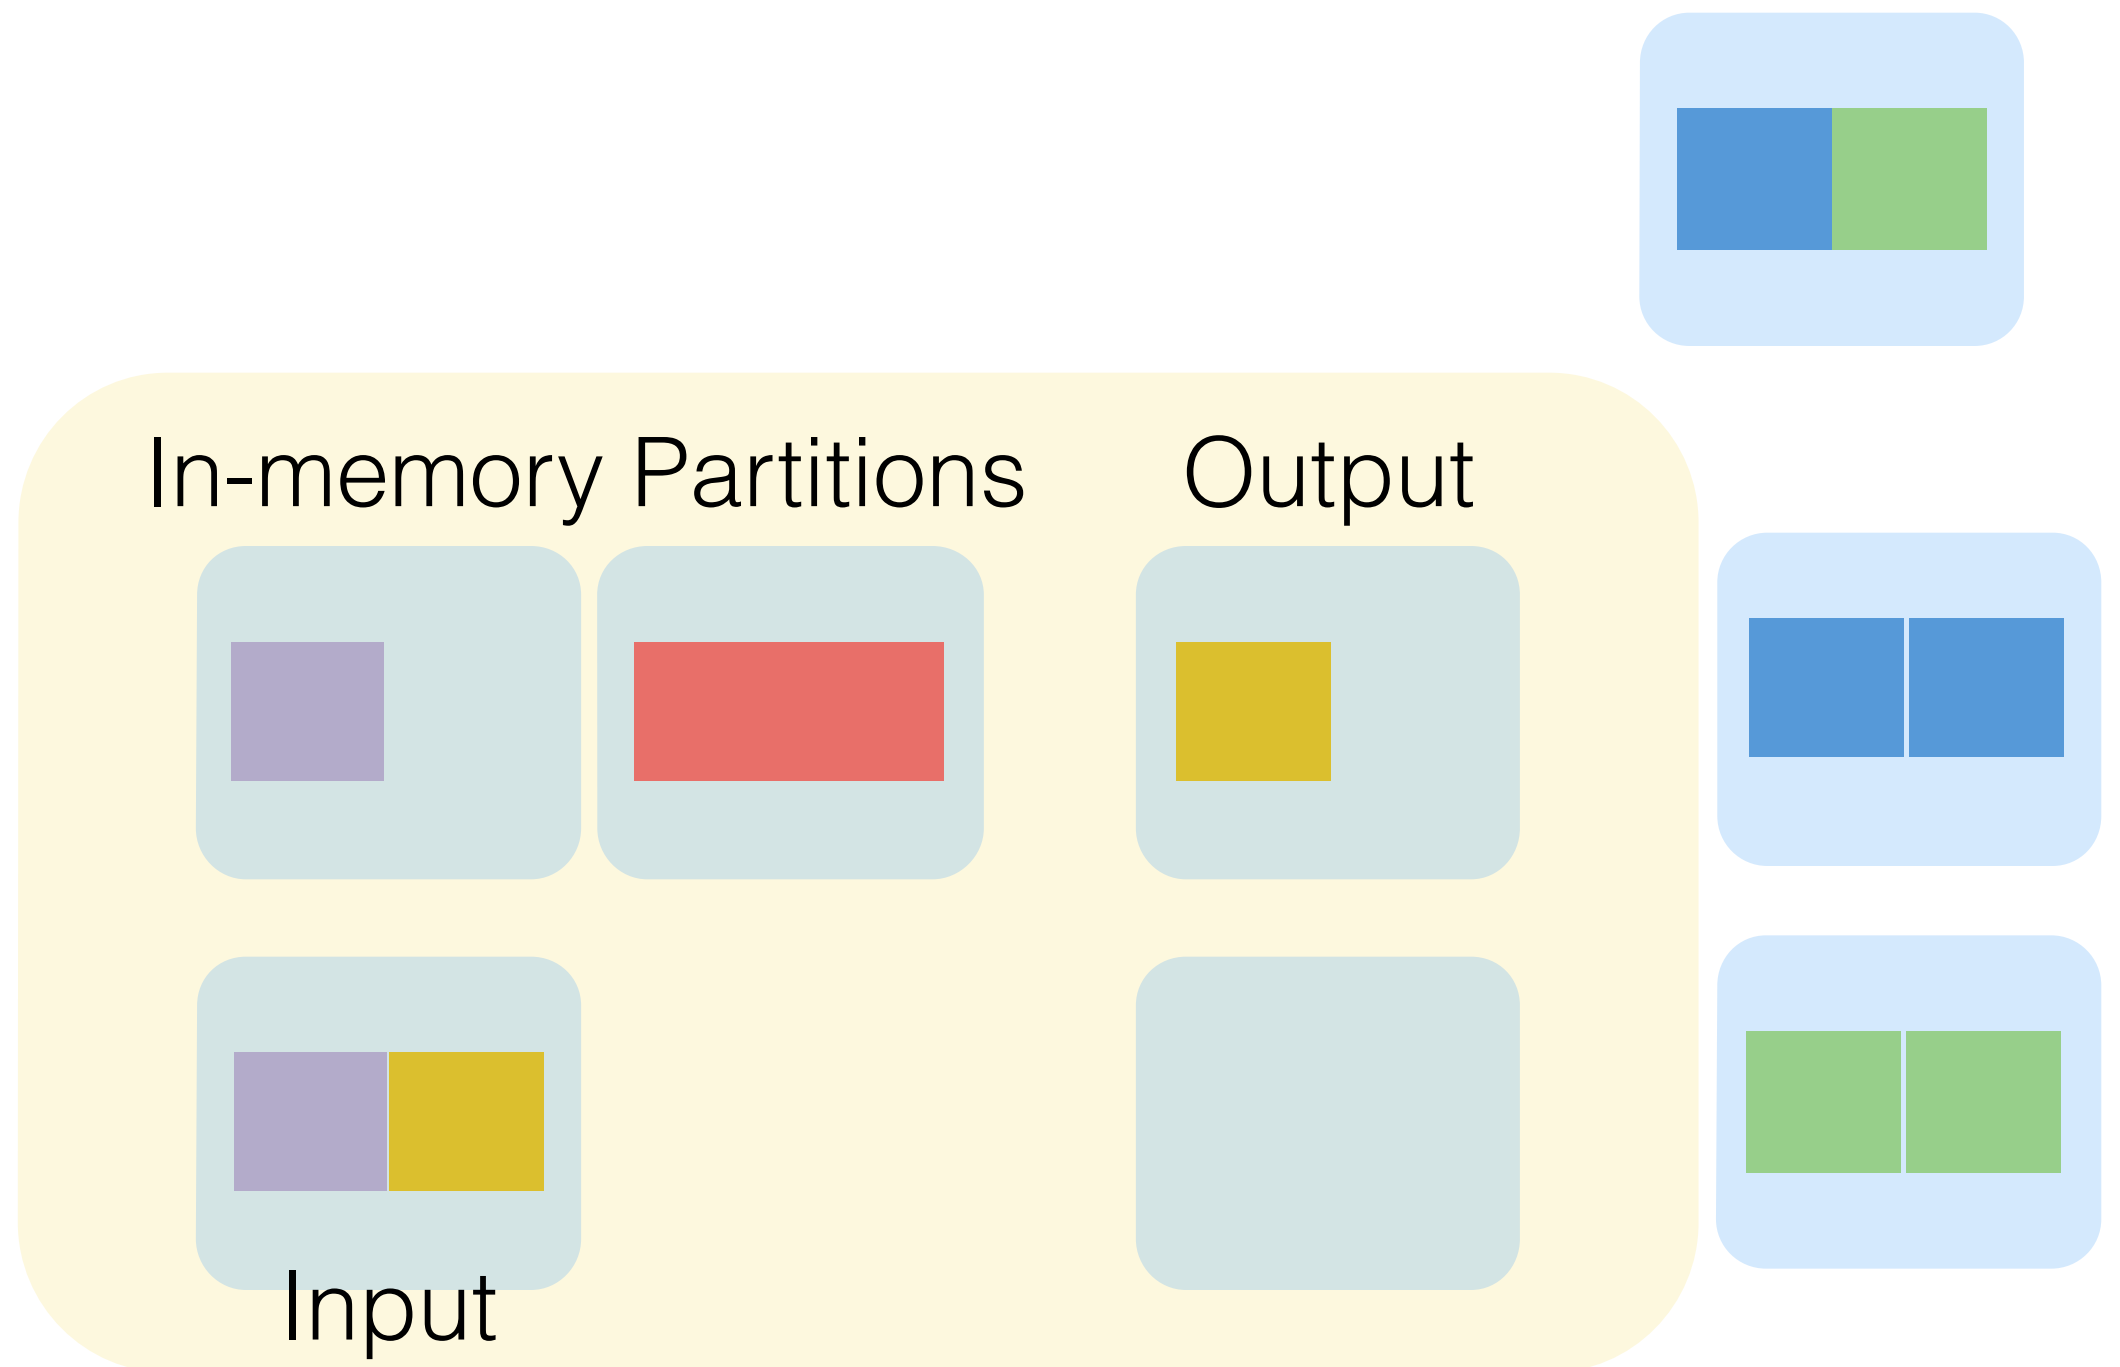
Solution: Keep some partitions in memory

Hybrid Hashing



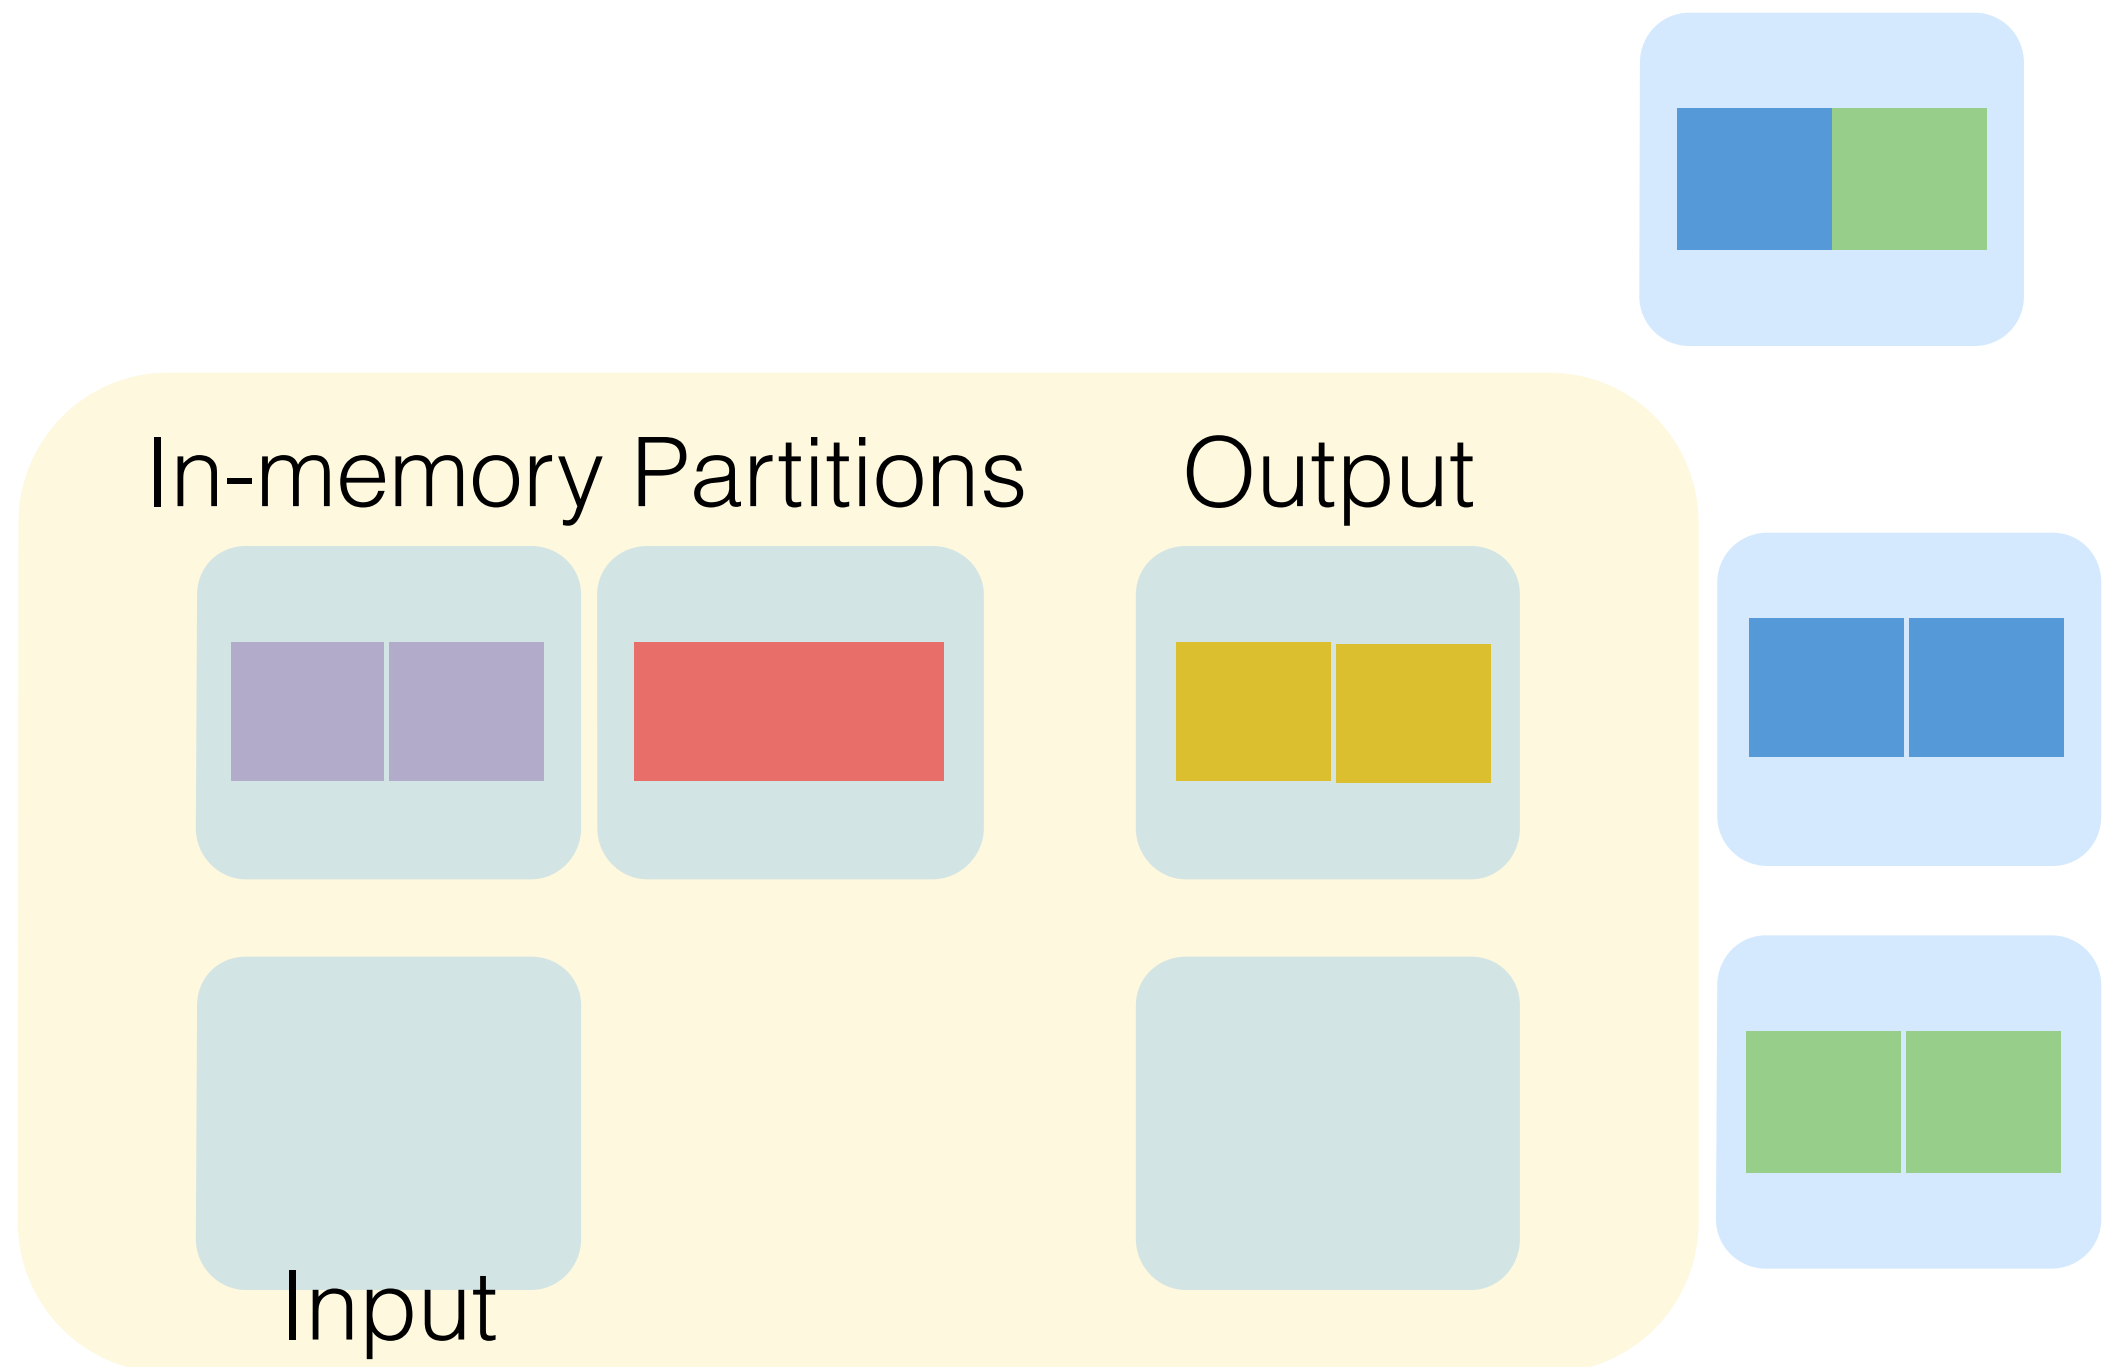
Solution: Keep some partitions in memory

Hybrid Hashing



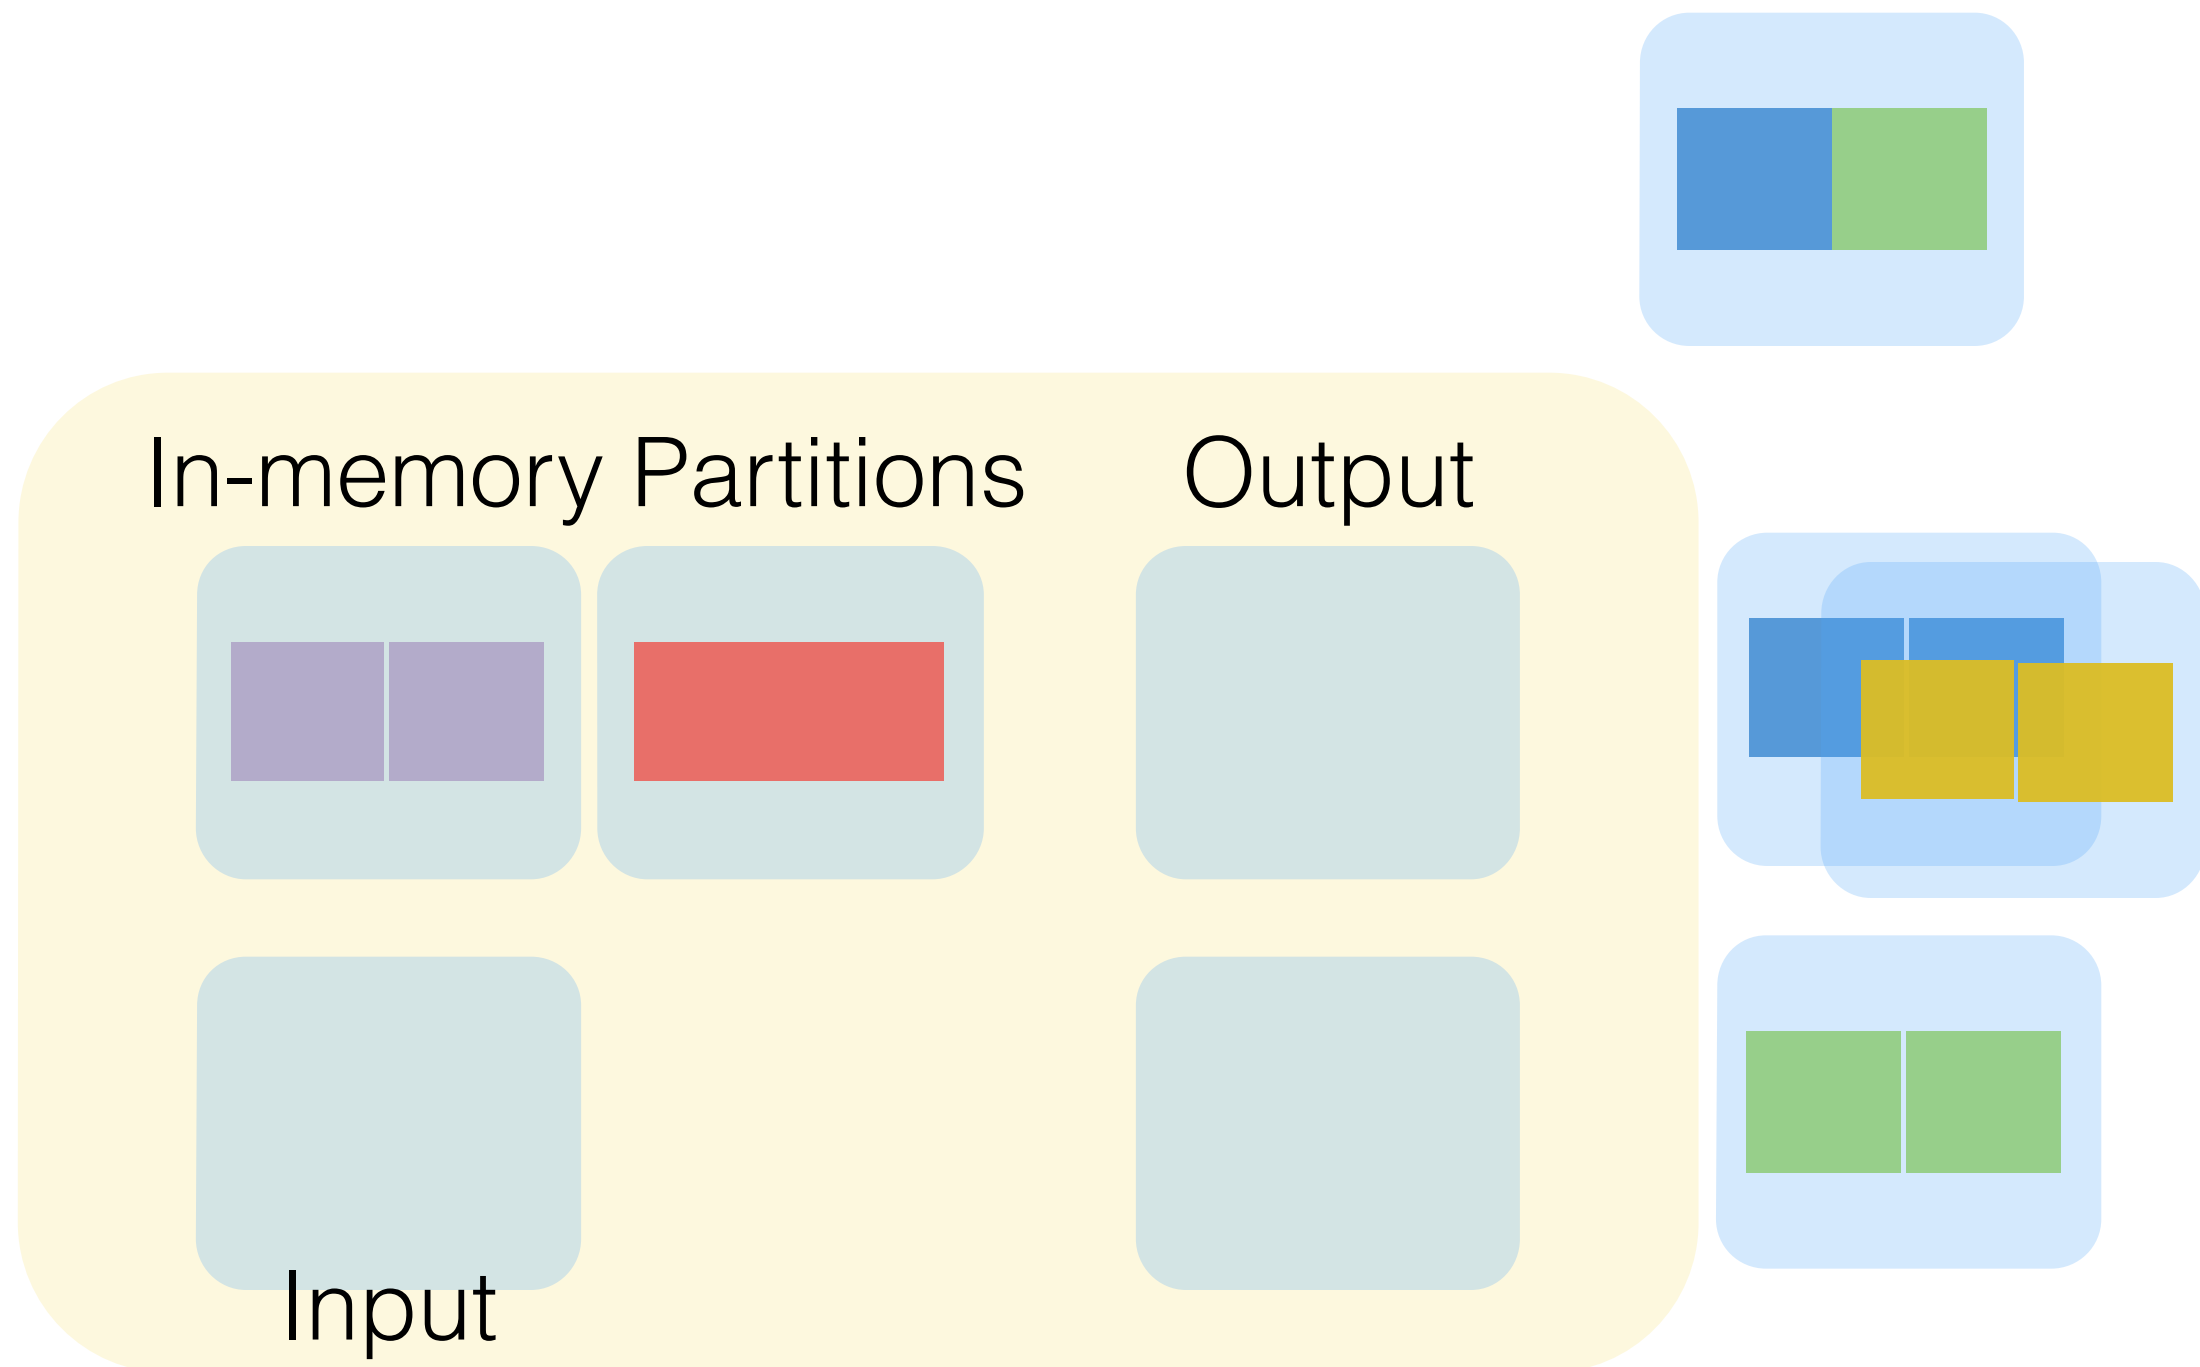
Solution: Keep some partitions in memory

Hybrid Hashing



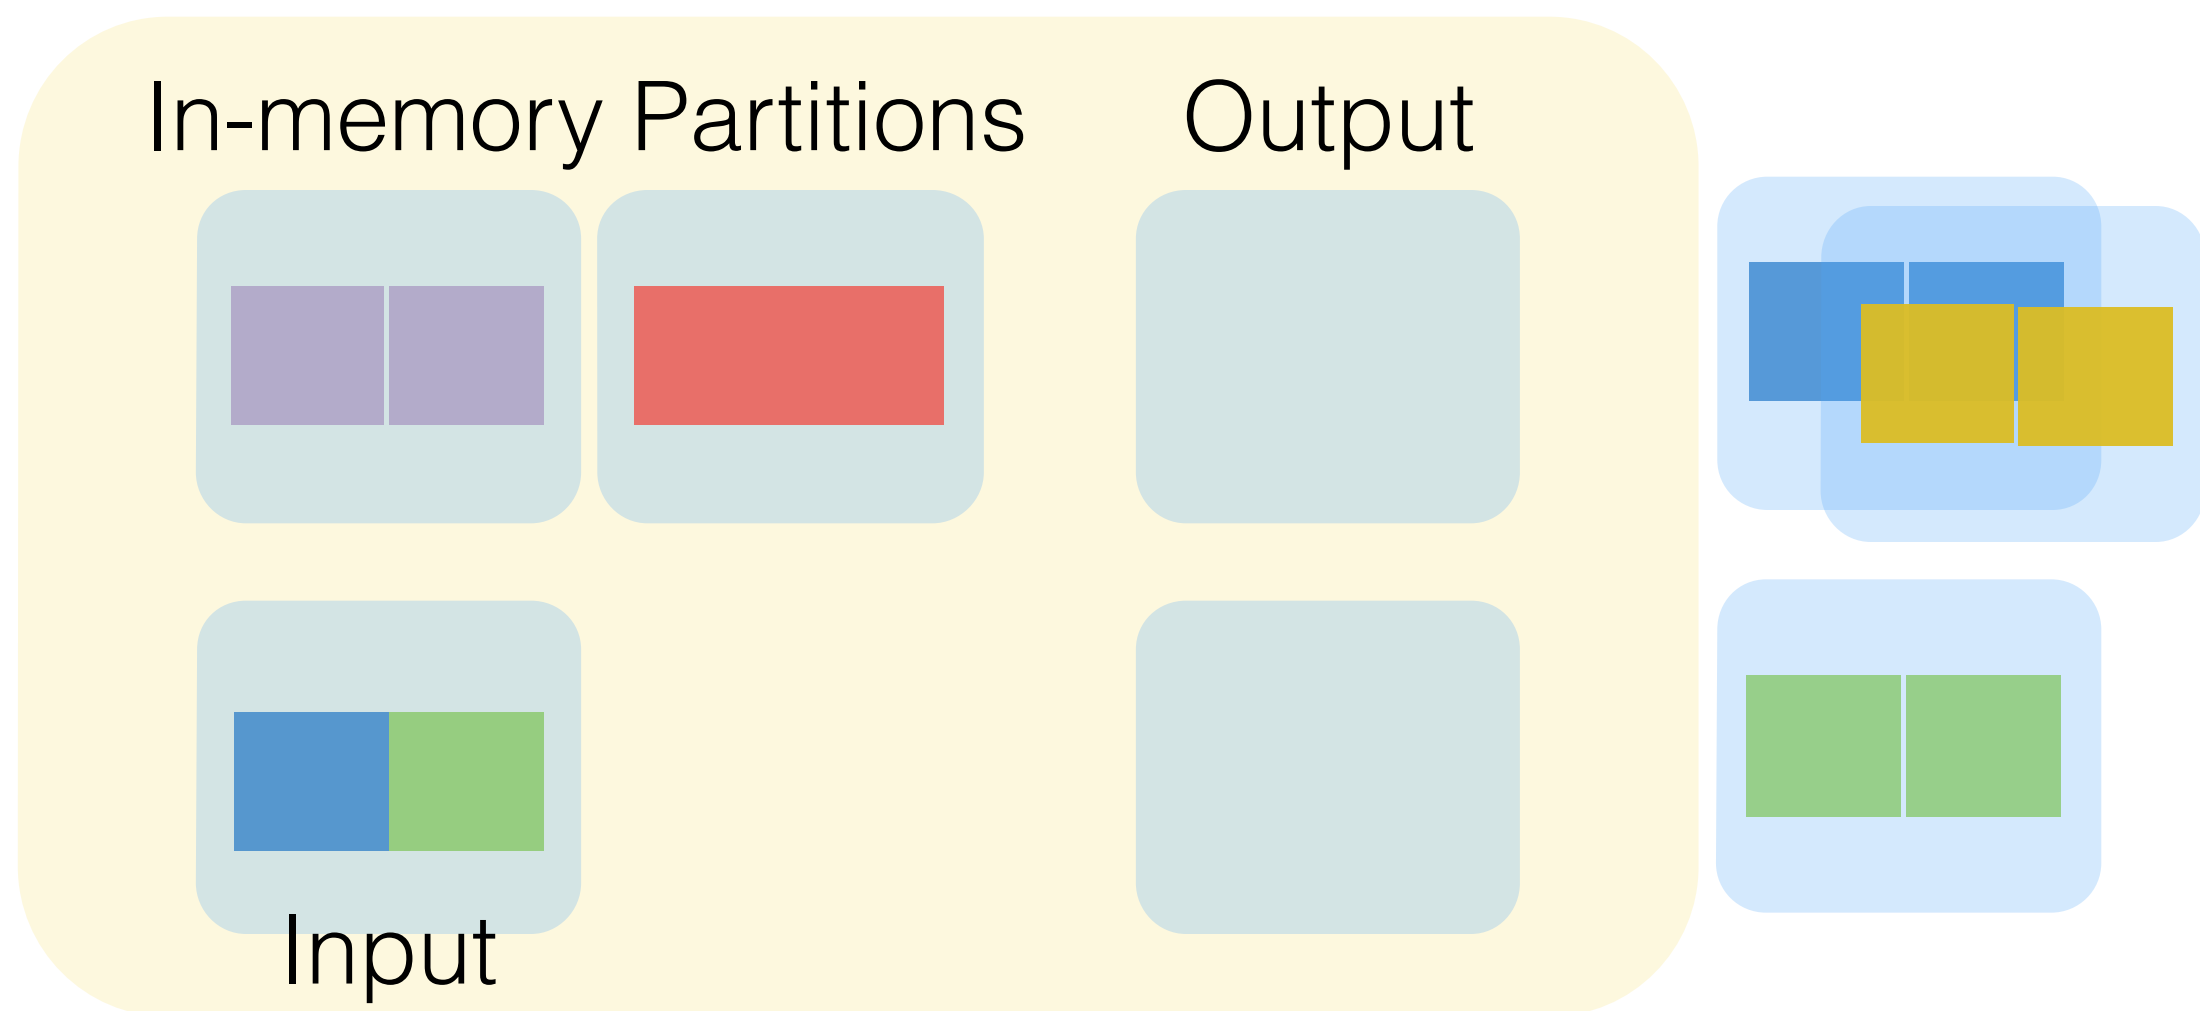
Solution: Keep some partitions in memory

Hybrid Hashing




Solution: Keep some partitions in memory

Hybrid Hashing



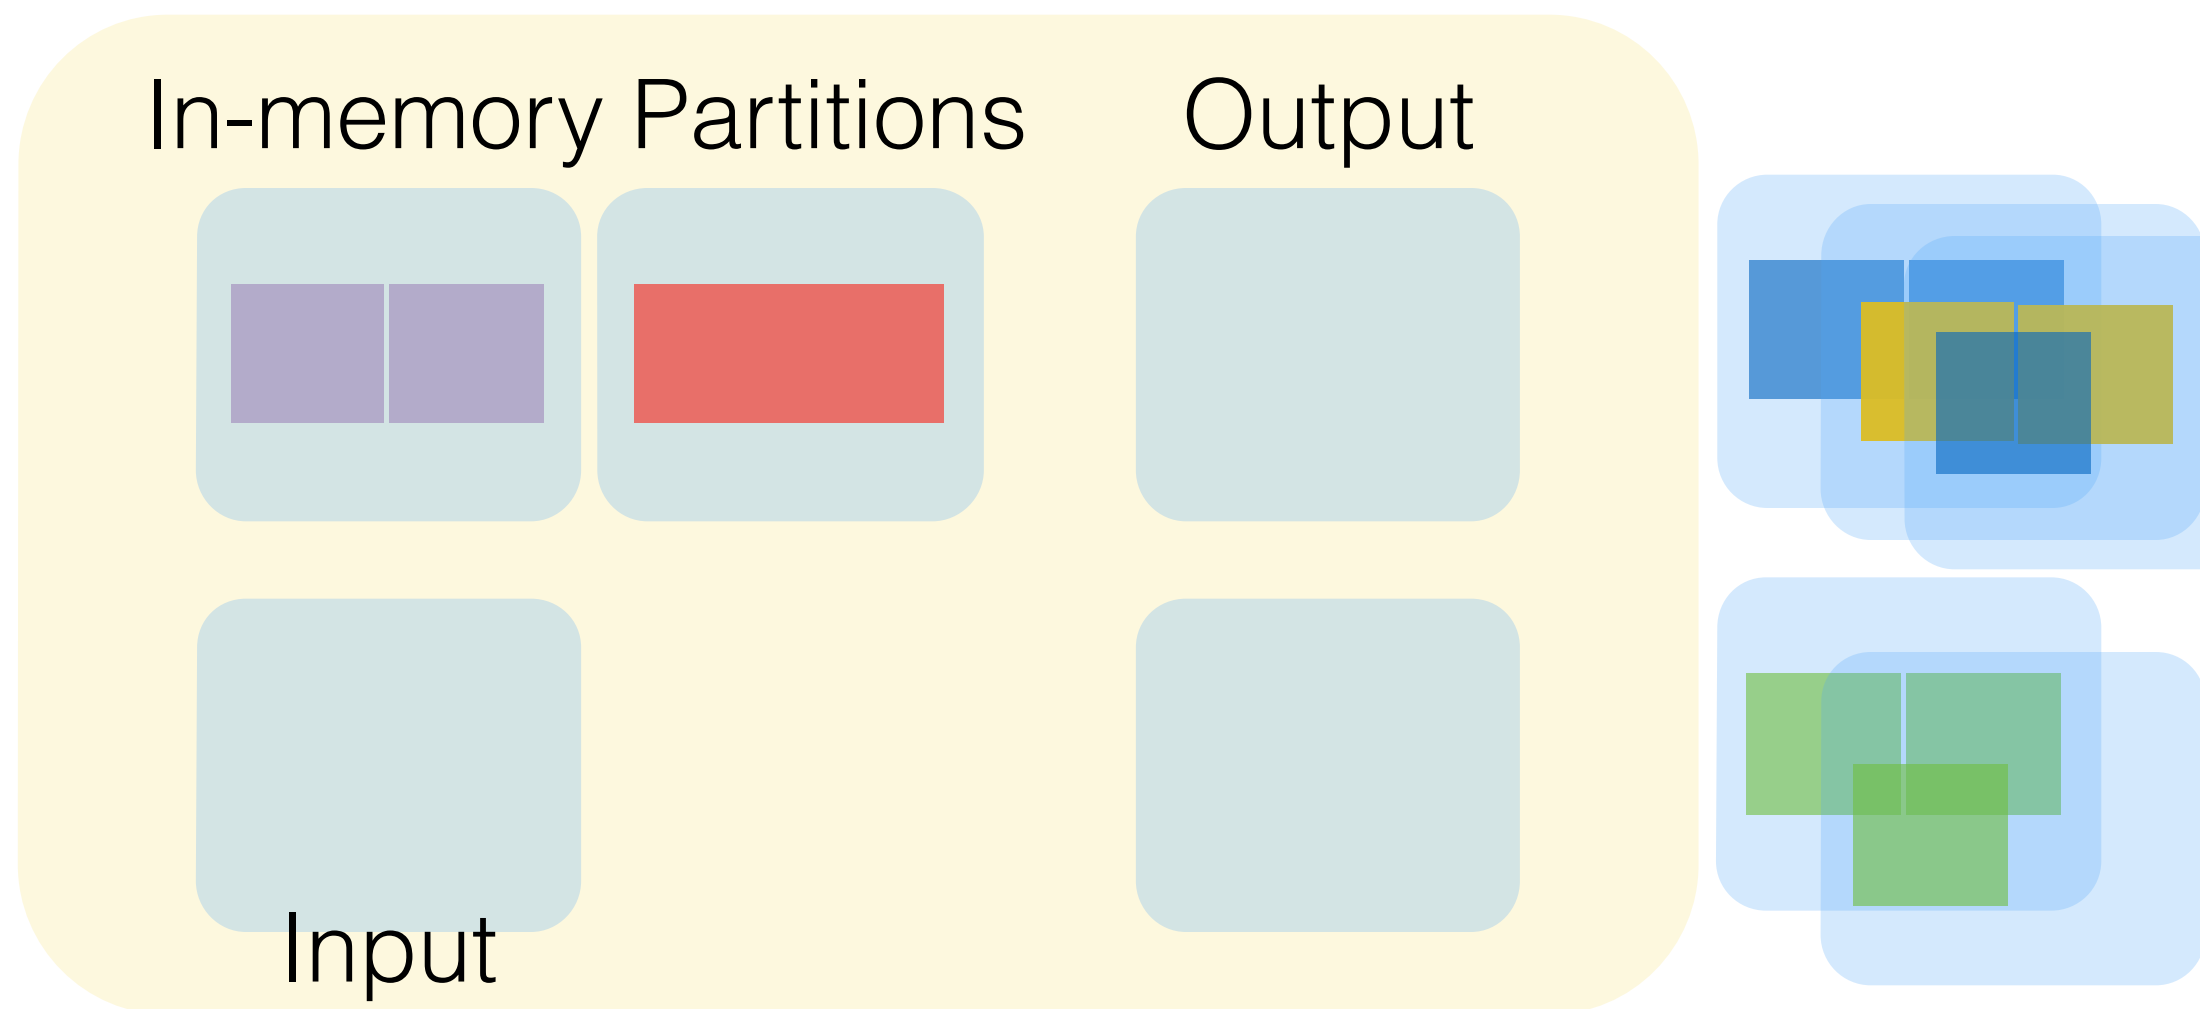
Solution: Keep some partitions in memory

Hybrid Hashing



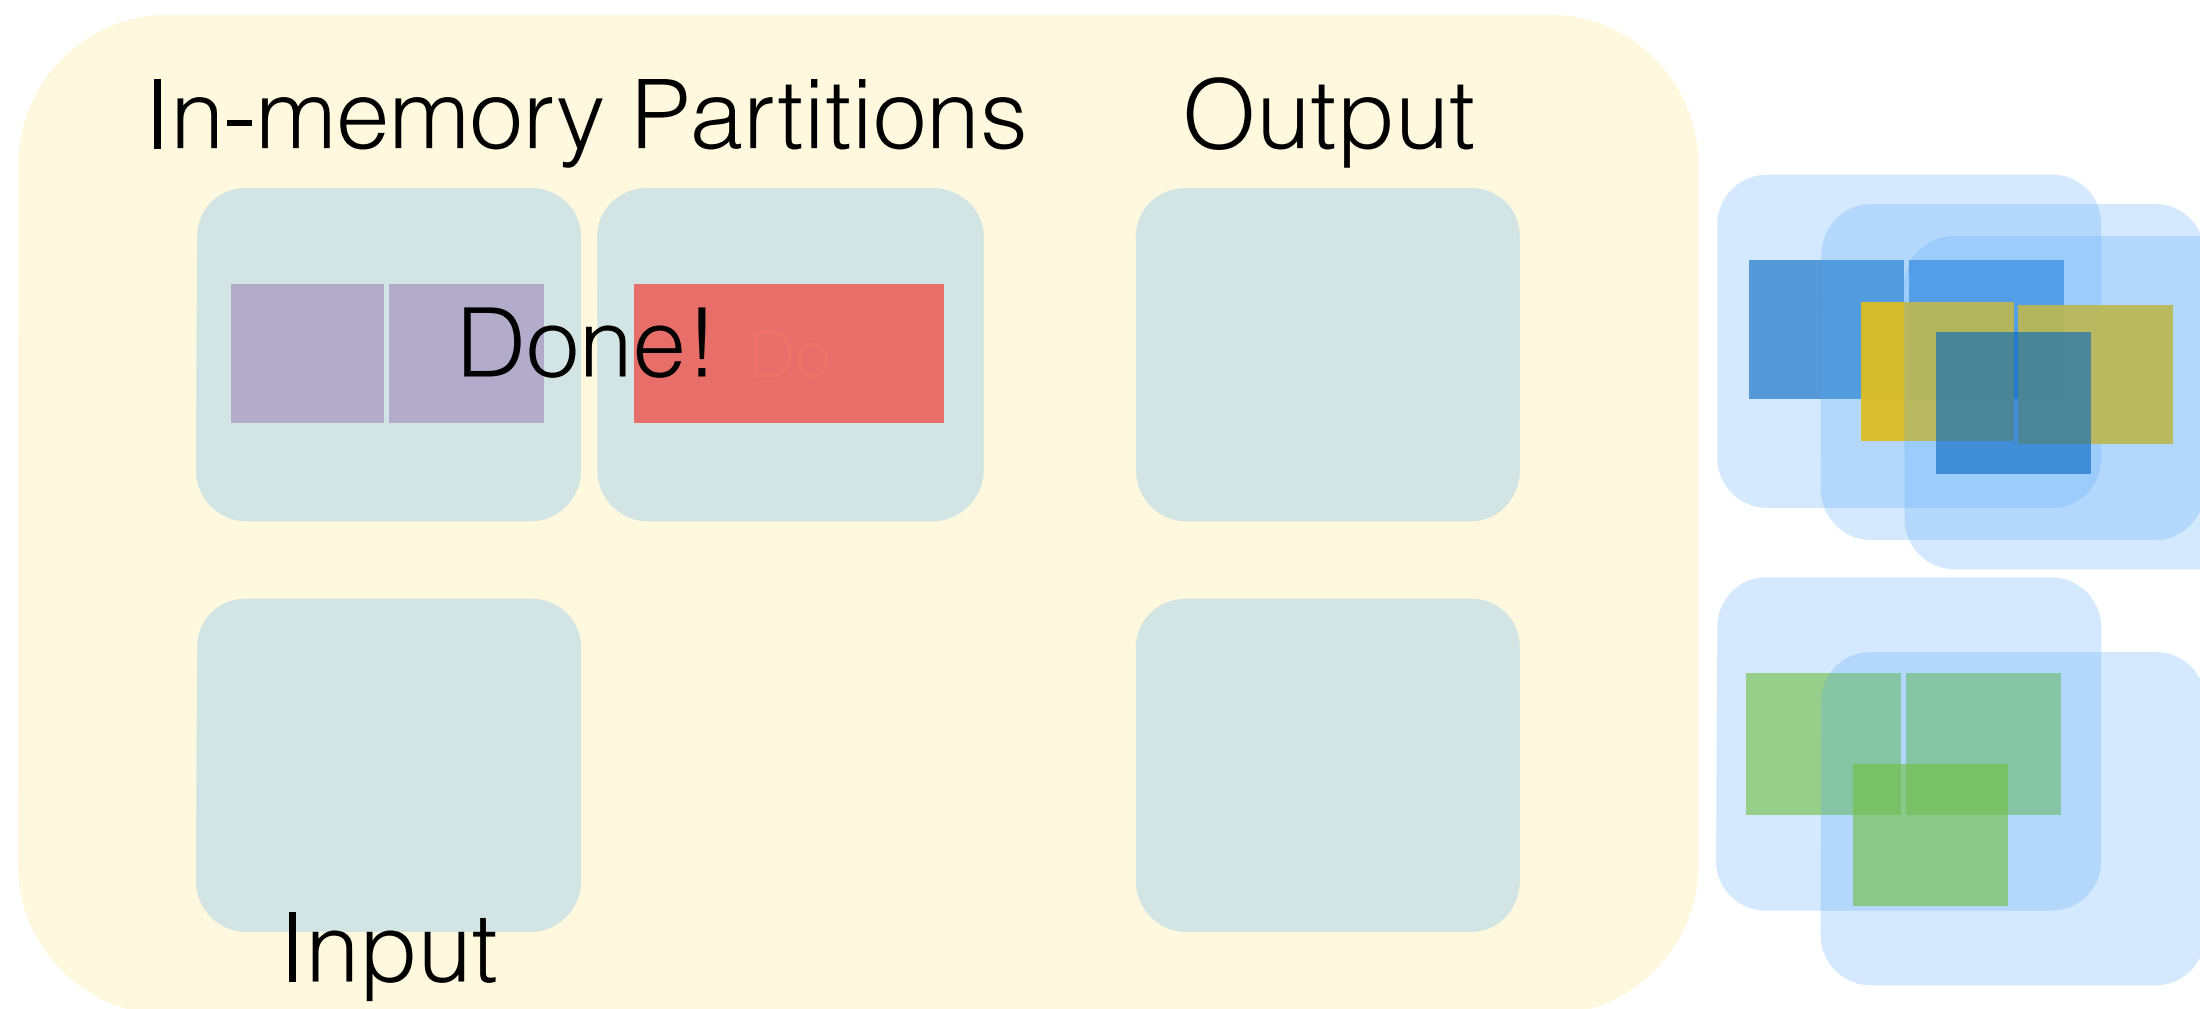
Solution: Keep some partitions in memory

Hybrid Hashing




Solution: Keep some partitions in memory

Hybrid Hashing



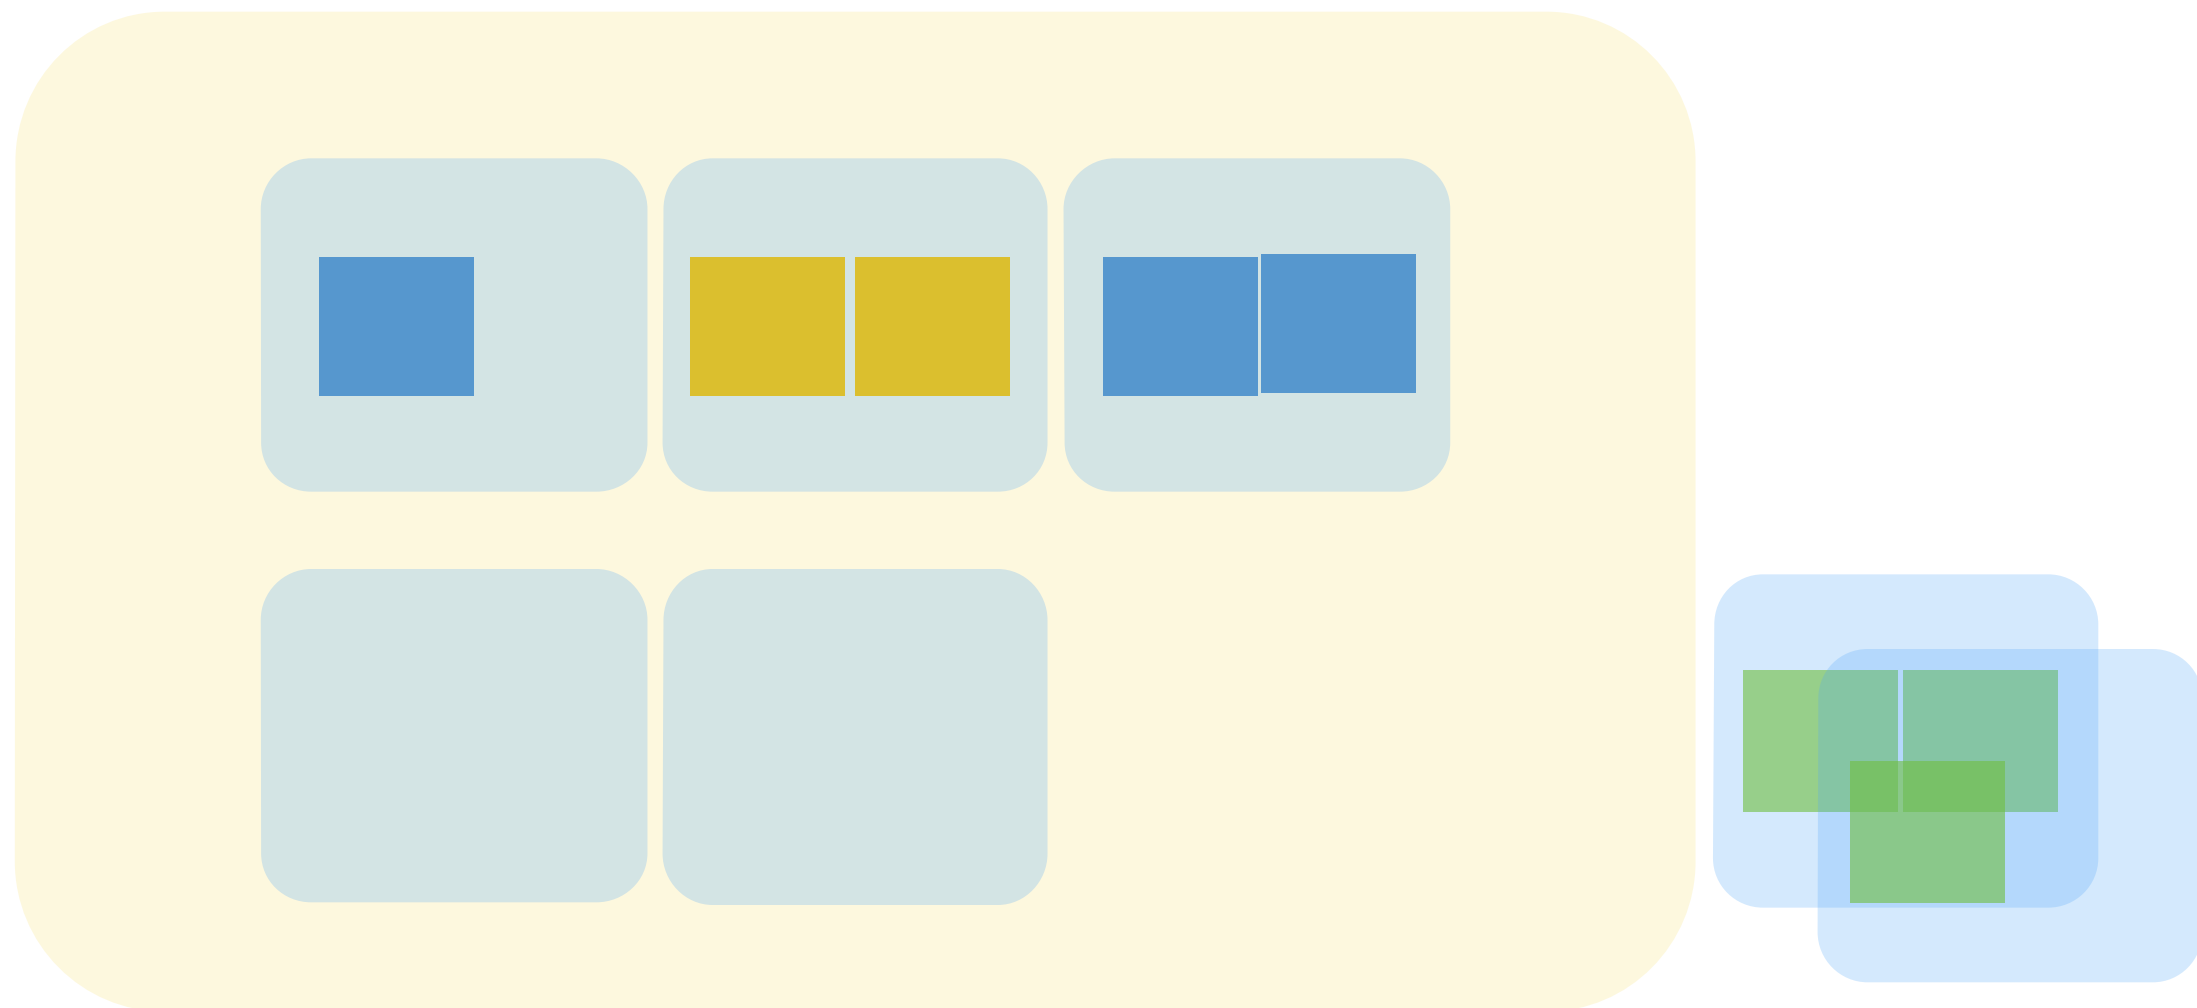
Save cost of reading those partitions from memory.

Hybrid Hashing



Read other partitions in memory.

Hybrid Hashing



Join Cheatsheet

Notation: $[S]$ == “# pages in S” ;

$|S|$ == “# tuples in S”

- Chunk nested loop join
 - Take **k pages** of S and match with each page of R.
 - Total Cost: $[S] + ([S] / k) * [R]$
- Sort merge join
 - Sort S and R **on join column**, then merge them!
 - Total Cost: $\sim 5([S] + [R])$
- Hash join
 - Partition S and R using same hash fn, then collect same partitions
 - Total Cost: $\sim 3([S] + [R])$
 - Assuming $\text{len}(\text{partition}) \leq B$ pages

When is a chunk-nested loops join the best?

When is a chunk-nested loops join the best?

- Not using an equality predicate
- Join is just a cross product

When is a sort-merge join the best?

When is a sort-merge join the best?

- Skewed input data
 - Could prompt recursive hashing
- Want sorted output or input already sorted

When is a hash-join the best?

When is a hash-join the best?

- One table is large, and the other small (can keep perform hybrid-hashing)

We have 12 pages of memory, and we want to join two tables [R] and [S] where [R] is 100 pages and [S] is 50 pages. [R] holds 100 tuples per page and [S] holds 50 tuples per page.

How many disk reads are needed to perform a Simple Nested Loops join?

We have 12 pages of memory, and we want to join two tables [R] and [S] where [R] is 100 pages and [S] is 50 pages. [R] holds 100 tuples per page and [S] holds 50 tuples per page.

How many disk reads are needed to perform a Simple Nested Loops join?

Using S as the outer relation yields the lowest I/O count.

$(\text{\#tuples in S}) * [R] * [S] + [S] = 50 * 100 * 50 + 50 = 250050 \text{ I/O's}$

We have 12 pages of memory, and we want to join two tables [R] and [S] where [R] is 100 pages and [S] is 50 pages. [R] holds 100 tuples per page and [S] holds 50 tuples per page.

How many disk reads are needed to perform Chunk Nested Loops Join?

We have 12 pages of memory, and we want to join two tables [R] and [S] where [R] is 100 pages and [S] is 50 pages. [R] holds 100 tuples per page and [S] holds 50 tuples per page.

How many disk reads are needed to perform Chunk Nested Loops Join?

$(\# \text{ of pages in smaller relation}) + ((\# \text{ of pages in smaller relation}) / (\# \text{ of pages in memory} - 2 \text{ for I/O})) * (\# \text{ of pages in larger relation})$

$$= 50 + (50/10) * (100) = 550$$

We have 12 pages of memory, and we want to join two tables [R] and [S] where [R] is 100 pages and [S] is 50 pages. [R] holds 100 tuples per page and [S] holds 50 tuples per page.

How about a Sort Merge Join? (Assume the join column is unique in both tables and that both tables are sorted in 2 passes)

We have 12 pages of memory, and we want to join two tables [R] and [S] where [R] is 100 pages and [S] is 50 pages. [R] holds 100 tuples per page and [S] holds 50 tuples per page.

How about a Sort Merge Join? (Assume the join column is unique in both tables and that both tables are sorted in 2 passes)

$$5[R] + 5[S] = 750 \text{ I/O's}$$

We have 12 pages of memory, and we want to join two tables [R] and [S] where [R] is 100 pages and [S] is 50 pages. [R] holds 100 tuples per page and [S] holds 50 tuples per page.

How about a Hash Join? (Assume no recursive partitioning)

We have 12 pages of memory, and we want to join two tables [R] and [S] where [R] is 100 pages and [S] is 50 pages. [R] holds 100 tuples per page and [S] holds 50 tuples per page.

How about a Hash Join? (Assume no recursive partitioning)

Partitioning Phase: $2([R] + [S])$

Matching Phase: $[R] + [S]$

Total = $3([R] + [S]) = 3 * 150 = 450$ I/O's

Midterm Practice

S with 11,000 pages, 2000 records per page

R with 500 pages, 500 records per page

Memory: 100 pages

How many I/O's are required to perform a simple, “unoptimized” sort-merge join?

Midterm Practice

S with 11,000 pages, 2000 records per page

R with 500 pages, 500 records per page

Memory: 100 pages

How many I/O's are required to perform a simple,
“unoptimized” sort-merge join?

cost to sort S + cost to sort R + cost to merge

Midterm Practice

S with 11,000 pages, 2000 records per page

R with 500 pages, 500 records per page

Memory: 100 pages

How many I/O's are required to perform a simple,
“unoptimized” sort-merge join?

cost to sort S + cost to sort R + cost to merge

$$2 \cdot 3 \cdot 11,000 + 2 \cdot 2 \cdot 500 + 11,000 + 500 = 79,500 \text{ I/Os}$$