

# CS186 Discussion #9

(More Concurrency)

# Midterm Review Session

Saturday, 11/7, 1-3pm in 100 GPB

# Lock Granularity


sid	points	grade
Bob	43	C
Joe	99	C
Alice	87	C
Suzy	50	C
Ted	73	C
Tim	12	C

T1: UPDATE students SET grade='A' WHERE points >= 70;

T2: UPDATE students SET grade='F' WHERE points < 70;

# Lock Granularity

T1(X)



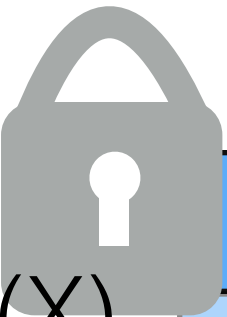
sid	points	grade
Bob	43	C
Joe	99	C
Alice	87	C
Suzy	50	C
Ted	73	C
Tim	12	C

T1: UPDATE students SET grade='A' WHERE points >= 70;

T2: UPDATE students SET grade='F' WHERE points < 70;

# Lock Granularity

T1(X)



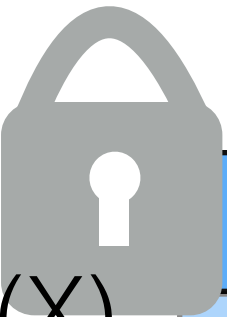
sid	points	grade
Bob	43	C
Joe	99	A
Alice	87	A
Suzy	50	C
Ted	73	A
Tim	12	C

T1: UPDATE students SET grade='A' WHERE points >= 70;

T2: UPDATE students SET grade='F' WHERE points < 70;

# Lock Granularity

T2(X)



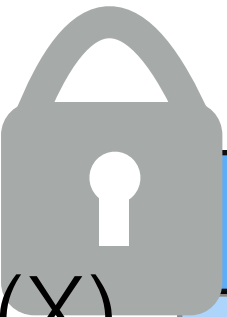
sid	points	grade
Bob	43	C
Joe	99	A
Alice	87	A
Suzy	50	C
Ted	73	A
Tim	12	C

T1: UPDATE students SET grade='A' WHERE points >= 70;

T2: UPDATE students SET grade='F' WHERE points < 70;

# Lock Granularity

T2(X)




sid	points	grade
Bob	43	F
Joe	99	A
Alice	87	A
Suzy	50	F
Ted	73	A
Tim	12	F

T1: UPDATE students SET grade='A' WHERE points >= 70;

T2: UPDATE students SET grade='F' WHERE points < 70;

# Lock Hierarchy

- Each Xact starts at root of hierarchy
- Gets locks in top-down order
- Releases locks in bottom-up order





# Locks

- IS: intent to get S lock(s) at finer granularity
- IX: intent to get X lock(s) at finer granularity
- SIX: shared lock, with intent to get X lock(s) at finer granularity


# Lock Compatibility Matrix

	IS	IX	SIX	S	X
IS					
IX					
SIX					
S				✓	-
X				-	-

# Lock Compatibility Matrix

	IS	IX	SIX	S	X
IS	✓	✓	✓	✓	—
IX	✓	✓	—	—	—
SIX	✓	—	—	—	—
S	✓	—	—	✓	—
X	—	—	—	—	—

# Lock Granularity



sid	points	grade	
Bob	43	F	T2(X)
Joe	99	A	T1(X)
Alice	87	A	T1(X)
Suzy	50	F	T2(X)
Ted	73	A	T1(X)
Tim	12	F	T2(X)

# Worksheet - Lock Granularity

Suppose a transaction, T1, wants to scan a table R and update a few of its tuples. What kind of locks should T1 have on R, its pages, and the tuples that are updated?

Suppose a transaction, T1, wants to scan a table R and update a few of its tuples. What kind of locks should T1 have on R, its pages, and the tuples that are updated?

- SIX lock on R
- SIX or IX lock on pages
- X lock on updated tuples
- S lock on other tuples

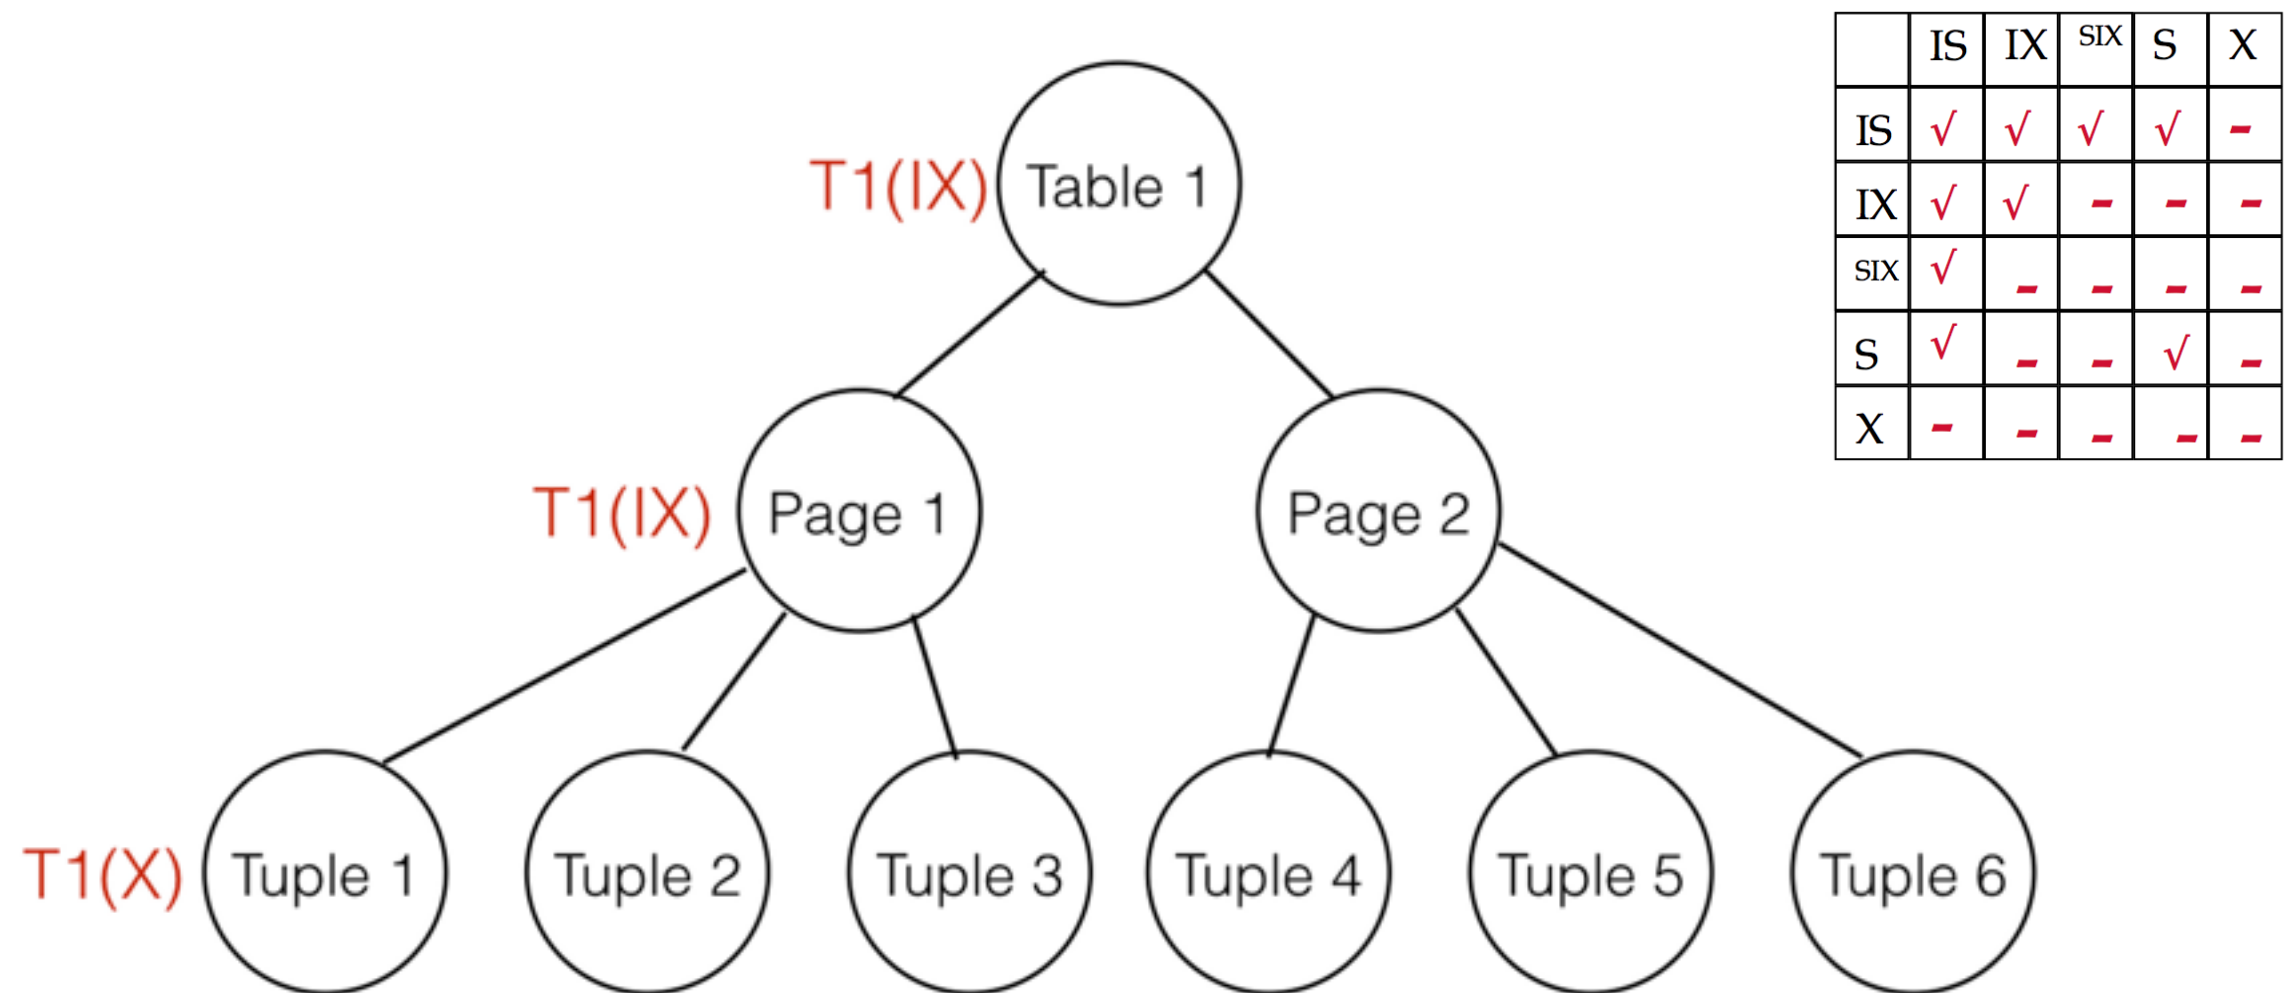
Is an S lock compatible with an IX lock?  
Explain why or why not. Make your  
description as simple as possible.




Is an S lock compatible with an IX lock?  
Explain why or why not. Make your  
description as simple as possible.

- Incompatible:
  - T1 has S lock on Students table to calculate average grade
  - T2 wants IX lock to change some grades

# Which locks could be granted to transaction T2 for tuple 2?




# Which locks could be granted to transaction T2 for tuple 2?



X, S


Can get IX, IS, but does not make sense at tuple-level.

# What locks could be granted to T2 for page 1?



	IS	IX	SIX	S	X
IS	✓	✓	✓	✓	-
IX	✓	✓	-	-	-
SIX	✓	-	-	-	-
S	✓	-	-	✓	-
X	-	-	-	-	-

# What locks could be granted to T2 for page 1?




	IS	IX	SIX	S	X
IS	✓	✓	✓	✓	-
IX	✓	✓	-	-	-
SIX	✓	-	-	-	-
S	✓	-	-	✓	-
X	-	-	-	-	-

S, IS


# Multiversion Concurrency Control

- Alternative to 2PL
  - Less waiting, but more aborts
- Timestamp Ordered MVCC:
  - Each transaction gets timestamp upon entry
  - Keep timeline of read timestamps and versions

# TO-MVCC




# TO-MVCC



Reads: Read version with biggest timestamp smaller than current timestamp




# TO-MVCC



Reads: Read version with biggest timestamp smaller than current timestamp


# TO-MVCC



Reads: Read version with biggest timestamp smaller than current timestamp

$R(X)$  @ 11 will read  $V_2$


# TO-MVCC



Writes: Find interval from given timestamp  $X$  to smallest timestamp bigger than  $X$

If there a read is in the interval, **abort!**

# TO-MVCC




Writes: Find interval from given timestamp  $X$  to smallest timestamp bigger than  $X$

If there a read is in the interval, **abort!**

$W(X)$  @ 11

# TO-MVCC

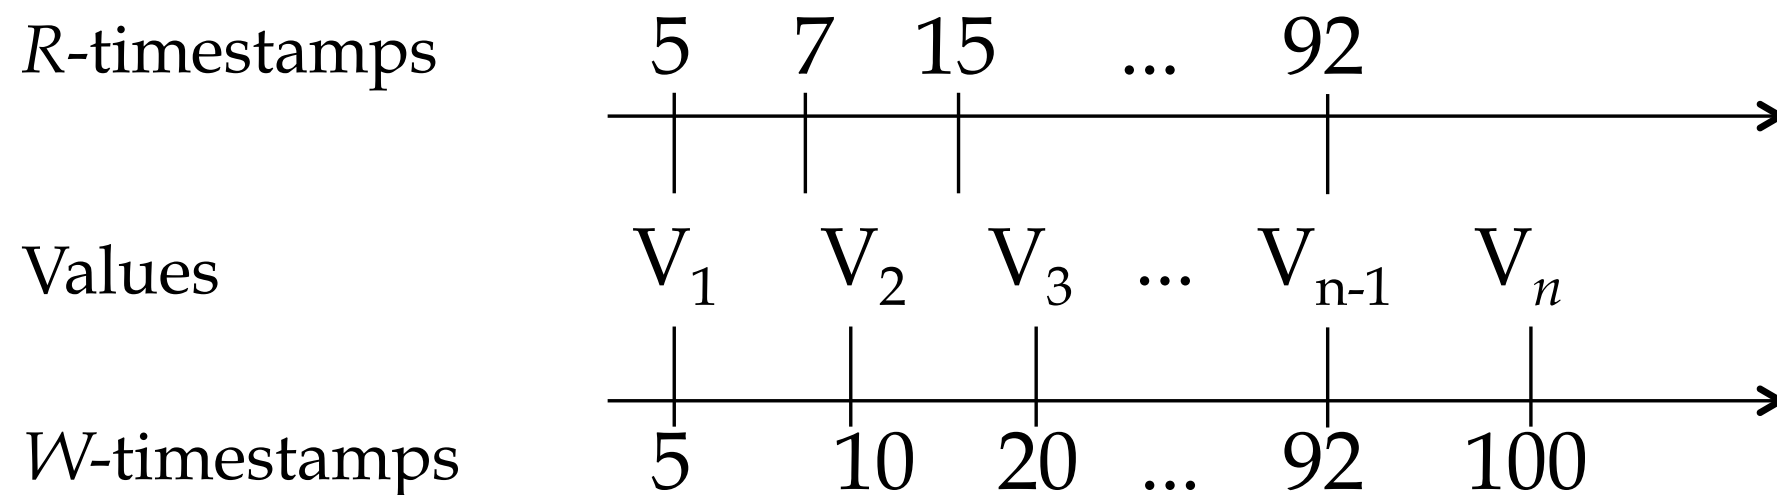


Writes: Find interval from given timestamp  $X$  to smallest timestamp bigger than  $X$

If there a read is in the interval, **abort!**

$W(X)$  @ 11

# TO-MVCC




Writes: Find interval from given timestamp  $X$  to smallest timestamp bigger than  $X$

If there a read is in the interval, **abort!**

$W(X)$  @ 8

# TO-MVCC



Writes: Find interval from given timestamp  $X$  to smallest timestamp bigger than  $X$


If there a read is in the interval, **abort!**

$W(X) @ 8$


# Worksheet - MVCC




R(X)@17, W(X)@18, W(X)@14, W(X)@12, R(X)@20, R(X)@19,  
R(X)@23, W(X)@26, W(X)@24




$R(X)@17$ ,  $W(X)@18$ ,  $W(X)@14$ ,  $W(X)@12$ ,  $R(X)@20$ ,  $R(X)@19$ ,  
 $R(X)@23$ ,  $W(X)@26$ ,  $W(X)@24$




R(X)@17, W(X)@18, W(X)@14, W(X)@12, R(X)@20, R(X)@19,  
R(X)@23, W(X)@26, W(X)@24




R(X)@17, W(X)@18, ~~W(X)@14~~, W(X)@12, R(X)@20, R(X)@19,  
R(X)@23, W(X)@26, W(X)@24



R(X)@17, W(X)@18, ~~W(X)@14~~, W(X)@12, R(X)@20, R(X)@19,  
R(X)@23, W(X)@26, W(X)@24



$R(X)@17$ ,  $W(X)@18$ ,  ~~$W(X)@14$~~ ,  $W(X)@12$ ,  $R(X)@20$ ,  $R(X)@19$ ,  
 $R(X)@23$ ,  $W(X)@26$ ,  $W(X)@24$



# Distributed Concurrency Control

- Parallel database or distributed database
- Each transaction has one coordinator node
- Every node handles its own CC

# Two Phase Commit

- Decides whether to commit/abort
- Ensures atomicity and durability
  - Writes happen in all replicas or none



# Two Phase Commit

- Phase 1:
  - Coordinator tells participants to prepare
  - Participants respond with “yes” commit or “no” commit
- Phase 2:
  - If **all** participants say “yes”, issue GLOBAL COMMIT
  - Else, abort

# 2PC & Logging

- Phase 1:
  - Each slave writes an entry to its undo and redo logs
- Phase 2:
  - If abort, the slaves use their undo log to rollback its effects