


Review Session

EXTERNAL SORTING

General External Merge Sort

- More than 3 buffer pages. How can we utilize them?
- To sort a file with N pages using B buffer pages:
 - Pass 0: use B buffer pages. Produce $\lceil N / B \rceil$ sorted runs of B pages each.
 - Pass 1, 2, ..., etc.: merge $B-1$ runs.



Cost of External Merge Sort


- Number of passes: $1 + \lceil \log_{B-1} \lceil N / B \rceil \rceil$
- Cost = $2N * (\text{\# of passes})$
- E.g., with 5 buffer pages, to sort 108 page file:
 - Pass 0: $\lceil 108 / 5 \rceil = 22$ sorted runs of 5 pages each (last run is only 3 pages)
 - Pass 1: $\lceil 22 / 4 \rceil = 6$ sorted runs of 20 pages each (last run is only 8 pages)
 - Pass 2: 2 sorted runs, 80 pages and 28 pages
 - Pass 3: 1 run => Sorted file of 108 pages

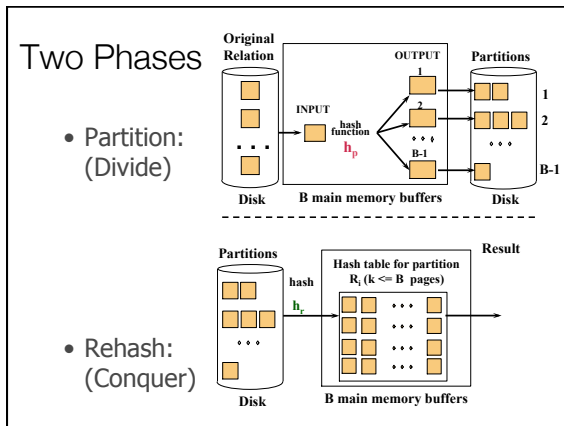
Formula check: $1 + \lceil \log_4 22 \rceil = 1 + 3 \rightarrow 4 \text{ passes} \checkmark$

EXTERNAL HASHING

Two Phases

- Partition: (Divide)






TOURNAMENT SORT

Tournament Sort (Heapsort)



First, load a heap with B-2 pages of records




This is a priority heap, so it is in sorted order (hence called Heapsort)

Tournament Sort (2)



```
while (records left on input) {
    m = H1.removeMin(); // get smallest value
    Output(m)           // put m in output buffer;
    if (H1 NOT empty)
        r = InputRecord()
        if (r < m) H2.insert(r);
        else      H1.insert(r);
    else
        H1 = H2; H2.reset();
        start new output run;
}
```



TRICKY QUESTION



Optimizing Sort-Merge Join

Sailors
(name = Bob, sid = 1)
(name = Jill, sid = 2)
(name = Yue, sid = 4)
(name = Sue, sid = 8)
(name = Jack, sid = 18)
(name = Cat, sid = 22)
...
(name = Sam, sid = 3)
(name = Sue, sid = 7)
(name = Joe, sid = 12)
...

Reserves
(sid = 1, bid = 4)
(sid = 1, bid = 7)
(sid = 4, bid = 3)
(sid = 8, bid = 1)
(sid = 8, bid = 13)
(sid = 12, bid = 1)
...
(sid = 3, bid = 6)
(sid = 8, bid = 15)
...

Key idea:
Internal Sort on both.
Perform merge on all runs!

Steps:

1. Internal sort S and R. (Pass 0)
2. Merge all runs.

I/Os:

~3([S] + [R])
Pass 0: 2([S]+[R])
Merging: [S]+[R]

Using General Merge Sort, when can you do Optimized Sort Merge join? Express with an inequality.

STEP 1:

- # runs for R – ceil([R]/B)
- # run for S – ceil([S]/B)

STEP 2:

- We need to fit EACH run into buffer

Optimizing Sort-Merge Join

Sailors
(name = Bob, sid = 1)
(name = Jill, sid = 2)
(name = Yue, sid = 4)
(name = Sue, sid = 8)
(name = Jack, sid = 18)
(name = Cat, sid = 22)
...
(name = Sam, sid = 3)
(name = Sue, sid = 7)
(name = Joe, sid = 12)
...

Reserves
(sid = 1, bid = 4)
(sid = 1, bid = 7)
(sid = 4, bid = 3)
(sid = 8, bid = 1)
(sid = 8, bid = 13)
(sid = 12, bid = 1)
...
(sid = 3, bid = 6)
(sid = 8, bid = 15)
...

Key idea:
Internal Sort on both.
Perform merge on all runs!

Steps:

1. Internal sort S and R. (Pass 0)
2. Merge all runs.

I/Os:

~3([S] + [R])
Pass 0: 2([S]+[R])
Merging: [S]+[R]

$$\text{ceil}([S]/B) + \text{ceil}([R]/B) \leq B - 1$$

Using tournament sort, when can you do Optimized Sort Merge join?
Sort Merge join? Express with an inequality.

STEP 1:

- # runs for R – $\text{ceil}([R]/2(B-2))$
- # run for S – $\text{ceil}([S]/2(B-2))$

STEP 2:

- We need to fit EACH run into buffer

Optimizing Sort-Merge Join

Sailors

Reserve = Bob, bid = 31
Reserve = Sam, bid = 21
Reserve = Tom, bid = 41
Reserve = Sue, bid = 61
Reserve = Jack, bid = 181
Reserve = Carl, bid = 221

Boats

bid = 1, bid = 41
bid = 3, bid = 71
bid = 4, bid = 91
bid = 8, bid = 131
bid = 12, bid = 171
bid = 3, bid = 41
bid = 8, bid = 151

Key Idea:
Internal Sort on both.
Perform merge on all runs!

Steps:
1. Internal sort S and R.
(Pass 0)
2. Merge all runs.

I/Os:
 $\sim 3([S] + [R])$
Pass 0: $2([S] + [R])$
Merging: $[S] + [R]$

$\text{ceil}([S]/2(B-2)) + \text{ceil}([R]/2(B-2)) \leq B - 1$

NOW, let's assume [R] is equal to [S]. Express the same inequality with [R] and [B].

STEP 1:

- # runs for R – $\text{ceil}([R]/2(B-2))$
- # runs for S – $\text{ceil}([R]/2(B-2))$

STEP 2:

- We need to fit EACH run into buffer

Optimizing Sort-Merge Join

Sailors

Reserve = Bob, bid = 31
Reserve = Sam, bid = 21
Reserve = Tom, bid = 41
Reserve = Sue, bid = 61
Reserve = Jack, bid = 181
Reserve = Carl, bid = 221

Boats

bid = 1, bid = 41
bid = 3, bid = 71
bid = 4, bid = 91
bid = 8, bid = 131
bid = 12, bid = 171
bid = 3, bid = 41
bid = 8, bid = 151

Key Idea:
Internal Sort on both.
Perform merge on all runs!

Steps:
1. Internal sort S and R.
(Pass 0)
2. Merge all runs.

I/Os:
 $\sim 3([S] + [R])$
Pass 0: $2([S] + [R])$
Merging: $[S] + [R]$

$\text{ceil}([R]/2(B-2)) + \text{ceil}([R]/2(B-2)) \leq B - 1$

Tournament Sort +
Optimized Sort Merge

$\text{ceil}([R]/(B-2)) < B$

$[R] < (B)(B-2)$

As B becomes very large, we can approximate this:
 $\text{SQRT}([R]) < B$

SQL

Vitamin Question

```

SELECT sname
FROM
  (SELECT sid
   FROM Reserves
   EXCEPT
   (SELECT sid
    FROM
      (SELECT Reserves.sid, PinkBoats.bid
       FROM Reserves,
            (SELECT bid FROM Boats WHERE color='pink') PinkBoats
       EXCEPT
        (SELECT sid, bid FROM Reserves)
      )
   ) R, Sailors S
 WHERE R.sid = S.sid;
  
```

Vitamin Question Step 1

```

### STEP 1
SELECT sname
FROM
  (SELECT sid
   FROM Reserves
   EXCEPT
   (SELECT sid
    FROM
      (SELECT Reserves.sid, PinkBoats.bid
       FROM Reserves, [All possible pink boats] PinkBoats
       EXCEPT
        [All (sid, bid) of reservations]
      )
   )
  ) R, Sailors S
 WHERE R.sid = S.sid;
  
```

```

#### STEP 1b
...
(SELECT Reserves.sid, PinkBoats.bid
FROM Reserves, [All possible pink boats] PinkBoats
EXCEPT
  [All (sid, bid) of reservations]
)
...

```

[All sailors in the reservation table X All pinkboats]

- [All (sailors, pink boat) existing reservations]

= [All (sailor_R, pink boats) combos that do not exist]

...(sailor_R are all sailors that have made a reservation)

Step 2

```

#### STEP 2
SELECT sname
FROM
  (SELECT sid
   FROM Reserves
   EXCEPT
     (SELECT sid
      FROM [All (sailor_R, pink boats) combos that do not exist]
     )
  )

```

*do not exist in Reservations

Step 3

```

#### STEP 3
SELECT sname
FROM
  (SELECT sid
   FROM Reserves
   EXCEPT
     All sailors_R that have not reserved all pink boats
  )

```

Step 4

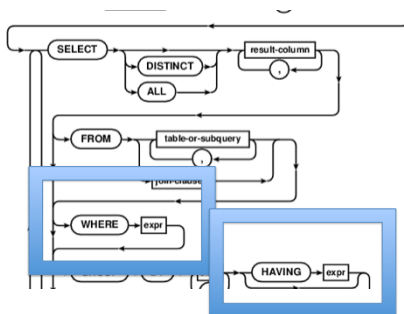
```

#### STEP 4
SELECT sname
FROM
  All Sailors that have reserved all pink boats

```

This actually only true sometimes (it has some implicit assumptions) – what happens when certain tables are NULL?


HAVING



Record id = <page id, slot #>.

FORMATS


Fixed vs Variable Length

RECORD FORMATS**Record Formats: Fixed Length**


- Field types same for all records in a file.
 - Type info stored separately in *system catalog*.
- Finding *i*'th field done via arithmetic like arrays

**Record Formats: Variable Length**

- Two alternative formats (# fields is fixed):




Second offers direct access to *i*'th field, efficient storage of *nulls* (special *unknown* value); small directory overhead.

PAGE FORMATS**Page Formats: Fixed Length Records**

Record id = <page id, slot #>.

In first alternative, moving records for free space management *changes rid*; may be problematic!

**"Slotted Page" Format: Variable Length Records**

Can move records on page without changing rid!
So, attractive for fixed-length records too.




Unordered (Heap) Files

- Collection of records in no particular order.
- As file shrinks/grows, disk pages (de)allocated
- To support record level operations, we must:
 - keep track of the *pages* in a file
 - keep track of *free space* on pages
 - keep track of the *records* on a page
- There are many alternatives for keeping track of this.
 - We'll consider 2




Heap File Implemented as a List



- Header page ID and Heap file name stored elsewhere
 - Database "catalog"
- Each page contains 2 "pointers" plus data
 - Problem for multi-page objects (blobs) – how to read blobs?



Better: Use a Page Directory



- Directory entries include #free bytes on the page.
- Directory is a collection of pages; linked list implementation is just one alternative.
 - *Much smaller than linked list of all HF pages!*
 - Can also point to groups of pages (say 64k chunks)